

Quiz 0

- some things to note
 - `xorq` an element with itself always give 0 (you can do this instead of `irmovq` if you want to clear the register)
 - note that it will set the condition code so might not be what you want
 - `andq` an element with itself is always itself (you can use this to check for the register's content and set the condition code for jump)

Corrections

Question 1: Interpreting Little Endian Values

- part 1

Question 1: Interpreting Little Endian Values

0xa348 48 44 6b 59 79 58 56 00

What is the value, in hexadecimal for the **little endian** `uint32_t` variable located at address 0xa34c?

0x	00565879	?	✓ 100%
----	----------	---	--------

- since it's little endian, you read the value backwards from right to left
- also since it's a `uint32_t` (4 bytes) - we only read 4 bytes
- note that the ordering of the individual bytes themselves stay the same (i.e it's 56 58 and not 65 85)

- part 2

What is the string that begins at address 0xa34c? (You might find [the ascii tables](#) handy.)

yXV\0	?	✗ 0%
-------	---	------

- note: endianness does not matter here as string is an array of chars/bytes
- the string that starts at `0xa34c` in bytes is `0x 79 58 56 00`
- translates into ascii and you get `yXV \0` but we do not print the null character at the end

Question 2: Interpreting Big Endian Values

- same thing as above but just big endian

Question 2: Interpreting Big Endian Values

0xd530 42 57 4b 49 6e 49 4c 00

What is the value, in hexadecimal for the **big endian** uint32_t variable located at address 0xd534?

0x 6e494c00 ? ✓ 100%

What is the string that begins at address 0xd534? (You might find [the ascii tables](#) handy.)

nIL\0 ? ✗ 0%

This question is complete and cannot be answered again.

- again, don't include the null character at the end

Question 3: Understanding sign extension

- question

Question 3: Understanding sign extension

Assume that the variable **k** is a signed char (`char`) and that the variable **w** is a signed 4-byte integer (`int32_t`), what is the value, in hexadecimal, of **w** after these C statements execute?

```
k = 0xed;  
w = k;
```

0x ffffffed ? ✓ 100%

This question is complete and cannot be answered again.

- so you are upcasting which in C requires **signed extension**
- since `e` in binary is `0b1110` so when extend, it'll be all `f`'s at the front
 - in fact, starting from `0xA` will have leading 1s

Question 4: Bit Shifting

- question

Question 4: Bit Shifting

Let **i** be a signed 4-byte integer (`int32_t`). What is the value, in hex, of **i** after this C statement executes?

```
i = 0x3c << 16;
```

0x 0x3C0000 ? ✓ 100%

This question is complete and cannot be answered again.

- shifting 16 bits is the same as shifting by **4 hex digits**
- `0x3c` in full is `0x 00 00 00 3c`
- so we remove the first 4 hex digits (which were all 0s anyways) and add 0s to the end

- becomes `0x 00 3c 00 00`

Question 5: Bit Manipulation

- this one was a bit of a tricky one - there are some things to note
 - the bit position is 0-indexed starting from the very right (i.e the rightmost bit is index 0)
 - the only answer that will work for this is ones that will make every other bit but the bit in question into 0 and only leave the star bit as the deciding bit between 0 and non zero values
 - we'll indicate the position in question with a star `*` and the other positions with an `X` as they can be both 0s or 1s
 - the breakdown of this question is kinda complicated, but you can do some of these in your head → I'll only go over the correct answers
- question

Question 5: Bit Manipulations

Consider a 4 byte unsigned variable `b`. Select all of the C expressions below that will evaluate to true if the bit at position 8 is 1 and will evaluate to false if the bit at position 8 is 0.

Recall:

- In C, any non-zero value evaluates as true.
- We number bits starting with 0 for the least significant bit, 1 for the next least significant bit, etc.

`((b << 7) & 0x8000)`

`((b << 7) & 0x8fff)`

`((b / 256))`

`((b >> 7) & 0x1) X`

`((b & 0x100) / 256)`

`((b * 16) & 0x1000)`

Select all possible options that apply. ?

○ 33%

- `b = 0b XXXXXXXX XXXXXXXX XXXXXXXX* XXXXXXXX`
- `((b << 7) & 0x8000)`
 - `b << 7 : 0b XXXXXXXX XXXXXXXX *XXXXXX XXXXXXXX`
 - `0x8000 : 0b 00000000 00000000 10000000 00000000`
 - `((b << 7) & 0x8000) : 0b 00000000 00000000 *00000000 00000000`
 - so if the star bit is "on" - we get a non zero value, if not we get a zero value
- `((b & 0x100) / 256)`
 - `0x100 : 0b 00000000 00000000 00000001 00000000`
 - `b & 0x100 : 0b 00000000 00000000 00000000* 00000000`
 - just the and itself would have worked but `/ 256` is the same as right shifting by $\log_2(256) = 8$ bits which just brings the star bit to the very end (index 0)
- `((b * 16) & 0x1000)`

- `b * 16` is the same as left shifting by $\log_2(16) = 4$ bits
- `b * 16 : 0b XXXXXXXX XXXXXXXX XXX*XXXX XXXX0000`
- `0x1000: 0b 00000000 00000000 00010000 00000000`
- `((b * 16) & 0x1000) : 0b 00000000 00000000 000*0000 00000000`

Question 6: Passing by Value and Reference

- question: Consider the following program:

```

1 void f1(int16_t *arg) {
2     *arg = 1107;
3 }
4 void f2(int64_t arg) {
5     arg = 3402;
6 }
7
8 int main(int argc, char *argv[]) {
9     // Suppress compiler warnings
10    (void) argc;
11    (void) argv;
12
13    int16_t ckiuc = 13;
14    int64_t avyvyl = 88;
15    f1(&ckiuc);
16    f2(avyvyl);
17    printf("Line A: %hi\n", ckiuc);
18    printf("Line B: %ll\n", avyvyl);
19 }
```

- part 1

What number will be printed after "Line A"?

1107



✓ 100%

- the function `f1` actually changes the number by dereferencing it

- part 2

What number will be printed after "Line B"?

88



✓ 100%

- the function `f2` doesn't actually change anything (something about arguments being passed by pointer so you've really just changed the pointer)

Question 7: Truncate a String

- question

Question 7: Truncate A String

Consider the following code.

Your job is to determine the appropriate values for `x` and `y` in line 3, so that the `printf` in line 4 outputs "Hello CPSC " (without the enclosing quotes; we put them there to make clear if there were any blanks).

Hint: There are multiple correct values for `y`. We will award full marks only for an answer that compiles without warnings under all conditions.

```
1 char str[] = "Hello CPSC 313 students.";
2 // what are the correct values for x and y so that it prints "Hello CPSC "
3 str[x] = y;
4 printf("%s", str);
5 return 0;
```

What is the value of `x`?

x 11 ? ✓ 100%

What is the value of `y`?

y \0 ? ✗ 0%

- you want to cut off the string at index 10, so you need to place a terminating string at index 11
- the answer for `y` is technically correct but they wanted explicitly '\0' (with the quotes)

Question 8: C Types

- question

Consider the following C declarations

```
char *buf;
uint8_t v1, *v2, **v3;
int8_t i1, *i2, **i3;
short e1, *e2, **e3;
int64_t n1, *n2, **n3;
```

For each expression below, pick the correct type.

(uint8_t *)buf pointer to a uint8_t ✓ 100%

&i2 pointer to a pointer to a int8_t ✓ 100%

*(short *)buf short ✓ 100%

**n3 int64_t ✓ 100%

- first: you cast it as a `uint8_t` pointer now
- second: you address a pointer so it becomes a pointer to a pointer
- third: you cast it to a `short` pointer, then dereference it, so it'll be a `short`
- fourth: it's a double pointer so first dereference you get pointer to `int64_t`, dereference that again and you get the `int64_t` itself

Question 9: Variable-sized Structures

- question

Question 9: Variable-sized Structures

Consider the following structure:

```
#define MAXBYTES      55

struct var {
    uint32_t nbytes;
    char array[MAXBYTES];
};
```

The `nbytes` field indicates how many bytes in the `array` contain valid data.

- part 1

The `nbytes` field indicates how many bytes in the `array` contain valid data.

What is the value of `sizeof(struct var)`?

bytes   100%

- `nbytes` is 4 bytes long (4-aligned)
- `array` is 55 bytes long (1-aligned)
- since the whole struct needs to be 4 aligned, we need to round 59 to 60

- part 2

Now, let's assume that we want to represent instances of this structure in as few bytes as possible. To do so, we want to allocate only enough space to hold the number of valid bytes in `array`, but still allowing us to pack the structures one after the other, while ensuring that each structure is properly aligned.

If the `nbytes` field contains the value 36, how many bytes must we allocate for the structure?

bytes   0%

- basically we don't allocate array of max size anymore, but only allocate size of `nbytes`
- `nbytes` is still 4 bytes long
- `array` is now 36 bytes long
- so in total it's 40 because it's already 4 aligned

- part 3

If `structPtr` is currently pointing at one structure, which of the following expressions will set `structPtr` to the address of the next structure in the buffer? (You may assume that the function `structSize(uint32_t nbytes)` properly computes the size of a structure when passed the value in the `nbytes` field as its parameter.

- (a) `structPtr = (struct var *)(((uint8_t *)structPtr + structSize(structPtr->nbytes)));` ✓
- (b) `structPtr = (structPtr + structSize(structPtr->nbytes));`
- (c) `structPtr = (structPtr + structPtr->nbytes);`
- (d) None of the above

✓ 100%

- first one: you cast `structPtr` to a `unit8_t` so any operation done on it will move by 1 byte only, then you add it by the size of the struct, so it'll move it by 40 bytes - which is correct
- second one: `structPtr` is already a `struct*` so adding 40 to it will move it `40 * 40` bytes
- third one: same problem as above it'll just move it by `36 * 40` bytes instead

Question 10: C Printfs

- question

Question 10: C Printfs

Consider the following C code snippet:

```
char *p;  
p = "b5bC";
```

Suppose `p` is at address 0x6598 and the string "b5bC" is at address 0x19a40. What will each `printf` statement below output?

Do not include extraneous punctuation, i.e., if `printf` will output quotation marks, include them; if it will not output quotation marks, do not include them.

```
printf("%p", *p);
```

0x19a40

?

✓ 100%

```
printf("%p", p);
```

0x19a40

?

✓ 100%

```
printf("%c", *p);
```

b

?

✓ 100%

```
printf("%c", *p);
```

b

?

✓ 100%

This question is complete and cannot be answered again.

- first: `%p` prints the pointer address, in this case we get the address of `p`, then dereference it, which gets us the value of `p` - the value of `p` is the address of the string `0x19a40`

- second: printing the value of `p` again → same above
- third: `%c` prints a character, we're deref `p` so we're accessing the string - the first char is `b`
- fourth: we deref `p`, get the address of that (so `p` again), then deref it again, so just basically the same as deref `p`
→ same as above

Question 11: Pointer Figure Question

- question

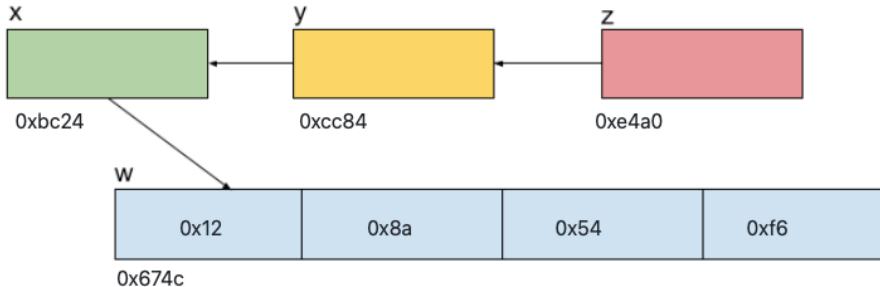
Question 11: Pointer Figure Question

In the figure, we have the following variables:

- `int w[4]; // An integer array with four elements`
- `int *x = w; // A pointer to the array w.`
- `int **y = &x; // A pointer to x.`
- `int ***z = &y; // A pointer to y.`

Assume that `sizeof(int)` is 4.

Values in boxes represent the values stored in memory; values under boxes represent addresses.



What is the value, in hex, of `x`?

0x 674c ? ✓ 100%

After executing the statement `++x`, what is the value of `***z`?

0x 8a ? ✓ 100%

This question is complete and cannot be answered again.

- `x` points to an array - so `x` itself holds the start address of the array - which is `0x674c`
- when you increment `x`, you move it by 8 bytes (assuming ints are 8 bytes), so now it's pointing to the second element of the array
 - you deref `z` three times and it's the same as deref `x` which is `0x8a` now

Question 12: C Pointer Operations

- question

Question 12: C Pointer Operators

Consider the following C code snippet:

```
int eep = 50;  
int *hf = &eep;
```

What do the following expressions refer to?

If an expression is invalid, select **Invalid**.

If an expression is valid but does not have a corresponding answer, select **None of the above**.

You may use answers more than once or not at all.

- | | |
|---|----------------------|
| c <input type="checkbox"/> *eep <input checked="" type="checkbox"/> | a. Value of eep |
| a <input type="checkbox"/> eep <input checked="" type="checkbox"/> | b. Address of eep |
| d <input type="checkbox"/> &hf <input checked="" type="checkbox"/> | c. Invalid |
| a <input type="checkbox"/> *hf <input checked="" type="checkbox"/> | d. Address of hf |
| | e. None of the above |

✓ 100%

This question is complete and cannot be answered again.

- only interesting one here is `*eep` since it's not a pointer, you're trying to access memory address 50 which may be invalid

Quiz 1

In Class Assignment

In Class 3

- I3.1: y86 Address of Next Instruction
 - variation 1

I3.1. y86 Address of Next Instruction

If you are executing an addq instruction at address `0x1405`, what is the address of the next instruction to be executed?

If there is not enough information provided to answer the question, enter 0.

0x



Save & Grade Single attempt

Save only

Additional attempts available with new variants ⓘ

- `addq` is 2 bytes long so you just add that
- result is `0x1407`

- variation 2

I3.1. y86 Address of Next Instruction

If you are executing a jge instruction at address `0x1b8d`, what is the address of the next instruction to be executed?

If there is not enough information provided to answer the question, enter 0.

0x ?

Save & Grade Single attempt **Save only** Additional attempts available with new variants ⓘ

- there's not enough information to tell us whether we will jump or not → if we do jump, the address is whatever `valc` is, if not, the address is `0x1b8d + 9`
- so the answer is 0 (because we don't know)

- I3.2: y86 Address Instruction after Jump instructions

- variation 1

I3.2. y86 Address Instruction after Jump instructions

If you are executing the following instruction:

`jle 0x34a8`

at address `0x1ff`, what is the address of the next instruction you will execute when the jump is not taken?

If there is not enough information provided to answer the question, enter 0.

0x ?

Save & Grade Single attempt **Save only** Additional attempts available with new variants ⓘ

- jump is not taken so we just go to the instruction after the jump (so `0x1ff + 9`)
- answer is `0x208`

- variation 2

I3.2. y86 Address Instruction after Jump instructions

If you are executing the following instruction:

`jg 0x3846`

at address `0x11d`, what is the address of the next instruction you will execute when the jump is taken?

If there is not enough information provided to answer the question, enter 0.

0x ?

Save & Grade Single attempt **Save only** Additional attempts available with new variants ⓘ

- jump is taken so you jump to `0x3846`

- I3.3: y86 Register File Read/Write

- question

I3.3. y86 Register File Read/Write

Typically **rA** and **rB** indicate which registers an instruction reads or writes. Which of the following instructions must read or write a register that is not specified in either **rA** or **rB**?

For the purposes of this question, the PC is a register, but select it only if the instruction updates it with a value other than that of the instruction that immediately follows the current instruction in memory.

mulq
 call ✓
 ret ✓
 rrmovq
 popq ✓
 mrmovq

Select all possible options that apply. ?

✓ 100%

[Try a new variant](#)

- these instructions manipulate the stack pointer **%rsp** which is only implied and not specified
- **jxx** would also count here since it could potentially set the PC with value that's NOT just **valP**

- I3.4: y86 Sources of Register Updates

I3.4. y86 Sources of Register Updates

When the y86 writes values into registers, there are multiple places on the processor from which those values might come. Which of the following might produce a value that will get written into the register file?

From the ALU
 From the PC
 From the condition codes
 From data memory

Select all possible options that apply. ?

[Save & Grade Single attempt](#) [Save only](#) Additional attempts available with new variants ?

- answer is from the ALU and from data memory
- if another option was "From the register file" - this would also be WRONG
 - because even values from register values are passed through the ALU - we just add 0 to it so it's basically ALU output by the time it gets written into registerse

- I3.5: y86 ALU for Other Instructions

- part 1

I3.5. y86 ALU for Other Instructions

1. It seems pretty clear that you will need to use the ALU to implement arithmetic and logical operations (e.g., addq, subq, mulq, divq, modq, xorq, andq), but perhaps we can use the ALU to perform other operations as well.

Which of the following instructions might also benefit from being able to use the ALU? (Select all that apply.)

Note: we do not want to use the ALU to compute the address of the next instruction, because a) that would mean that we need the ALU to do two different things on arithmetic/logical instructions and that's not possible, and b) the ALU is a more complex circuit than that necessary to calculate the next value of the PC.

- nop
- jump instruction (j, jle, jg, jge, jl, jeq, jne)
- pushq/popq ✓
- mrmovq/rmmovq ✓
- call, ret ✓
- halt

Select all possible options that apply. [?](#)

✓ 100%

- `call/ret` and `pushq/popq` works with the stack so they increment and decrement by 8
- `mrmovq/rmmovq` feeds the values through and add 0 to it

- part 2

2. If we use the ALU to implement these other instructions, should they set the condition codes?

- (a) No ✓
- (b) Yes

- no, only `opq` should set the condition codes

- part 3

3. Which of the following are potential inputs for one of the two ALU inputs AluA and AluB?

- (a) PC
- (b) icode
- (c) CC
- (d) 0 ✓
- (e) +/- 8 ✓
- (f) +/- 16

Select all possible options that apply. [?](#)

✓ 100%

- I3.6: y86 Instructions Memory Read/Write

- question

I3.6. y86 Instructions Memory Read/Write

Which of the following instructions read from memory (excluding the bytes representing the instruction)?

- popq
- ret
- rmmovq
- halt
- rrmovq

Select all possible options that apply. [?](#)

[Save & Grade Single attempt](#)

[Save only](#)

Additional attempts available with new variants [?](#)

- **popq/ret** both reads from the stack (which is in memory)
- note: if they didn't have the notes saying "excluding the bytes representing the instructions" then all instructions require a read from memory to fetch the instruction itself

In Class 4

- I4.1: y86 Buffer Overflow: Getting Started

Step 1

1. Starting with the file below, write instructions to create an infinite loop in the exploit function.
2. Add a **nop** so that when you run your code in the simulator, you can visually confirm that it is running.
3. Copy your code into the simulator and run it to confirm that it works.

```
1 # Function that simply loops forever
2 .pos 0x1000
3 exploit:
4   nop
5   jmp exploit
```

- we just make it jump back to the start forever
- the **nop** is to make the jumping a bit more obvious

- I4.2: Reading y86 code for the Buffer Overflow Attack

Step 2: Warmup: Reading the y86 Code

You will see that the file below, **01-exploit2.ys** contains a driver program and some data to exercise a function named mystery. You may find it useful to copy this into the simulator and run it to help you perform the following steps.

1. Add comments describing what the mystery function is doing.
2. Rename mystery to a more descriptive name.
3. Save the file.

```
1 # This function DOES WHAT?
2 # Input %rdi
```

```

3 # Input %rsi
4 # Output 0
5
6 # THIS IS ALL PART OF THE DRIVER PROGRAM
7 # Initialize the stack
8 irmovq 0x5000, %rsp
9
10 # Driver program to test mystery
11 # Note that we pass parameters to functions
12 # in registers %rdi and %rsi (in order)
13 irmovq a, %rdi    # a is the first parameter
14 irmovq b, %rsi    # b is the second parameter
15 call copy        # mystery(int[2] a, int[2] b)
16 halt
17
18 a:      .quad 0xDEADBEEF
19         .quad 0
20
21 b:      .quad 0
22         .quad 0
23 # END OF DRIVER PROGRAM
24
25 .pos 0x2000
26 mystery:
27     irmovq 8, %r10          # put 8 into r10 (will likely be an index incrementer)
28 loop:
29     mrmovq 0(%rdi), %rax   # put a[i] (dereference) into rax -> rax = *rdi
30     rmovq %rax, 0(%rsi)    # move rax into address of rsi -> *rsi = rax (copy value of
31     an index in a into the same index at b)
32     andq %rax, %rax       # it does nothing to rax -> BUT IT SETS THE CONDITION CODE
33     je done                # basically, if rax = a[i] == 0 then we're done
34
35     addq %r10, %rdi        # increment the index -> rdi++ or rdi += 8
36     addq %r10, %rsi        # increment the index -> rsi++ or rsi += 8
37     jmp loop               # jump
38 done:
39     ret

```

- I4.3. Writing the Exploit

Step 3: Writing the Exploit

In this step you are going to write an exploit where a benign program (the caller) calls a function, but that function is malicious. When the function returns, the processor will be stuck in an infinite loop.

Copy the code from the file below into the simulator.

1. Use the simulator to run this program.
2. To what address (in hex) will the function `evil` return? (I.e., what is the address of the instruction that should execute immediately after the `return` statement.)

0x 13 ✓100%

3. Where is that address stored? (call this the return address)

0x 0000000000 ✓100%

4. Now, copy the code you wrote in Step 1 (your infinite loop) into the simulator, adding it to the program you already have.
5. **Add lines of code to the `evil` function that will overwrite the return address with the address of the exploit function.** (This should take only 2 instructions.)
6. If you have done everything correctly, the simulator should end up in an infinite loop.
7. Save the code that you've modified in the simulator into your editor window below and submit it.

```
1 # Attack 1: A benign caller invokes a function
2 # who is malicious; upon return the processor
3 # will be in an infinite loop.
4
5 # Initialize the stack
6 irmovq 0x2000, %rsp
7
8 main:
9   call evil # call function
10  halt    # If we return
11    # successfully, halt
12
13 # Return address is just above the stack pointer:
14 evil:
15   # WRITE YOUR CODE HERE
16   irmovq  exploit, %rax
17   rmmovq  %rax, 0(%rsp)
18   ret
19
20 # Place the code you wrote in Step 1 here
21
22 # We put this here so you can observe the contents
23 # of the stack (why is this at 0x1FF0 and not 0x2000?)
24 .pos 0x1FF0
25 .quad 0x0
26 .quad 0x0
27
28 .pos 0x1000
29 exploit:
```

```
30    nop
31    jmp exploit
```

- the stack pointer was pointing at the return address for the caller, and we've overwritten that

In Class 5

- refer to the `ic5_logic_blocks.c` file

In Class 6

- refer to the notes

Corrections

Question 1: Using and changing the PC

- part 1

All instructions use and can change the PC in some way (i.e., because we read the instruction at the PC's address and then we update the PC to refer to the next instruction).

Which of the following may *use the PC's existing value* in some other way than these?

- opq
- ret
- rmmovq
- call ✓
- mrmovq

Select all possible options that apply.

✓ 100%

- `call` is the only instruction that uses the PC in a different way as it pushes the PC onto the stack

- part 2

Which of the following may *set the PC's value* in some other way than these?

- rrmovq
- mrmovq
- rmmovq
- jmp/jxx ✓
- call

Select all possible options that apply.

✗ 0%

- `jmp/jxx` is correct because we set the PC value using what's specified in the instruction (i.e. `valC`)
- in the same vein, `call` also sets the PC using the function address specified in the instruction

Question 2: Jump conditions and condition codes

- part 1

The jump instructions use the opcode 7. In the table below, the *first* byte of a jump instruction is given on the left of each row. For each combination of the condition codes ZF and SF given in the scenario columns, indicate whether the jump would be taken for those values of the ZF and SF flags.

First Byte	Scenario 1			Scenario 2			Scenario 3		
	ZF	SF	Branch Taken?	ZF	SF	Branch Taken?	ZF	SF	Branch Taken?
0x72	0	0	No	0	1	Yes	1	0	No
0x75	0	0	Yes	0	1	No	1	0	Yes
0x73	0	0	No	0	1	No	1	0	Yes
0x70	0	0	Yes	0	1	Yes	1	0	Yes

- you were supposed to fill in the "Branch Taken?" part yourself - but it is pretty simple logic

Question 3: y86 Condition Codes: Affected Instructions

- part 1

Which of the following instructions will behave differently depending on the value of the condition codes?

- pushq
- modq
- irmovq
- rrmovq
- jl ✓
- divq

Select all possible options that apply. ?

✓ 100%

- easy option here but overall nstructions that depend on the values of the condition codes are conditional instructions: conditional jump instructions (`jl`, `jlt`, `je`, `jne`, `jg`, `jge`) and condition move instructions (`cmovl`, `cmove`, `cmovne`, `cmovg`, `cmovge`)

Question 4: Instructions and registers

- part 1

Considering only the registers that can be named (so not including the PC), which of the following instructions can result in a change to more than one register's value?

- pushq x
- call x
- popq ✓
- rrmovq

Select all possible options that apply. ?

✗ 0%

- only `popq` changes > 1 registers value (it changes the destination register `rA` as well as `%rsp` since we increment the stack pointer)
- wrong:
 - `pushq` changes the stack pointer `%rsp` only
 - `call` does change `%rsp` (because we push return address to the stack) as well as PC but PC is said to not be included so it only technically changes 1

- part 2

Considering only the registers that can be named (so not including the PC), which of the following instructions read at least one register's value?

- irmovq
- cmovxx ✓
- call ✓
- jxx

Select all possible options that apply. ?

✓ 100%

- `cmovxxx` is a register to register move, so it needs to read from the source register to write to the destination register
- `call` needs to write the return address to the stack, so it reads the stack pointer `%rsp` to decrement

- part 3

Considering only the registers that can be named (so not including the PC), which of the following instructions refer to at least one register as part of the instruction?

- rmmovq ✓
- pushq ✓
- call
- jxx

Select all possible options that apply. ?

✓ 100%

- this question is asking which instructions require you to **explicitly specify** a register
 - so `call` there's an implicit register which is `%rsp` but that doesn't count
- `rmmovq` you have to specify `rA` and `rB`
- `pushq` you have to specify a source register `rA`

Question 5: ISA Basics

- part 1

How many valid opcodes are there in the y86 ISA?

13

opcodes



x 0%

- `opcodes` here refers to `icode` (so doesn't include the `ifunc`) of which there are 12 (0 to 11)



Question 6: y86 Calling Conventions

- part 1

A y86 program calls a function with zero arguments. Assuming that the function adheres to the y86 calling conventions, which of the following registers' values may have changed when the function returns?

%r13

%rbp

%r12

%rbx

%r14

%rax ✓

Select all possible options that apply.

✓ 100%

- basically anything that's a "scratch" register or a "caller-saved" register
- most options here are callee-saved

Question 7: y86 Register Usage

- part 1

A y86 program calls a function with six arguments. Assuming that the function adheres to the y86 calling conventions, which of the following registers are guaranteed to contain the same value when the function returns as they did when the function was called?

- %rdx
- %r8
- %rcx
- %rax
- %rdi
- %r12 ✓

Select all possible options that apply. [?](#)

✓ 100%

- basically the callee-saved registers

Question 8: Parameter Passing

- part 1

A y86 program calls a function with four integer arguments. Assuming that the function adheres to the y86 calling conventions, which of the following registers will be used to pass parameters?

- %rsi ✓
- %rcx ✓
- %rbp
- %r9
- %rbx
- %r13

Select all possible options that apply. [?](#)

✓ 100%

- in order registers used to pass parameters are: %rdi, %rsi, %rdx, %rcx, %r8, %r9
- but we only have 4 args so we can only use %rdi, %rsi, %rdx, %rcx and only 2 of those are here
- note: if we had %rsp or %rbp - that would also be one we consider a callee-saved register

Question 9: Interpreting execution notation

- part 1

The y86 architects have decided to add a new instruction to the y86. The instruction is written as: `call* D(rB), rA` and its behaviour is:

- `call*` treats `rB` as the index in a table of addresses.
- The table starts at address `D` and is 0-indexed.
- `call*` calls the function whose address is at the `rBth` index into the table.
- The return address is stored in register `rA` instead of on the stack.

Both in the text and in the videos that introduce each instruction, we used a formal notation to precisely describe the behaviour of an instruction. Using that notation, select the option that precisely matches the description above.

- `rA ← PC + 10`
`PC ← D + M8[8 * rB]`
- `R[rA] ← PC + 10`
`PC ← M8[D + 8 * R[rB]]`
- `rA ← PC + 10`
`PC ← D + 8 * rB`
- `rA ← PC + 10`
`PC ← M8[D + rB]`
- `R[rA] ← PC + 10`
`PC ← D + M8[8 * R[rB]]`
- `rA ← PC + 10`
`PC ← M[D + 8 * rB]`

 100%

- firstly, process of elimination
 - since return address is stored in `rA` - `rA` would have to take on the value of `PC + size(instruction) = PC + 10` (because `call` is 10 bytes long)
 - we know proper notation to store something into a register is `R[ra] <- val` so anything that doesn't have this is eliminated (1, 3, 4, 6)
- now we only have 2 and 5
 - 2 makes the most sense since we want to read start from D so D has to go within `M[]` clause
- note: `M8[addr]` means read 8 bytes starting from `addr`

Question 10: Understanding y86 Programs

- part 1

In this problem, we're introducing a new instruction into the y86 architecture: `MAXQ rA, rB` which assigns the larger of `R[rA]` and `R[rB]` to `R[rB]`.

Consider the following y86 program

```

1  irmovq  0x5000, %rsp  # initialize the stack
2  irmovq  4, %rbp       # rbp = 4
3  irmovq  1, %rbx       # rbx = 1
4  irmovq  8, %r13        # r13 = 8 (increment size?)
5  xorq   %r14, %r14      # clean out r14?? -> r14 will be the pointer

```

```

6    xorq %r12, %r12      # clean out r12??
7  loop:
8    rrmovq %rbx, %rdi      # rdi = rbx
9    subq %rbp, %rdi        # rdi = rdi - rbp
10   irmovq 1, %rax         # rax = 1
11
12   jg done                # if (rdi - rbp) > 0 => if rdi > rbp then done
13   subq %rax, %rbp        # rbp -= 1 (since rax = 1 right now)
14   # this portion above is an insanely complicated way to traverse the list
15   # we do (rdi - rbp) every time where rbp is shrinking by 1 each time
16   # since rdi < rbp at the start, we go until rdi > rbp basically
17   # overall, we will run *the loop* (the actual work) %rbp time
18
19   mrmovq A(%r14), %rdi    # rdi = A[r14]
20   call check               # call check which returns rax = max(A[r14], 5)
21   addq %rax, %r12          # r12 += rax
22   addq %r13, %r14          # r14 += r13 += 8 => increment the pointer
23   jmp loop                 # loop
24 done:
25   rrmovq %r12, %rax       # result = r12, return result
26   halt
27
28 check:
29   irmovq 5, %rax
30   MAXQ %rdi, %rax
31
32   # this is my own code to represent the max function
33   rrmovq %rax, %r8        # r8 = rax
34   subq %rdi, %r8           # rax - rdi
35   cmovl %rdi, %rax        # rax = rdi if rax - rdi < 0
36
37 ret
38
39 .pos 0x2000
40 A:
41   .quad 8
42   .quad 2
43   .quad 2
44   .quad 10

```

- basically, the code is traversing through the array `A` and adding `max(A[i], 5)` to `res` and returning that

```

1 res = 0
2 for i in range(length): # length is rbp
3     res += max(A[i], 5)
4 return res

```

- part 1

How many times will line 8 execute?

4 ? ✘ 0%

- at the start, `rbp = 4` and `rdi = 1` and we're doing `rdi - rbp` everytime
 - first iteration: $1 - 4 = -3$
 - second iteration: $1 - 3 = -2$
 - third iteration: $1 - 2 = -1$
 - fourth iteration: $1 - 1 = 0$
 - fifth iteration: $1 - 2 = 0$
 - so basically, the content of the loop itself only run 4 times (and thus `rbp = number of elements we want to read`) but the loop condition check runs 5 times, the last time it will succeed and we will jump to `done`
 - correct answer is 5

- part 2

What value (in decimal) will be in register %rax when this function halts?

28 ? ✓ 100%

- look at the python code above, we'll get $\max(8, 5) + \max(2, 5) + \max(2, 5) + \max(8, 10) = 28$

- part 3

If we changed line 25 to `rrmovq %rdi, %rax`, what would happen to the value in %rax at the end of the program?

It would get smaller. ✓
 It would get larger.
 It would be unchanged.
 It is impossible to tell.

✓ 100%

- then we're basically just summing up the array
- $8 + 2 + 2 + 10 = 24$ so <u>**smaller**

- this bitch was hard as fuck

Question 11: Stack Smash

- part 1

Here is a C signature and documentation for the hand-written assembly `arrayincr` function below:

```
// Given index, k, and array, performs: array[index] += index*k.  
void arrayincr(uint64_t index, int64_t k, int64_t *array);
```

`arrayincr` is called in the assembly program below by `main`. You will supply values for `%rdi` and `%rsi` in order to force the call to `arrayincr` to return to `evilfun` rather than `main`.

Hint: Violate an assumption that `arrayincr` makes about its parameters to create unexpected behaviour.

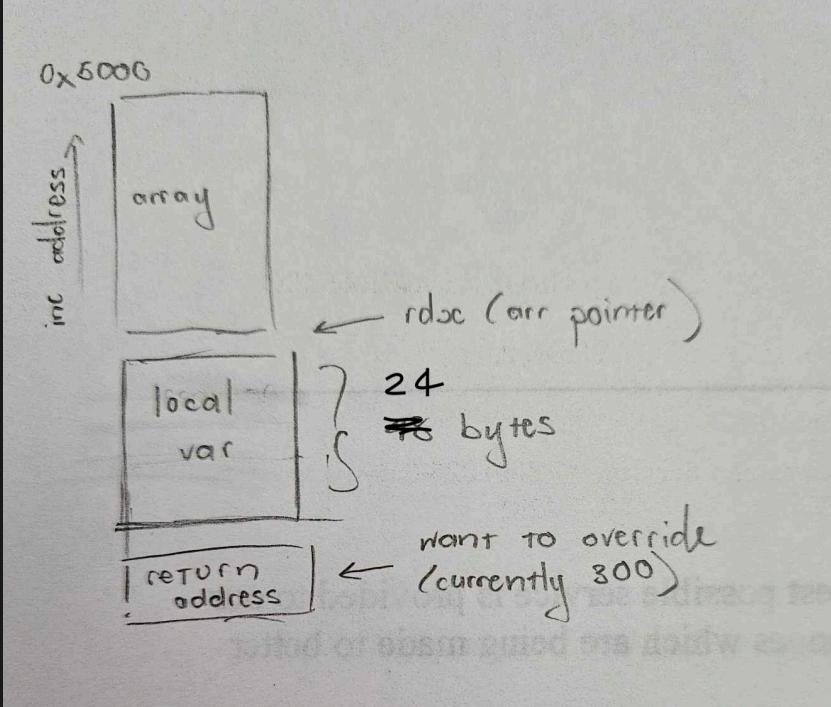
Assembly program:

```
1 .pos ????
```

Intentionally left out

```
2  
3         irmovq 0x5000, %rsp      # Initialize the stack pointer  
4 main:    irmovq 0x48, %rax  
5         subq %rax, %rsp      # Make space for an array of 9 8-byte quantities  
6         rrmovq %rsp, %rdx      # Initialize the array pointer  
7         irmovq 0x18, %rax  
8         subq %rax, %rsp      # Make space for another local variable of 24 bytes.  
9         irmovq $____, %rdi      # Select the index. VALUE SUPPLIED BELOW.  
10        irmovq $____, %rsi      # Select k. VALUE SUPPLIED BELOW.  
11        call arrayincr  
12        irmovq 0x60, %rax      # ASSUME the address of this instruction is 300.  
13        addq %rax, %rsp  
14  
15 arrayincr: irmovq 8, %r8      # constant 8  
16          mulq %rdi, %rsi      # compute index*k  
17          mulq %r8, %rdi      # compute the offset to array[index]  
18          addq %rdi, %rdx      # compute array + offset  
19          mrmovq 0(%rdx), %rcx      # Load value at address %rdx  
20          addq %rsi, %rcx      # compute array[index] + index*k  
21          rmmovq %rcx, 0(%rdx)      # store updated value back into array[index]  
22          ret  
23  
24  
25  
26 .pos 1500  
27 evilfun: # Goal is to force a ret from arrayincr into this function  
28         nop  
29         jmp evilfun
```

- basically, the program first initialize an array on the stack → array pointer is in `%rdx`
- then `arrayincr` will access that and based on `i` and `k` we will modify the array
- diagram



- what we want to do is override the return address (that's current 300) - we want to make it 1500
- the usual assumption is that with a positive index, we will index `arr` on the stack, however, with a negative index, we can access local var and beyond to get to the return address
 - so we go down 24 bytes to get the beginning of the return address, but remember that we write to memory (upwards) - so we need to go down an additional 8 bytes to the end of the return address
 - since the indices are multiplied by 8 → this means that `i = -4`
- now we need to actually make that value that's currently 300 to 1500 we have

$$\begin{aligned}
 \text{arr}[i] &= \text{arr}[i] + (i \times k) \\
 \underbrace{1500}_{\text{new addr}} &= 300 + (-4k) \\
 4k &= -1200 \\
 k &= -300
 \end{aligned}$$

- so the final answer is `%rdi = -4, rsi = -300`

Question 12: Iterating through an array

- part 1

The following y86 function sums up all of the elements of an array whose value is larger than that of a given threshold. The sum is placed in register %rax. **The function does not follow the y86 calling conventions** (to avoid giving away answers to other questions on this quiz).

The function takes three parameters:

- %rdi contains the address of the first element of the array.
- %r8 contains the number of values in the array.
- %rcx contains the threshold value.

Use the dropdowns to fill in an appropriate instruction names for OPCODE1 and OPCODE2, and appropriate instructions for INSTRUCTION1, INSTRUCTION2 and INSTRUCTION3.

```

1 addarray: xorq    %rax, %rax
2             irmovq 8, %r9
3             rrmovq %r8, %r11
4             mulq   %r9, %r11
5 loop:      OPCODE1 %r11, %r11
6             INSTRUCTION1
7             INSTRUCTION2
8             rrmovq %r11, %rsi
9             addq   %rdi, %rsi
10            mrmovq 0(%rsi), %rdx
11            rrmovq %rdx, %r10
12            OPCODE2  %rcx, %r10
13            INSTRUCTION3
14            addq   %rdx, %rax
15            jmp    loop
16 end:       ret

```

Select None for any unused register values.

You may use either a symbolic name, a DECIMAL value, or blank in the offset.

Location	Instruction	rA	rB	offset
OPCODE1	andq			✓100%
INSTRUCTION1	je	None	None	end ? ✓100%
INSTRUCTION2	subq	%r9	%rdi	Leave blank if no off ? ✓100%
OPCODE2	subq			✓100%
INSTRUCTION3	jle	None	None	loop ? ✓100%

This question is complete and cannot be answered again.

- **OPCODE1 :**
 - since it's a operation between `%r11` and itself, this is a very common pattern to AND a register with itself so that the value within the register itself remains the same, but it'll set the ALU flag
- **INSTRUCTION1 :**
 - it stands to reason that after setting the ALU flag, we would want to jump somewhere - and since it's at the beginning of the loop, makes sense that it would want to jump to the end
 - what kind conditional jump we do needs additional context (see below for instruction 2)
- **INSTRUCTION2 :**

- so the very first iteration in the loop, `%r11 = number elements * 8 = size of array`
 - and later on in the code, we actually basically try to read from `start_address + r11`
 - so if `start_address = 0` and `r11 = 32` (total size), we would try to read from `M[32]` which is actually invalid because that's past the space we've allocated
 - at this point it's important to realize that we're going to be accessing the array backwards
 - so it makes sense that we want to decrease `r11` by 8 each time (even on the first iteration so we don't go out of bounds)
 - `r9` has 8
 - so the solution is `subq r9, r11`
- **OPCODE2 :**
- we know that `%r10 = %rdx = element accessed at m[%rsi]`
 - from the question, we only want to add this element to `%rax` only if it's larger than a threshold
 - so we basically want to compare `r10 > rcx => r10 - rcx > 0` so we want to do `subq`
- **INSTRUCTION3 :**
- from the formula above, if `r10 - rcx < 0` meaning `r10 < rcx`, we don't want to add - so we return to the beginning of the loop early
 - hence we use `jle`
 - if we don't jump early, we'll end up adding the element to the solution `rax`

Quiz 2

In Class Assignment

In Class 7

- I7.1: Pipeline Efficacy

I7.1. Pipeline Efficacy

Pipelining will be most effective when:

The processor has a large number of registers.
 The maximum delay of a pipeline stage is significantly larger than the register delay. ✓
 There are thousands of pipeline stages.

✓100%

[Try a new variant](#)

- I suppose it's because we won't "feel" the overhead of the pipeline registers so much
- I7.2: Pipelining Benefits

I7.2. Pipelining Benefits

Select the statements that are **always** true.

Pipelining...

- Makes it possible for more than one instruction to be in-flight at any one time ✓
- Makes it possible for a processor to fetch multiple instructions from memory at the same time
- Is the only way to achieve parallelism in hardware
- Increases instruction throughput ✓
- Enables different parts of the processor to be used concurrently ✓
- Increases instruction latency ✓

Select all possible options that apply. ?

✓ 100%

[Try a new variant](#)

1. many instructions can be in the pipeline at once (in flight)
2. not true, we can only fetch 1 instruction at a time, but we can do it more often now instead of waiting for the entire cycle
3. not true, we can do multi-core stuff as well
4. this is just true
5. yes, we are allowing all the different stages and parts of the processor to be used at all time
6. this is also just true

- I7.3: Sequential Latency

I7.3. Sequential Latency

Assume a processor architecture with 4 stages with execution times of 50 ps, 65 ps, 40 ps, 55 ps and whose register delay time is 10 ps.

What is the latency of a single instruction in a **sequential** implementation of this architecture?

integer ps ?

[Save & Grade](#) [Single attempt](#)

[Save only](#)

Additional attempts available with new variants ?

- you incur the register delay time once at the beginning

$$\text{latency} = 10 + 50 + 65 + 40 + 55 = 220 \text{ ps}$$

- I7.4: Pipelined Latency

I7.4. Pipelined Latency

Assume a processor architecture with 4 stages with execution times of 65 ps, 45 ps, 40 ps, 75 ps and whose register delay time is 7 ps.

What is the latency of a single instruction in a **pipelined** implementation of this architecture?

integer ps ?

Save & Grade Single attempt

Save only

Additional attempts available with new variants ?

- you incur the pipeline latency at every stage AND you incur the maximum stage delays every time

$$\text{latency} = (7 \times 4) + (4 \times 75) = 325 \text{ ps}$$

- I7.5: Sequential Retire

I7.5. Sequential Retire

Assume a processor architecture with 6 stages with execution times of 70 ps, 60 ps, 45 ps, 50 ps, 75 ps, 65 ps and whose register delay time is 11 ps.

Let's say that you have two **andq** instructions; the first instruction completes execution at time T ps. In a **sequential** implementation, at what time will the second **andq** instruction complete?

T + integer ps ?

Save & Grade Single attempt

Save only

Additional attempts available with new variants ?

- since it's sequential, you need to go through all the stages from the beginning, so it's
 $11 + 70 + 60 + 45 + 50 + 75 + 65 = 376$

- I7.6: Pipelined Retire

I7.6. Pipelined Retire

Assume a processor architecture with 6 stages with execution times of 65 ps, 40 ps, 55 ps, 45 ps, 60 ps, 50 ps and whose register delay time is 7 ps.

Let's say that you have two **mulq** instructions; the first instruction completes execution at time T ps. In a **pipelined** implementation, at what time will the second **mulq** instruction complete?

T + integer ps ?

Save & Grade Single attempt

Save only

Additional attempts available with new variants ?

- since it's pipelined, by the time the first instruction finishes, the second instruction would have been in the last stage - so it only has to incur the cost of 1 stage AND the cost of the register delay (which is the maximum of all the cost - in this case 65)
- so the answer is $65 + 7 = 72$

- I7.7: Sequential Throughput

I7.7. Sequential Throughput

Assume a processor architecture with 6 stages with execution times of 75 ps, 65 ps, 50 ps, 60 ps, 40 ps, 70 ps and whose register delay time is 8 ps.

What is the throughput, in GIPS, of the **sequential** implementation? Give your answer to 2 decimal places. If there is insufficient information to answer the question, enter 0.

number (2 digits after decimal)

GIPS



Save & Grade Single attempt

Save only

Additional attempts available with new variants

- so you need first to calculate the retirement latency
- trick: if you are given numbers in ps, just do $10^3/\text{retire latency (ps)}$

$$\text{retirement latency} = 8 + 75 + 65 + 50 + 60 + 40 + 70 = 368 \text{ ps}$$

$$\text{throughput} = \frac{10^3}{368 \text{ ps}} = 2.71 \text{ GIPS}$$

- I7.8: Pipelined Throughput

I7.8. Pipelined Throughput

Assume a processor architecture with 6 stages with execution times of 40 ps, 45 ps, 50 ps, 55 ps, 65 ps, 60 ps and whose register delay time is 5 ps.

What is the maximum throughput possible, in GIPS, of a **pipelined** implementation? Give your answer to 2 decimal places. If there is insufficient information to answer the question, enter 0.

number

GIPS



Save & Grade Single attempt

Save only

Additional attempts available with new variants

- same thing as above just that the retirement latency calculation is different

$$\text{retirement latency} = 5 + 65 = 70$$

$$\text{throughput} = \frac{10^3}{70 \text{ ps}} = 14.29 \text{ GIPS}$$

In Class 8

- version 1

I8.1. Mitigating Hazards with NOPs

Assume a pipelined architecture **with no hardware hazard prevention**. Given the following instruction sequence:

```
Ln_0: subq %rsi, %rax
Ln_1: addq %r10, %rdx
Ln_2: rrmovq %r10, %r11
Ln_3: pushq %rcx
Ln_4: divq %rcx, %r12
Ln_5: mrmovq 0x9f2(%rbx), %rdi
Ln_6: xorq %rsi, %r12
```

Assume the following:

1. any registers used were initialized earlier,
2. the stack has also been set up and contains values that may be popped off, and
3. if there are conditional instructions, assume they are taken

Indicate the line number of the first instruction that could produce incorrect behavior due to a data hazard

- line 0
- line 1
- line 2
- line 3
- line 4
- line 5
- line 6
- There is no data hazard

How many nops must you insert to mitigate this hazard? (Enter 0 or leave blank if there was no hazard.)

[Save & Grade](#)[Single attempt](#)[Save only](#)*Additional attempts available with new variants*

- line 6 is trying to read `%r12` that line 4 writes to
 - since there's already 1 instruction between them you need 2 bops
- variation 2

I8.1. Mitigating Hazards with NOPs

Assume a pipelined architecture **with no hardware hazard prevention**. Given the following instruction sequence:

```
Ln_0: popq %rcx
Ln_1: xorq %rdi, %r11
Ln_2: andq %rsp, %rbx
Ln_3: andq %r13, %r12
Ln_4: addq %r10, %rsi
Ln_5: andq %rdx, %rdi
Ln_6: divq %r13, %r8
```

- line 2 tries to read from `%rsp` that line 0 implicitly writes to

- you'll need 2 `nop`s between them

In Class 9

- I9.1: ALU Forwarding

I9.1. ALU Forwarding

Recall that forwarding `e_valE` into the Decode stage allows execution of consecutive ALU instructions where instruction i writes to a register that instruction $i+1$ reads. Perhaps surprisingly, this does not solve the problem in all situations.

Consider the following sequence of instructions:

```
addq %rax, %rbx
nop
nop
nop
addq %rbx, %rcx
```

From where would we need to forward a signal to make this case execute properly without stalling?

- A: you don't have to stall at all - because 3 nops have already occurred. By the time that we get to the second `addq`, the register would have been updated correctly already

- I9.2 ALU Forwarding (Pt 2)

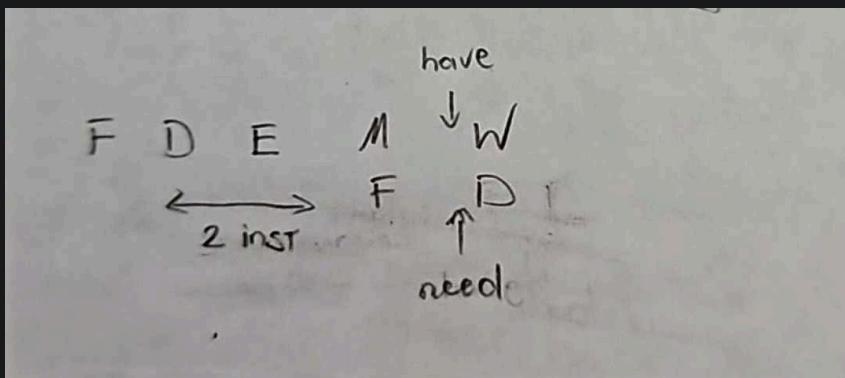
- variation 1

I9.2. ALU Forwarding

Imagine that the only forwarding we do forwards the signal `m_valE` to the Decode register (assuming that we extend the Decode register to contain these values and then add wires to direct them into valA and valB in the Execute register).

What is the minimum number of instructions needed between two ALU operations if the second operation uses the `rB` register from the first one so that the processor does not need to stall.

- note that this question says that we're forwarding to the Decode register (the register BEFORE the decode) stage - NOT the same as what we do in class
- diagram



- so you can see that we can forward it correctly with no stall if there are 2 instruction between them

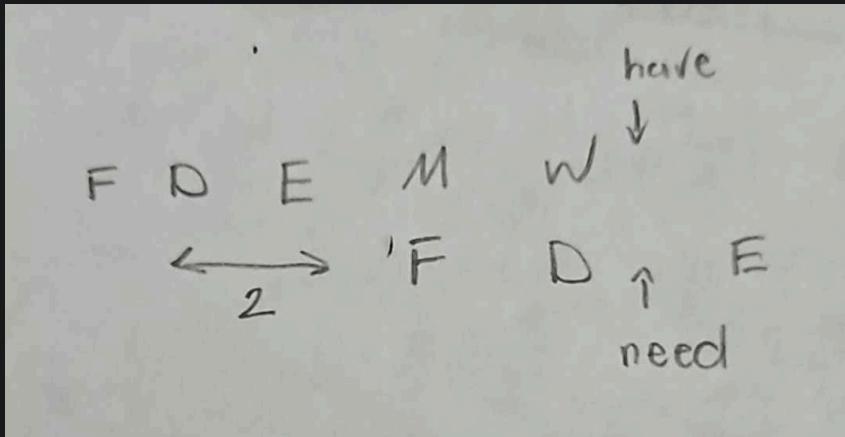
- variation 2

19.2. ALU Forwarding

Imagine that the only forwarding we do forwards the signal `w_valE` to the logic blocks that determine valA and valB.

What is the minimum number of instructions needed between two ALU operations if the second operation uses the `rB` register from the first one so that the processor does not need to stall.

- note: now we are forwarding to between the decode and execute phase (so just before the execute phase) - like what we did in class
- diagram



- so as you can see - need 2 instruction between them

- 9.3: MEM Forwarding

- variation 1

19.3. MEM Forwarding

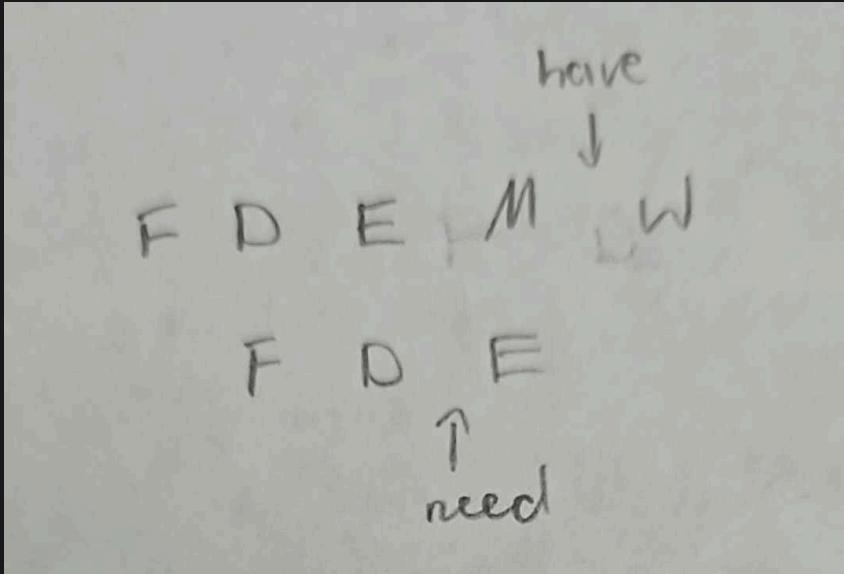
We also use forwarding for values read from memory. Recall that forwarding `e_valE` into the Decode stage allows execution of consecutive ALU instructions where instruction n writes to a register that instruction $n+1$ reads. However, this does not address all data hazard scenarios.

Consider the following sequence of instructions:

```
mrmovq 0(%rax), %rbx  
addq %rbx, %rcx
```

From where would we need to forward a signal to make this case execute properly without stalling?

- here the value we need won't be available until after the memory phase
- however, `addq` comes straight after with no instructions in between - there is no way to do this, you NEED a stall
- diagram



- variation 2

19.3. MEM Forwarding

We also use forwarding for values read from memory. Recall that forwarding `e_valE` into the Decode stage allows execution of consecutive ALU instructions where instruction i writes to a register that instruction $i+1$ reads. However, this does not address all data hazard scenarios.

Consider the following sequence of instructions:

```
mrmovq 0(%rax), %rbx
nop
nop
nop
addq  %rbx, %rcx
```

From where would we need to forward a signal to make this case execute properly without stalling?

- as we know, if you put 3 `nop` between instructions, the value we need will get populated into the registers correctly
- so we need no forwarding

- variation 3

19.3. MEM Forwarding

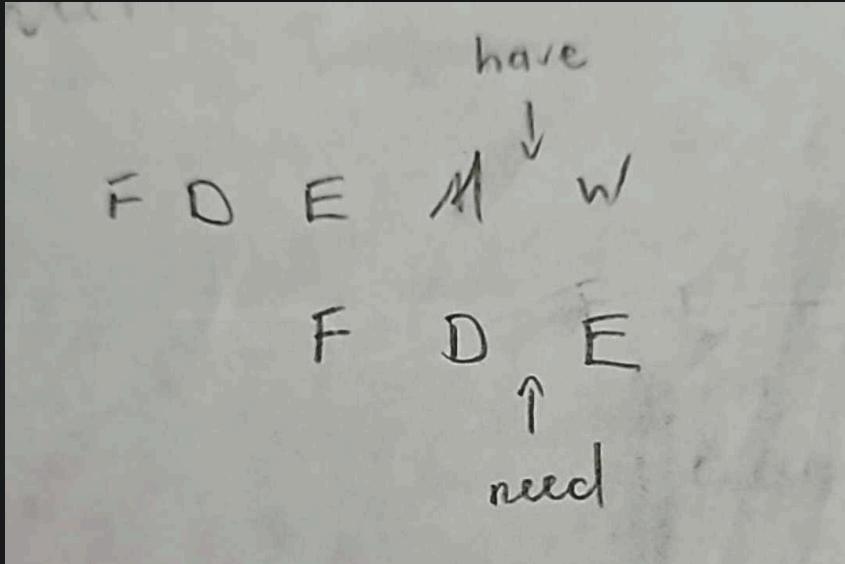
We also use forwarding for values read from memory. Recall that forwarding `e_valE` into the Decode stage allows execution of consecutive ALU instructions where instruction i writes to a register that instruction $i+1$ reads. However, this does not address all data hazard scenarios.

Consider the following sequence of instructions:

```
mrmovq 0(%rax), %rbx
nop
addq  %rbx, %rcx
```

From where would we need to forward a signal to make this case execute properly without stalling?

- diagram



- so we have `valM` (what we need) as soon as we're done with Memory stage and that lines up nicely with where we need it
- note: `m_valM` is the value that is about to go into the W stage - it's also the first time that we actually have `valM`
- so we want to forward `m_valM`
- I9.4: Memory Forwarding (Pt2)
 - very similar to 9.2, the difference is that we don't have the value until the end of the memory stage
 - but the question already dictates which signal we're forwarding so we don't really have to worry about that

In Class 10

- 10.1: Counting Stall/Quash cycles in y86
 - variation 1

Given the implementation with pipelining with forwarding and a branch prediction algorithm that predicts taken for forward branches and not taken for backward branches, for how many cycles in total will the pipeline stall or spend executing instructions that must be squashed/Cancelled?

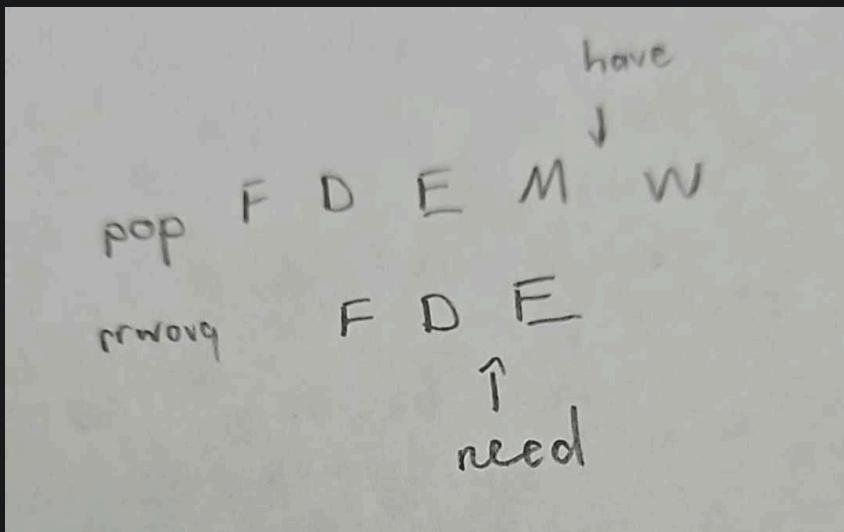
```

Ln_1:  irmovq    $-10, %rbx
Ln_2:  irmovq    $-11, %rdx
Ln_3:  divq     %rbx, %rdx
Ln_4:  jne   L1
Ln_5:  mulq     %rbp, %rsi
Ln_6:  subq     %rbp, %rdi
Ln_7:  addq     %rsi, %rsp
Ln_8:  mulq     %rdi, %rbp
Ln_9:  pushq    %rsp

L1:
Ln_12: xorq    %r9, %r10
Ln_13: popq    %r10
Ln_14: rrmovq   %r10, %rdx
Ln_15: mrmovq   0xef8(%r10), %r11
  
```

- note the prediction logic that's baked in - this changes between question BE CAREFUL
- so we always jump (but in this case we should because `-11 / -10 != 0`)

- so no squashes for guessing wrong
- within the inner function, we incur 1 stall for line 14



- we need the info 1 cycle before we actually have it, so we need to stall
- variation 2

M0.1. Counting Stall/Quash cycles in y86

Given the implementation with pipelining with forwarding and a branch prediction algorithm that predicts never taken, for how many cycles in total will the pipeline stall or spend executing instructions that must be squashed/cancelled?

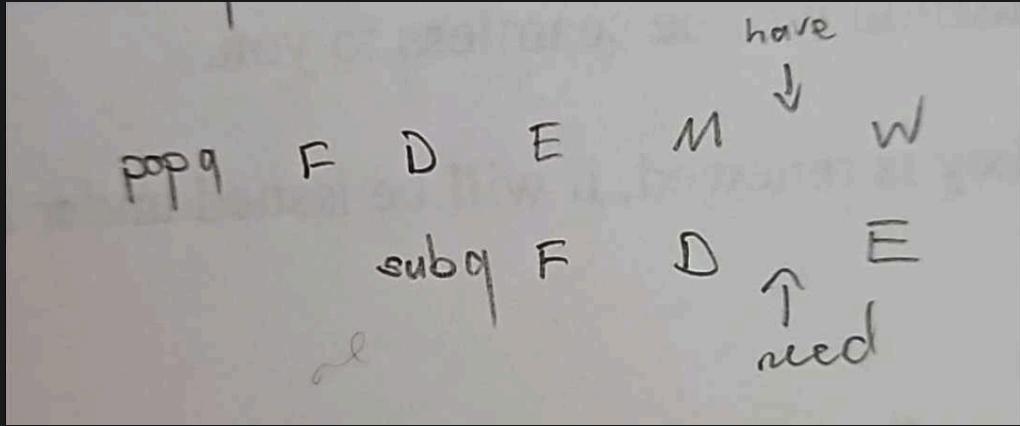
```

Ln_1:  irmovq    $-1, %rbx
Ln_2:  irmovq    $15, %rcx
Ln_3:  addq     %rbx, %rcx
Ln_4:  jle   L1
Ln_5:  irmovq    $0x1189, %rdi
Ln_6:  subq     %rbp, %rsi
Ln_7:  xorq     %rdi, %rbp
Ln_8:  rrmovq    %rsp, %rdi
Ln_9:  mrmovq    0x1345(%rsi), %rbp

L1:
Ln_12: popq    %r11
Ln_13: subq    %r9, %r8
Ln_14: subq    %r10, %r11

```

- this is not exactly clear but this is a "skip forward" type of conditional
 - i.e if it satisfy some condition, do line 5 - 9 in **addition** to line 12 to 14; if not, then just skip and do line 12 to 14 only
- prediction
 - we predict: don't jump
 - actually: don't jump
 - so we guess correctly - don't incur any squash cycles
- line 14 can possibly be problematic but since we have forwarding, there's no stall needed



- we can just forward `m_valM`
- variation 3

Given the implementation with pipelining with forwarding and a branch prediction algorithm that predicts never taken, for how many cycles in total will the pipeline stall or spend executing instructions that must be squashed/cancelled?

```

Ln_1:  irmovq    $19, %rcx
Ln_2:  irmovq    $-16, %rdx
Ln_3:  subq     %rcx, %rdx
Ln_4:  jg       L1
Ln_5:  addq     %rsi, %rsp
Ln_6:  xorq     %rsi, %rbp
Ln_7:  subq     %rsp, %rdi
Ln_8:  xorq     %rsp, %rsi
Ln_9:  andq     %rbp, %rdi

L1:
Ln_12: popq    %r11
Ln_13: popq    %r10
Ln_14: andq    %r10, %rcx
Ln_15: andq    %r10, %r11
Ln_16: mrmovq   0xe39(%r8), %r9
  
```

- prediction:
 - predict: don't jump
 - actual: don't jump
 - so don't incur any squash cycle
- interesting is line 12 to 15
 - line 14 and 15 is problematic - we will need 1 stall
 - line 15 and line 12 can possibly be problematic too - but because we already have that stall from line 14 and 13, this one can execute no problem
- do more practice with this

Practice Quiz

R2.2: Pipeline Instruction Processing

- variation 1

R2.2. Pipeline Instruction Processing

Consider the following sequence of instructions:

```
Ln_1: irmovq    $0xd753, %r11
Ln_2: mulq     %rsi, %rcx
Ln_3: subq     %r11, %rsi
```

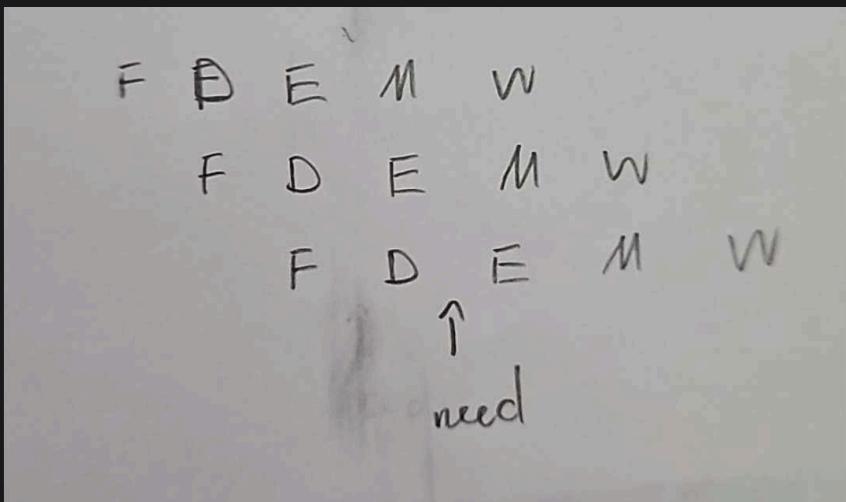
which are executed on our pipelined processor that does not deal with data hazards. When `subq` needs the value of `%r11`, in what stage is the `irmovq` instruction?

Fetch
 Decode
 Execute
 Memory ✓
 Writeback

✓ 100%

[Try a new variant](#)

- diagram



- so when we need the the data - the first instruction is in its memory stage
- note: even though the diagram show that we are "after" the memory stage, in actuality we're still in the memory stage
- variation 2

R2.2. Pipeline Instruction Processing

Consider the following sequence of instructions:

```
Ln_1: irmovq    $0x7e7b, %r10
Ln_2: rrmovq    %r9, %rbx
Ln_3: andq      %rsi, %rcx
Ln_4: divq      %r10, %r11
```

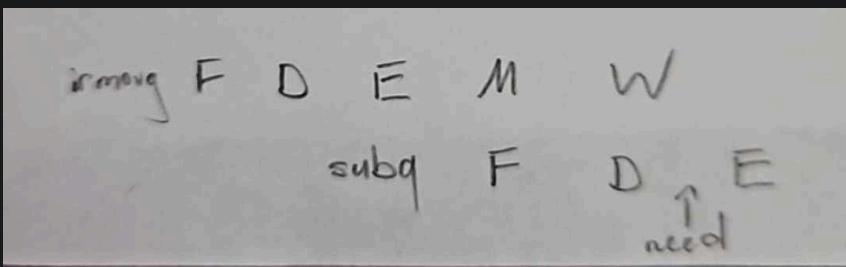
which are executed on our pipelined processor that does not deal with data hazards. When `divq` needs the value of `%r10`, in what stage is the `irmovq` instruction?

- Fetch
- Decode
- Execute
- Memory
- Writeback ✓

✓100%

[Try a new variant](#)

- diagram



- so we're in the writeback stage for `irmovq`

R2.3: Pipeline Stage Destination

- variation 1

R2.3. Pipeline Stage Destinations

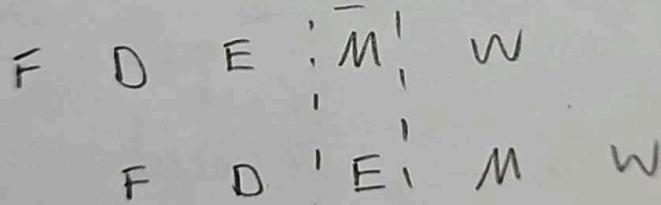
Suppose that the register `%r13` has `0xa000` in it and `%r12` has `0x7000` in it. Consider the following two instructions, which execute one after the other.

1. `mrmovq 0xd00(%r13), %r14`
2. `mrmovq 0x300(%r12), %rsi`

For our pipelined processor with data forwarding, when the first instruction is in the Memory stage, what value will the Execute stage's `dstE` register contain?

If there is insufficient information to answer the question or if the signal's value does not matter, select NONE.

- diagram



- so when first instruction is in the memory stage, the second instruction is in the execute stage → thus the question is asking about the 2nd instruction
- however, the 2nd instruction does not store `valE` into register (it stores `valM` instead) – so the answer is None
- note: if the question asked for `dstM` – the 2nd instruction's `dstM` is `%rsi`

- variation 2

R2.3. Pipeline Stage Destinations

Suppose that the register `%rbp` has `0xd000` in it and `%r12` has `0xf000` in it. Consider the following two instructions, which execute one after the other.

- `mrmovq 0x800(%rbp), %rbx`
- `mrmovq 0x100(%r12), %r9`

For our pipelined processor that uses stalling to deal with data hazards, when the first instruction is in the Memory stage, what value will the Memory stage's `M_dstM` register contain?

If there is insufficient information to answer the question or if the signal's value does not matter, select NONE.

- this is now asking about the memory stage → it's referencing the 1st instruction
- for the first instruction, we're storing `valM` to `%rbx` – thus `dstM` is `rbx`

- variation 3

R2.3. Pipeline Stage Destinations

Suppose that the register `%rbp` has `0xc000` in it and `%r8` has `0x9000` in it. Consider the following two instructions, which execute one after the other.

- `mrmovq 0x200(%rbp), %rdx`
- `mrmovq 0x400(%r8), %rsi`

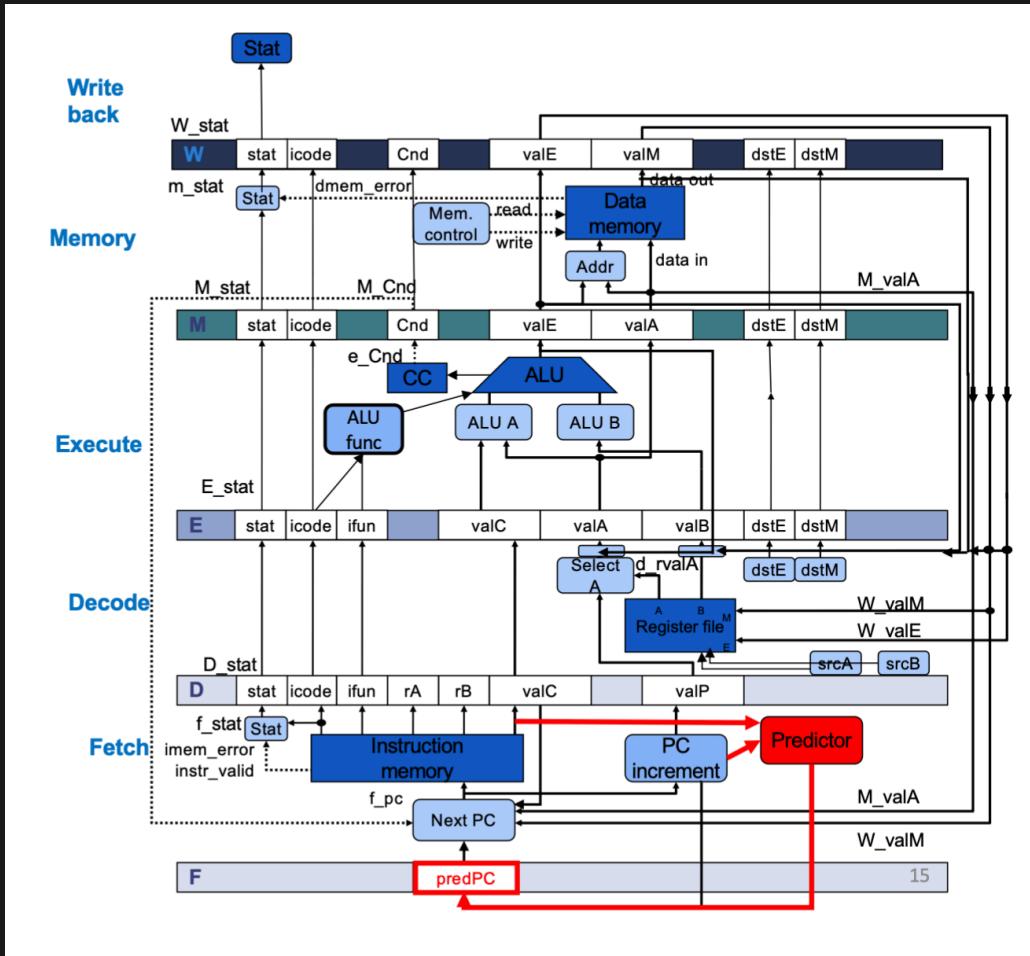
For our pipelined processor that uses stalling to deal with data hazards, when the first instruction is in the Memory stage, what value will the Decode stage's `d_dstM` register contain?

If there is insufficient information to answer the question or if the signal's value does not matter, select NONE.

- use the same diagram as variation 1
- but the decode stage here is referencing some third instruction – which we don't know about
 - (by the time the first instruction reaches memory stage, 2nd instruction reached execute stage which means it has PASSED the decode stage)
- since there is insufficient info – we say NONE

R2.4: What will valA be assigned?

- reference sheet



- variation 1

R2.4. What will valA be assigned?

Suppose the instruction

- popq %rbx

is being processed in our y86 pipelined processor. When this instruction is in the decode stage, which option best identifies the value that will be assigned to d_valA?

- for `pop/ret` - we need the old stack pointer (`%rsp`) AND the new stack pointer (`%rsp + 8`)
 - in the sequential implementation, `R[%rsp] = valB` and `R[%rsp] + 8 = valE` → we read from memory using `valB` and then put `valE` as the new `%rsp`
 - this is no longer the case in pipelined implementation
 - you can see that `valB` doesn't get passed up to the memory stage, but `valA` does
- so now, `valA` and `valB` both initially take on `R[%rsp]`
 - `valB` gets added with 8 to become `valE` (the same)

- `valA` has a wire which bypasses the ALU and pass it up to the Memory stage - we use this to read from memory i.e `valM = M8[valA]`
 - so the answer here is `R[%rsp]`
- variation 2

R2.4. What will valA be assigned?

Suppose the instruction

- `pushq %rbx`

is being processed in our y86 pipelined processor. When this instruction is in the decode stage, which option best identifies the value that will be assigned to `d_valA`?

- for `pushq` we actually only need the new stack pointer (`%rsp - 8`)
 - so `valB` still takes `%rsp`
- but now, we have a value that we're trying to push into the stack (`rA` or `%rbx` in this case) - we make this `valA`
 - this gets bypassed through the ALU (since we use ALU for the `-8`) and gets written to memory
- so the answer is `R[%rbx]`

R2.5: Finding Forwarding Opportunities

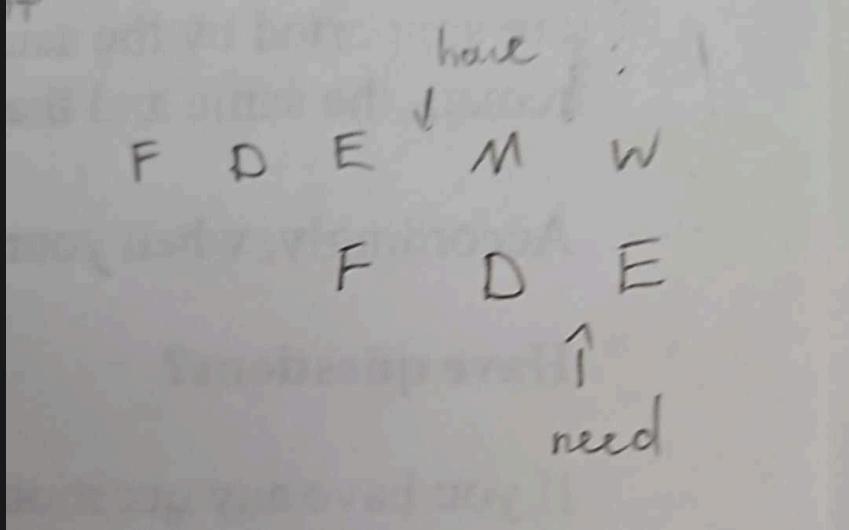
- variation 1

R2.5. Finding Forwarding Opportunities

In class, we determined that forwarding `e_valE` into either or both of `d_valA` and `d_valB` allowed us to execute consecutive ALU instructions where instruction N writes a register that instruction N+1 reads. Perhaps surprisingly, forwarding that signal does not solve the problem that arises from the following sequence:

```
subq    %rdi, %r10
nop
xorq    %r10, %rbx
```

- diagram



- so we have the data (**valE**) way earlier than we need it
- we can just pass it up until the 3rd instruction needs it → we can pass it down from the memory stage
- so the answer is **M_valE**

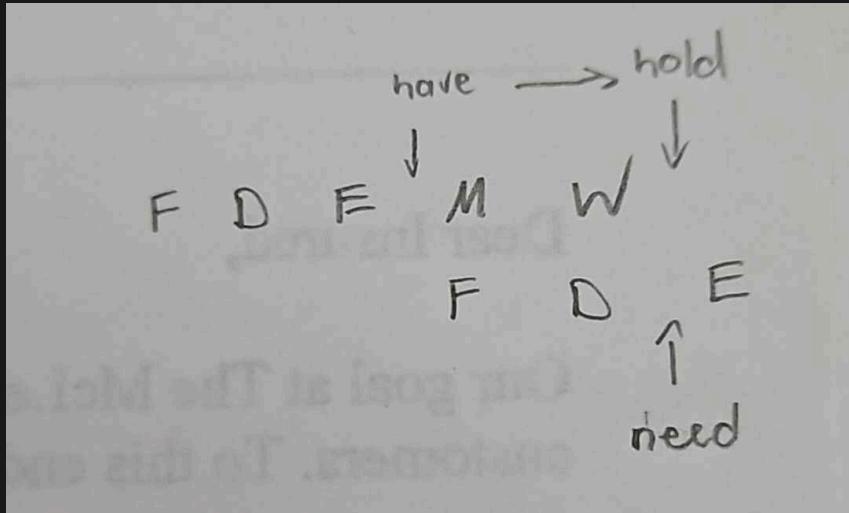
- variation 2

R2.5. Finding Forwarding Opportunities

In class, we determined that forwarding e_valE into either or both of d_valA and d_valB allowed us to execute consecutive ALU instructions where instruction N writes a register that instruction N+1 reads. Perhaps surprisingly, forwarding that signal does not solve the problem that arises from the following sequence:

```
subq    %rax, %rsi
nop
nop
divq    %rsi, %r10
```

- diagram



- again, we can hold it for a bit
- so the answer is **W_valM**

R2.6: Finding Forwarding Opportunities From Memory

- basically, for this question, we don't have the answer until after the memory phase (that is we need to forward `m_valM` at the earliest)
- variation 1

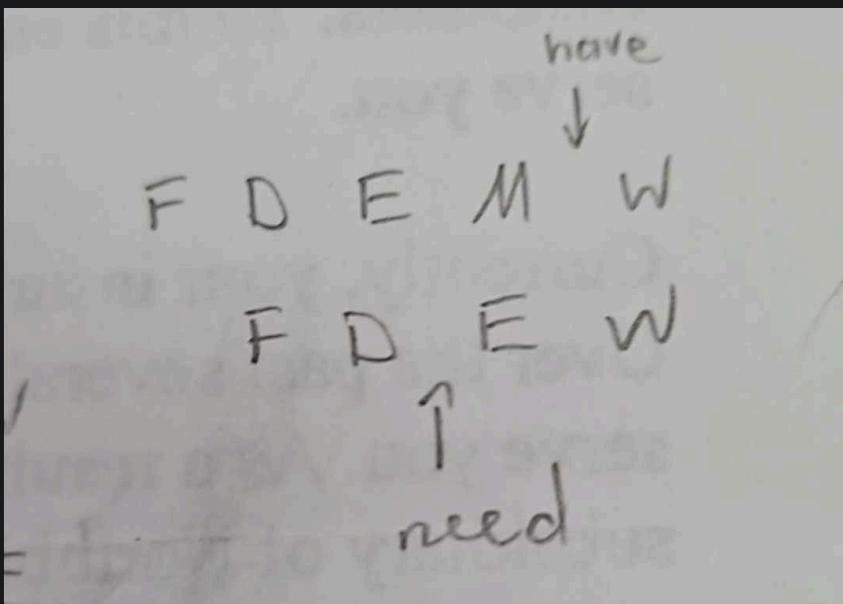
R2.6. Finding Forwarding Opportunities From Memory

Consider instructions that read from memory. In particular, consider the following sequence:

```
mrmovq 0x14c8(%r12), %rbp  
subq   %r9, %rbp
```

From where would we need to forward a signal to make this case execute properly without stalling?

- diagram



- there is no way to make this work without stalling, you simply get the answer too late compared to when you need it

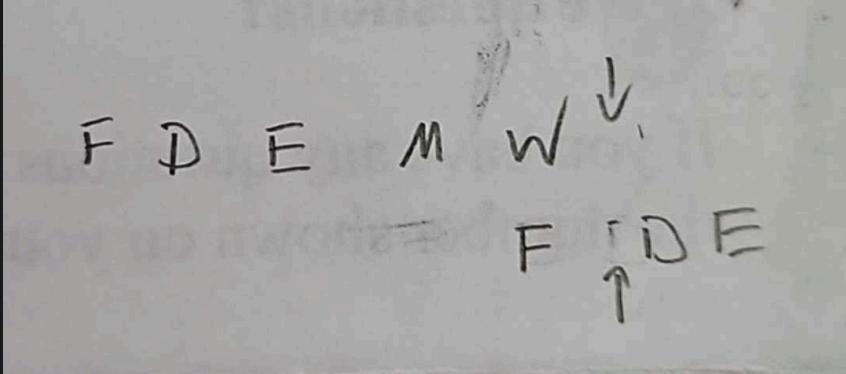
R2.7: Finding Forwarding Opportunities From Memory

- variation 1

R2.7. Finding Forwarding Opportunities From Memory

Imagine that we no longer forward `m_valM` into the forwarding logic in the Decode stage. If we instead forward `W_valM` into the Decode pipeline register, what is the minimum number of instructions needed between an `mrmovq` and an ALU operation that reads from the register written by the `mrmovq` so that the processor does not need to stall?

- note: this one forwards the info into the pipeline registers of the Decode stage (so before the decode stage)
- diagram



- top is where you have to forward, bottom is when you need
 - so you need 3 instructions between them

R2.8: Finding Forwarding Opportunities From Memory

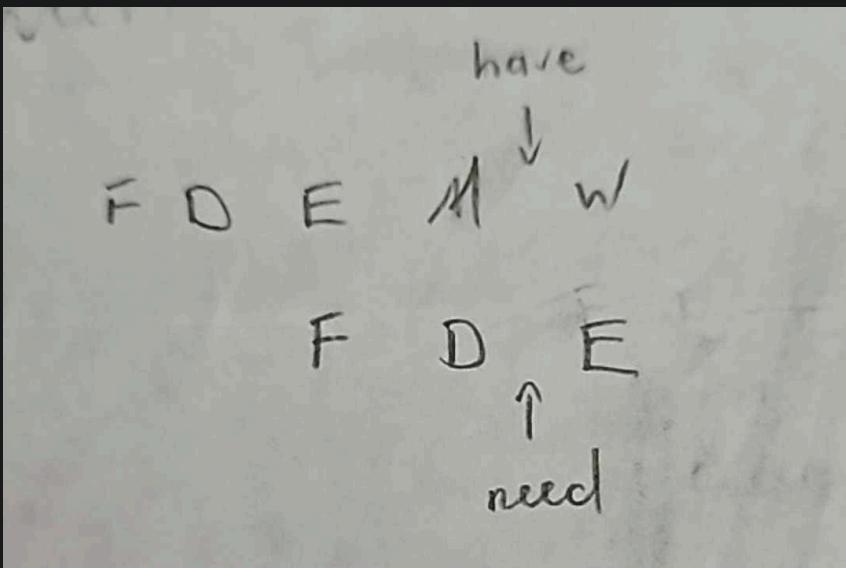
- again, the first value we can send is `m_valM`
- but this time, we are forward to between Decode and Execute
- variation 1

R2.8. Finding Forwarding Opportunities From Memory

What signal should we forward into the forwarding logic in the Decode stage to allow the following sequence of instructions to work correctly without stalling?

```
mrmovq 0x141c(%r10), %rbx
nop
divq    %rbx, %r8
```

- diagram



- so you can just straight up forward `m_valM`

R2.9: Pipeline Instruction Processing - Nops

- question

R2.9. Pipeline Instruction Processing - Nops

Consider the sequence of instructions executed in the **pipelined CPU that does not deal with data hazards (PIPE-)**:

```
Ln_0: mrmovq 0x8dae(%r11), %rsi
Ln_1: mrmovq 0xbbb74(%rsi), %rdi
Ln_2: divq   %rcx, %r12
Ln_3: mulq   %rcx, %rbx
Ln_4: mrmovq 0x9bba(%r12), %r9
Ln_5: mulq   %rbx, %rax
Ln_6: rmmovq %rcx, 0x56b8(%rcx)
```

Consider the following questions cumulatively. That is, if in Q1, you indicate that you need to insert 5 nops, when you answer Q2, answer assuming that those 5 nops have been inserted immediately before the line that triggers the hazard in Q1.

- part 1

Observe that in line 1, `%rsi` depends on the value of `%rsi` from line 0.

Q1: How many nops need to be inserted immediately before line 1 because of this potential hazard?

3



✓ 100%

- they came immediately after one another so you need the full 3 nops

- part 2

Observe that in line 4, `%r12` depends on the value of `%r12` from line 2.

Q2: How many nops need to be inserted immediately before line 4 because of this potential hazard?

2



✓ 100%

- this one is unrelated to the last one
- since there is already 1 instruction between them - you need 2 nops

- part 3

Observe that in line 5, `%rbx` depends on the value of `%rbx` from line 3.

Q3: How many nops need to be inserted immediately before line 5 because of this potential hazard?

0



✓ 100%

- this one is kinda related to part 2 since it occurs right after line 2
- since part 2 we already incurred 2 nops, the registers will be populated correctly by the time we get to line 5
- so no nops needed

Lab 4

L4.1: Convert from units of time to seconds

- question

L4.1. Convert from units of time to seconds

How many milliseconds are there in a second?

integer ?

Save & Grade Single attempt Save only Additional attempts available with new variants ?

- dumb question
- there are 10^3 ms in a second

L4.2: Compute latencies of pipelined and sequential processors

- question

L4.2. Compute latencies of pipelined and sequential processors

Consider a processor with a register latency of 18 ps and 4 stages, each of which takes 130 ps

What is the end to end latency of an instruction in the sequential implementation? integer ps ?

What is the end to end latency of an instruction in the pipelined implementation? integer ps ?

Save & Grade Single attempt Save only Additional attempts available with new variants ?

- for the sequential latency, you only incur the register latency once (at the start)
- for the pipelined latency, you incur the register latency at every stage

$$\text{seq latency} = 18 + 4(130) = 538 \text{ ps}$$

$$\text{pipe latency} = (18 + 140) \times 4 = 592 \text{ ps}$$

L4.3: Compute number of ps to retire an instruction

- this question is about **retirement latency**
 - for sequential it's the same as latency
 - for pipelined it's the `time for the longest stage + register latency`
- variation 1

L4.3. Compute number of ps to retire an instruction

Consider a processor with a register latency of 21 ps and 4 stages, each of which takes 180 ps

In the sequential implementation, an instruction will be retired every

integer ps ?

Save & Grade Single attempt

Save only

Additional attempts available with new variants ?

- same as latency

$$\text{retirement latency} = 21 + (4 \times 180) = 741 \text{ ps}$$

- variation 2

L4.3. Compute number of ps to retire an instruction

Consider a processor with a register latency of 20 ps and 5 stages, each of which takes 160 ps

In the pipelined implementation, an instruction will be retired every

integer ps ?

Save & Grade Single attempt

Save only

Additional attempts available with new variants ?

- it's just the time for the longest stage (including register latency) basically

$$\text{retirement latency} = (160 + 20) = 180 \text{ ps}$$

L4.4: Compute throughput of pipelined and sequential processors

- compute the retirement latency and use that to calculate throughput
 - if the number is given in picosecond, you can just do $10^3 / \text{retirement latency}$
- question

L4.4. Compute throughput of pipelined and sequential processors

Consider a processor with a register latency of 9 ps and 4 stages, each of which takes 100 ps

What is the throughput of the sequential implementation? (Round your answer to two decimal places.)

number (2 dig) GIPS ?

What is the throughput of the pipelined implementation? (Round your answer to two decimal places.)

number (2 dig) GIPS ?

Save & Grade Single attempt

Save only

Additional attempts available with new variants ?

- sequential

$$\text{retirement latency} = 9 + (4 \times 100) = 409 \text{ ps}$$

$$\text{throughput} = \frac{10^3}{409} = 2.44 \text{ GIPS}$$

- pipelined

$$\text{retirement latency} = 100 + 9 = 109 \text{ ps}$$

$$\text{throughput} = \frac{10^3}{109} = 9.17 \text{ GIPS}$$

L4.5: Analyzing Pipeline Behaviour

- overall question

L4.5. Analysing Pipeline Behavior

We are going to examine hazard behavior in the different y86 pipelined implementations.

There may be ways to brute force the answers to this problem, but you will want to make sure that you understand exactly what is happening. Spending time understanding what the simulator is doing should help you significantly on the next quiz and the next lab.

Consider the following y86 program

```
1  main:
2  1:  irmovq  stack, %rsp
3  2:  irmovq  b, %rdi
4  3:  irmovq  f, %rsi
5  4:  call swap
6  5:  halt
7
8  swap:
9  6:  mrmovq  0(%rdi), %r10
10 7:  mrmovq  0(%rsi), %r11
11 8:  rmmovq  %r10, 0(%rsi)
12 9:  rmmovq  %r11, 0(%rdi)
13 10: ret
14
15 .pos 0x1000
16 data:
17 a:    .quad 0xCAFE
18 b:    .quad 0xFACE
19 c:    .quad 0xFEED
20 d:    .quad 0xDECAF
21 e:    .quad 0xBAD
22 f:    .quad 0xBEAD
23 g:    .quad 0xACE
24 h:    .quad 0xFADE
25
26 .pos 0x2000
27     .quad 0
```

```
28      .quad 0
29      .quad 0
30      .quad 0
31      .quad 0
32      .quad 0
33  stack:
```

- part 1

Run the program as-is in the y86 simulator using the Sequential implementation

Enter the number of clock cycles and retired instructions below.

10	clock ticks	?	✓ 100%	10	instructions	?	✓ 100%
----	-------------	---	--------	----	--------------	---	--------

- you just run it in the simulator
- point is: clock cycles are different in sequential vs pipelined
 - 1 clock cycle in a sequential means you retire 1 instruction (all 5 phases)
 - in pipelined: you execute a phase → you can think of it as executing 1/5 of an instruction 5 times

- part 2

If you try to run this program using the naive pipelined implementation, things will go very badly! Even if you add nops, you will find that there is one instruction that will not work. Which one is it?

- ret ✓
- mrmovq
- irmovq
- call
- rmmovq

- `ret` will require stalling → but it has to be machine implemented stalling (i.e we can't add nops manually to make it work)
- this is because within the `ret` itself, it has to stall for 3 cycles so that it can get the address out of memory and feed it back into the PC

- part 3 and 4

Now run the implementation with pipelining and stalling

Enter the number of clock cycles and retired instructions below.

20	clock ticks	?	✓ 100%	10	instructions	?	✓ 100%
----	-------------	---	--------	----	--------------	---	--------

Now run the implementation with pipelining and data forwarding

Enter the number of clock cycles and retired instructions below.

17	clock ticks	?	✓ 100%	10	instructions	?	✓ 100%
----	-------------	---	--------	----	--------------	---	--------

- you can just run it

- part 5

Data forwarding allowed us to avoid some of the stalls and bubbles that happen in the stalling implementation

Which lines of code stall in the stalling implementation, but do not in the data forwarding one? (Note: if instruction `i` stalls because instruction `i-1` is stalled, do not count it as causing a stall unless it stalls after instruction `i-1` is no longer stalled.)

- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9
- 10

Select all possible options that apply.

100%

- first - which lines do we have to stall on if we just did stalling
 - on line 4: we need to stall for 1 cycle we mess with the stack pointer on line 1 and we need that for `call`
 - due to this stall, it means that line 6 no longer has to stall
 - also means that line 7 no longer has to stall (they both have 3 instructions between them now)
 - on line 8: we need to stall for 2 cycles because 8 uses `%r10` which is modified in line 6
 - due to these stalls, we don't have to stall for line 9 anymore which uses `%r11` which line 7 modifies
- after data forwarding
 - line 4: we can forward new `%rsp` value to `call`
 - not really what the question is asking but we also don't have to stall on line 6 and 7 any more BECAUSE OF DATA FORWARDING - not because of stalling (because we don't stall anymore)
 - line 8: we can also forward the dat from line 6

L4.6: Pipelining and Inlining

- question

L4.6. Pipelining and Inlining

We are going to examine the impact that function calls have on pipelining by examining what happens when you move the body of a function into the place from which it was called. (This is called *inlining*.)

There may be ways to brute force the answers on this, but you will want to make sure that you understand exactly what is happening. Spending time understanding what the simulator is doing should help you significantly on the next quiz and the next lab.

Consider the following y86 program

```
1 main:
```

```

2 irmovq stack, %rsp
3 irmovq b, %rdi
4 irmovq f, %rsi
5 call swap
6 halt
7
8 swap:
9 mrmovq 0(%rdi), %r10
10 mrmovq 0(%rsi), %r11
11 rmmovq %r10, 0(%rsi)
12 rmmovq %r11, 0(%rdi)
13 ret
14
15 .pos 0x1000
16 data:
17 a:     .quad 0xCAFE
18 b:     .quad 0xFACE
19 c:     .quad 0xFEED
20 d:     .quad 0xDECAF
21 e:     .quad 0xBAD
22 f:     .quad 0xBEAD
23 g:     .quad 0xACE
24 h:     .quad 0xFADE
25
26 .pos 0x2000
27     .quad 0
28     .quad 0
29     .quad 0
30     .quad 0
31     .quad 0
32     .quad 0
33 stack:

```

Inlining is a technique frequently used to improve code performance. In this problem, we'll explore the value of inlining.

Rewrite the program above to incorporate the swap function in the main program -- do not optimize the program beyond replacing the `call` instruction with the code that implements the `swap` functionality. (Hint: the new program should have 8 instructions.)

- part 1

Run this program in the PIPE- (naive) implementation

What is the final value left in %r10 (in hexadecimal)?

0x 2030F430

?

✓ 100%

- you can run this in the simulator
- but the explanation for it is that at the beginning `%rdi = 0`

- by the time we get to `mrmovq 0(%rdi), %r10` - we're actually reading from `M[0]` instead of `M[b]` if it was loaded correctly
- what's in memory at address 0 happens to be the instruction code for the first `irmovq` instruction

`000000 30 F4 30 20 00 00 00 00 00`

hence we load in `0x2030f430` into `r10`

- part 2

Now that you've inlined swap, which line of code could you safely remove from your program?

- `irmovq stack, %rsp` ✓
- `irmovq b, %rdi`
- `irmovq f, %rsi`
- `mrmovq 0(%rdi), %r10`
- `mrmovq 0(%rsi), %r11`
- `rmmovq %r10, 0(%rsi)`
- `rmmovq %r11, 0(%rdi)`

- inlined code

```

1 main:
2 irmovq stack, %rsp
3 irmovq b, %rdi
4 irmovq f, %rsi
5
6 mrmovq 0(%rdi), %r10
7 mrmovq 0(%rsi), %r11
8 rmmovq %r10, 0(%rsi)
9 rmmovq %r11, 0(%rdi)
10
11 halt

```

- we don't need to `call` anymore - so we don't need to the stack anymore

- part 3 and 4

Remove the instruction referred to in the previous question

Now, insert the minimal number of nop instructions to get this program to run correctly in the PIPE-implementation (use the `Start` button to run your program to completion, draining the pipeline at the end).

How many nops did you have to insert?

4



✓ 100%

- need 2 nops before `mrmovq 0(%rdi), %r10` so `%rdi` gets loaded correctly
- need 2 nops before `rmmovq %r10, 0(%rsi)` so `%r10` gets loaded correctly
- code

```

1 main:

```

```

2 irmovq stack, %rsp
3 irmovq b, %rdi
4 irmovq f, %rsi
5 nop
6 nop
7 mrmovq 0(%rdi), %r10
8 mrmovq 0(%rsi), %r11
9 nop
10 nop
11 rmmovq %r10, 0(%rsi)
12 rmmovq %r11, 0(%rdi)
13
14 halt

```

Between which pairs of instructions must you insert nops?

- `irmovq b, %rdi` and `irmovq f, %rsi`
- `irmovq f, %rsi` and `mrmovq 0(%rdi), %r10` ✓
- `mrmovq 0(%rdi), %r10` and `mrmovq 0(%rsi), %r11`
- `mrmovq 0(%rsi), %r11` and `rmmovq %r10, 0(%rsi)` ✓
- `rmmovq %r10, 0(%rsi)` and `rmmovq %r11, 0(%rdi)`
- `rmmovq %r11, 0(%rdi)` and `halt`

- part 5

Enter the number of clock cycles it takes to run your implementation to completion on the PIPE- (naive).

15	cycles	?	✓ 100%
----	--------	---	--------

If you really got the minimal number of nops, then you should be able to remove the nops and run the PIPE with stalling implementation.

Enter the number of clock cycles.

15	clock ticks	?	✓ 100%
----	-------------	---	--------

Now run the implementation with pipelining and data forwarding

Enter the number of clock cycles.

11	clock ticks	?	✓ 100%
----	-------------	---	--------

- for the first 2 - there are 11 instructions overall, so we need to do it in 11 clock cycles, and eat 4 additional cycles at the end

L4.7: Sequential v Pipelined Performance

- overall question

L4.7. Sequential v Pipelined Performance

Please make the following assumptions:

1. Each statement below is comparing two implementations of the same architecture.
2. You should ignore any cycles that happen after the address of a halt instruction appears in the program counter.

For each statement below, select the item from the dropdown that makes the statement true. If none of the selections make the statement true, select "None".

- clock cycles

- variation 1

A specific program is likely to **✓100%** take fewer clock cycles on a sequential implementation than on a pipelined implementation.

- because we count clock cycles kinda differently between pipeline and sequential implementations
- in the best case (and excluding the additional cycles at the end) - sequential and pipeline will have equal number of clock cycles (1 per instruction)
- however, if we have any kind of data or control hazards - we might need to stall so the pipelined instruction is likely to take more cycles

- variation 2

On a pipelined implementation, a given program is likely to **✓100%** require more clock cycles than the number of instructions executed.

- again, if there are hazards, we will require more cycles than just the instructions executed (due to stalls)

- variation 3

A specific program is unlikely to take more clock cycles on a pipelined implementation with 5 stages than on one with 6 stages.

- if you have more stages - you are more likely to get a data hazard - thus taking more cycles
- note: data hazards only affect cycle count - not cycle time

- cycle time

- variation 1

A pipelined implementation with 5 stages is likely to **✓100%** have a longer cycle time than a pipelined implementation with 6.

- general idea: longer pipeline LIKELY to have shorter cycle time
- this is because if you're doing the same work and you have more stages, it's likely that you're breaking the work up more evenly, so your maximum register delay could potentially be smaller - thus decreasing the overall cycle time

- other

- variation 1

A specific program is likely to **✓ 100%** complete more quickly on a pipelined implementation than on a sequential implementation.

- there's no reason to invent pipelined implementation if it's not going to be faster
- but there are some cases - maybe when there's 1 instruction or something that they take the same

- variation 2

On a sequential implementation, a given program will never **✓ 100%** require fewer clock cycles than the number of instructions executed.

- this will never happen - you literally need the same number of clock cycles as instructions in the sequential implementation

L4.8: Pipelined Behavior

- overall question

L4.8. Pipelined Behavior

All the following statements describe behavior of one or more pipelined implementations of the y86 architecture. For the purposes of this question, **ignore all behavior after the address of a halt instruction appears in the program counter**.

For each statement below, select the item from the dropdown that makes the statement true. If none of the selections make the statement true, select "None".

- we are supposed to ignore the the 4 halts at the end

- variation 1

In the 'PIPE with Data Forwarding and Branch Prediction' implementation, all programs will execute for **✓ 100%** clock cycles than there are instructions retired.

- data forwarding and branch prediction DOES NOT solve all stalls (i.e memory reads, ret and wrong jumps) - so there may be cases where we stall and thus programs will execute for more clock cycles than there are instructions
- since we don't actually know if there will be hazards or not - we say NONE

- variation 2

For any program in which the 'PIPE-' implementation returns incorrect results, to produce correct results the 'PIPE-' implementation requires inserting **✗ 0%** nops than the 'PIPE with Data Forwarding' implementation requires stalls.

- the answer is NONE
- PIPE- might return incorrect due to either control or data hazards
 - some data hazards can be solved with data forwarding (but not all)
 - control hazards may or may not be solved by branch prediction

- without knowing precisely what the program is doing, you can't know whether data forwarding alone will be helpful
- variation 3

For any program in which the 'PIPE-' implementation returns incorrect results, the number of instructions retired by the 'PIPE with Data Forwarding' implementation will be ✓ 100% than the number of clock cycles that it takes to execute.

- we know that there are some kind of hazards because PIPE returns incorrect result
- data forwarding may or may not solve this issue
 - if it does solve it, then cycle count = instruction retired
 - if it doesn't, then cycle count > instruction retired
- so since we don't know what kind of hazards, we say none

Corrections

Question 1: Moving Pointers in C

- question

By now you probably recall that arithmetic with pointers is different from arithmetic with integers in C.

Given the following code:

```
struct astruct {
    int8_t f1;
    uint16_t *f2;
};
struct astruct as[] = { /* some huge number of structs */ };
```

Which of the following code snippets will initialize a pointer `ptr` to be the address of `as[2]` in memory?

- `uint8_t *ptr = (uint8_t *)as;`
`ptr += 2;`
- `struct astruct *ptr = as;`
`ptr += 2;`
- `uint8_t *ptr = (uint8_t *)as;`
`ptr += 2 * sizeof(struct astruct);`
- `struct astruct *ptr = as;`
`ptr += 2 * sizeof(struct astruct);`

Select all possible options that apply.

✓ 100%

- `astuct` is 16 bytes in size
 - 1 byte for `f1`, 7 byte for padding, 8 byte for `f2` (because it's a pointer)
 - so `as[2]` would be 32 bytes from the original address
- first option

- `ptr` is now `uint8_t` pointer
- adding 2 to `ptr` will only move it 2 bytes
- second option
 - the `ptr` is a `astruct` pointer
 - adding to `ptr` will move it 32 bytes as needed
- third option
 - the `ptr` is a `uint8_t` pointer
 - so simply adding 2 to the `ptr` will only move it 2 bytes
 - but we actually add `2 * sizeof(struct astruct) = 32` to it
 - so we move down 32 bytes as needed
- last:
 - the `ptr` is a `astruct` pointer
 - we actually add `2 * sizeof(struct astruct) = 32` to it
 - so we move down `32 * 32 = 1024` bytes
- important thing to note about this question is pointer arithmetic
 - incrementing pointers will move it by the size of the type it is pointing to (i.e the pointer type)
 - ex. `uint16* ptr` - incrementing this will move `ptr` by 2 bytes
 - ex. `uint64* ptr` - incrementing this will move `ptr` by 8 bytes

Question 2: Assembling Instructions into Memory

- question

Consider the following sequence of instructions.

```
ret  
jle 0xc924  
nop
```

Assume that the first instruction of this program begins at address 0x1e37. Fill in the table below, with the encoding, in hexadecimal, of these instructions. You must fill in all the entries in the address column, but may leave any other boxes whose contents you do not know empty.

Address									
0x1e28									
0x									90
0x1E									✓ 100%
0x	71	24	c9	00	00	00	00	00	
0x1E	✓ 100%								
0x	00	10							
0x1E	✓ 100%	✓ 100%							
0x									
0x1E									

- just need to make sure you count the address properly
- fill out all the address column first
- then count to where you should start
 - note that first box on every row is address the row start address, and count from there

Question 3: Decoding instructions from memory

- question

Question 3: Decoding instructions from memory

For all parts of this question:

- Credit for the registers and offset will be given only if the instruction is correct.
- Select None for any unused register values.

A portion of the memory of a y86 processor is shown below (as a hex dump).

Address	Content
0xe0	50 ac 88 f4 00 00 00 00 00 00 00 40 10 23 ad 00 00
0xf0	00 00 00 00 74 5a 55 00 00 00 00 00 00 73 9d
0x100	25 00 00 00 00 00 00 50 1c 03 ff 00 00 00 00 00
0x110	00 80 96 af 00 00 00 00 00 00 a0 0f 90 30 00 00
0x120	00 00 00 00 00 00 00 30 30 00 00 00 90 90 30

- part 1

Select all the necessary parts for the y86 instruction that begins at address 0xe0.

Instruction	rA	rB	Offset/Address/Value
mrmovq	%r10	%r12	0x f488

- at `0xe0` the first byte is `50` which is an `mrmovq`
- the next byte specify the registers, `a = %r10` and `c = %r12`
- the next 8 bytes if to specify the offset - remember that this is little endian
 - so it's `0xf488`

- part 2 and 3

Select all the necessary parts for the y86 instruction that begins at address 0xf4.

Instruction	rA	rB	Offset/Address/Value
jne	None	None	0x 555a

Select all the necessary parts for the y86 instruction that begins at address 0x11a.

Instruction	rA	rB	Offset/Address/Value
pushq	%rax	None	0x Leave blank if no

- same kind of thing

Question 4:

- question

Question 4: What happens in the writeback stage?

Suppose the instruction `popq %rax` is being processed by the pipelined implementation of our y86 processor. In the writeback stage, which of the following actions take place?

Select all that apply.

- R[%rsp] = W_valE ✓
- R[%rsp] = W_valM
- R[rA] = W_valE
- R[rA] = W_valM ✓
- R[rB] = W_valE
- R[rB] = W_valM
- None of the above

Select all possible options that apply. ?

✓ 100%

This question is complete and cannot be answered again.

- at the writeback stage, we are writing into registers
- we want to update the stack pointer (we decremented it so its `valE`) so need to write to `%rsp`
- we want to put the popped value from memory (which is `valM`) into its destination `rA`

Question 5: Pipeline Stage Signals

- question

Question 5: Pipeline Stage Signals

Suppose that the register `%rbx` has `0xa000` in it and `%r11` has `0x6000` in it. Consider the following two instructions, which execute one after the other.

1. `mrmovq 0x900(%rbx), %r9`
2. `mrmovq 0x500(%r11), %r9`

For our pipelined processor that uses stalling to deal with data hazards, when the first instruction is in the Execute stage, what value will the Decode stage's `D_valC` register have? Give the value in hex and do not include the `0x`.

If there is insufficient information to answer the question, enter 0.

0x 0x500

? ✓100%

- when the first instruction is in the execute stage, the second instruction is in the decode stage → so this question is referring to the second instruction
- for the second instruction, `valC` is the offset given which is `0x500`

Question 6: What will e_valE be assigned?

- question

Question 6: What will e_valE be assigned?

Suppose the instruction

- call functionMisery

is being processed in our y86 pipelined processor. When this instruction is in the execute stage select the option that best describes what e_valE will be assigned?

- R[%rbx]
- E_valB + E_valC
- E_valB + 8
- E_valB - 8 ✓
- None of the above

✓100%

This question is complete and cannot be answered again.

- when you `call`, you are pushing the return address to the stack, so you are modifying the stack pointer `%rsp`
- also in `call`, `%rsp` is fed in as both `valA` and `valB`
- so as you are pushing something to the stack, you decrement it, hence `valB - 8`

Question 7: Counting Off Stages

- part 1

Question 7: Counting Off Stages

Imagine we are using our **naive pipelined processor that does not automatically stall instructions**. We run the processor for **exactly 8 cycles**. Indicate what stage each instruction is in. For instructions that have been retired, choose "retired"; for instructions that have not yet entered `Fetch`, choose "unstarted".

Note: If we ran the processor for exactly 1 cycle instead, the answer for `Ln_1` would be `Fetch` and for all other lines it would be `unstarted`.

- retired ↴ ✓100% `Ln_1: rrmovq %rbp, %r11`
- retired ↴ ✓100% `Ln_2: irmovq $0x340, %rdx`
- retired ↴ ✓100% `Ln_3: irmovq $0xab8e, %r10`
- Writeback ↴ ✓100% `Ln_4: subq %r10, %r12`
- Memory ↴ ✓100% `Ln_5: rrmovq %rsi, %rax`
- Execute ↴ ✓100% `Ln_6: divq %r12, %rbp`
- Decode ↴ ✓100% `Ln_7: rrmovq %rcx, %rdi`

- you can run this in the simulator

- but you also know that in the naive implementation, if we're starting fresh, the first instruction gets retired after 5 cycles (so on the 6)
 - so the 2nd instruction gets retired on 7 cycle
 - third instruction gets retired after 8 cycle
- obviously the fourth one would be on its final stage (Writeback)
- and since we don't have any stalling, the next instructions would just be one stage behind the instruction before

- part 2

Now, imagine we use our **pipelined processor with stalling but no data forwarding**. Again, we run the processor for **exactly 8 cycles**. Indicate what stage each instruction is in.

retired	✓ 100%	Ln_1: rrmovq %rbp, %r11
retired	✓ 100%	Ln_2: irmovq \$0x340, %rdx
retired	✓ 100%	Ln_3: irmovq \$0xab8e, %r10
Decode	✓ 100%	Ln_4: subq %r10, %r12
Fetch	✓ 100%	Ln_5: rrmovq %rsi, %rax
unstarted	✓ 100%	Ln_6: divq %r12, %rbp
unstarted	✓ 100%	Ln_7: rrmovq %rcx, %rdi

- identify data hazard: line 4 and line 3 because line 4 reads `%r10` and line 3 writes to it
- same idea for the first 3 as above since there's no hazard, there're all retired
- since we're stalling, that means that line 4 is currently in its decode stage and waiting for line 3 to get retired fully
- line 5 would be in fetch, and the following would be unstarted

- random attempt



- TODO: picture from phone

Question 8: From Forwarding Signals to Code

- question

Question 8: From Forwarding Signals to Code

This problem is about the **pipelined processor with data forwarding**. In each part below, you build a sequence of instructions such that:

- **The first instruction needs to forward** from a particular source to a particular destination in order **to avoid a data hazard** in a later instruction. Each part states its source and destination.
- The later instruction **uses the forwarded value at just the right moment**.
- **All other instructions are nops**.
- You may not use a `popq %rsp` instruction.
- Some instructions/registers may be unavailable in the dropdown boxes for some parts.

For condition codes, assume **SF starts true (1)** and **ZF starts false (0)**.

Some credit is available for solutions that cause any forwarding, with more credit as the forwarding shares certain properties with the forwarding we request. So, **carefully** ensure that your solution causes forwarding!

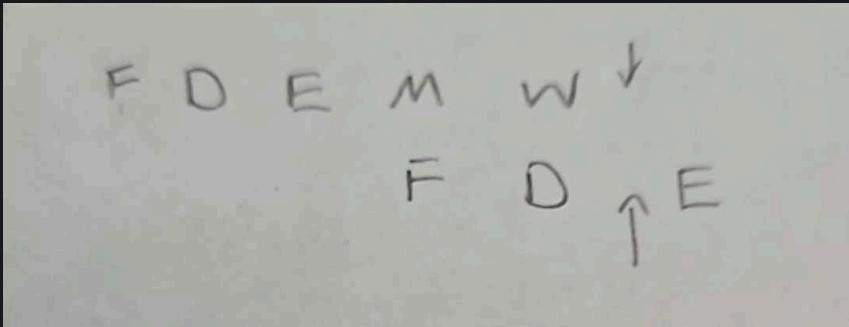
The pipelined processor diagram is available below for reference. Also review the y86 reference!

- so this is basically the opposite of all the practice we've done so far, we want to write some instruction set that causes the data forwarding
- part 1

For this problem, the first instruction needs to forward from `W_valM` to `d_valA`.

Instruction	rA	rB	Offset/Address/Value
mrmovq	x 0%	%rsp	x 0% %rax x 0% 0x 0 ? x 0%
nop	x 0%	None	x 0% None x 0% 0x Leave blank if no offset ? x 0%
nop	x 0%	None	x 0% None x 0% 0x Leave blank if no offset ? x 0%
andq	x 0%	%rax	x 0% %rsi x 0% 0x Leave blank if no offset ? x 0%
nop	x 0%	None	x 0% None x 0% 0x Leave blank if no offset ? x 0%

- for that signal forwarding to happen, so it would look something like



- first hint is that we need `valM` → we must use `mrmovq` to cause some kind of data hazard
- from the diagram, we can see that for it to forward in the way we want it to, the second instruction must be 2 stages behind → add 2 `nops` in between
- the answer above actually works, I just messed up the order of `rA` and `rB` in the `mrmovq` (`valM` is written to `rA`, not `rB`)

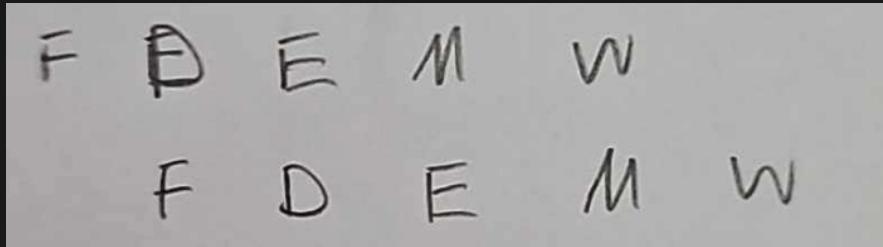
```
1 mrmovq 0(%rsp), %rax      # rA = %rax, rB = %rsp
2 nop
3 nop
4 andq %rax, %rax          # cause data hazard and correct forwarding
```

- part 2

For this problem, the first instruction needs to forward from `e_valA` to `d_valA`.

Instruction	rA	rB	Offset/Address/Value
popq	✓ 100%	%rax	✓ 100% None ✓ 100%
0x Leave blank if no offset	?	✓ 100%	
popq	✓ 100%	%rcx	✓ 100% None ✓ 100%
0x Leave blank if no offset	?	✓ 100%	
nop	✓ 100%	None	✓ 100% None ✓ 100%
0x Leave blank if no offset	?	✓ 100%	
nop	✓ 100%	None	✓ 100% None ✓ 100%
0x Leave blank if no offset	?	✓ 100%	
nop	✓ 100%	None	✓ 100% None ✓ 100%
0x Leave blank if no offset	?	✓ 100%	

- there is many way to cause this, but we were given very limited choices of instructions
- for this data forwarding to happen, the 2 instructions have to be back to back



- `popq` modifies the stack pointer as `valE` (i.e `e_valE = %rsp + 8`)
- so the next instruction needs to use `%rsp` in its ALU operation as well
 - `popq` also takes `%rsp` as both its `valA` and `valB`, so this again would work

Question 9: Pipeline Instruction Processing - Stalls - No Forwarding

- question

Question 9: Pipeline Instruction Processing - Stalls - No Forwarding

Consider the execution of the following instructions in the pipelined CPU. Assuming they are executed in the given order, on the 'pipeline with stalling' implementation, which instruction(s) will experience a data hazard?

Select all that apply.

- subq %r8, %rbx
- divq %r10, %rax
- divq %r12, %rdx
- mulq %rbp, %rax
- andq %r13, %rbp

Select all possible options that apply.

100%

- the fourth line reads from `%rbp` which 3rd line writes too

Next Quiz + Midterm

In Class Assignment

In Class 11

- I11.1: y86 Pipeline Branch Prediction Choices

Assume a pipelined implementation with no forwarding and a branch prediction algorithm that predicts always taken.

```
1 1:  
2 2:    irmovq 1, %rax  
3 3:    irmovq 14, %rbx  
4 4:    irmovq 1, %rcx  
5 5:  
6 6:    loop_start:  
7 7:    rrmovq %rax, %rdi      # rdi = rax  
8 8:    addq  %rcx, %rax       # rax += 1  
9 9:    rrmovq %rbx, %rbp      # rbp = rbx  
10 10:   subq  %rdi, %rbp      # rbp - rdi => rbx - i  
11 11:   jg    loop_start      # if rbp - rdi > 0, keep going  
12 12:  
13 13:   loop_end:  
14 14:   xorq  %r9, %r10  
15 15:   rrmovq %r12, %r9  
16 16:   divq  %r11, %r13  
17 17:   subq  %r8, %r11  
18 18:   addq  %r12, %r10
```

```
19 19: rrmovq %r11, %rsi
```

- first, figure out how many time the loop body runs (there's no top or bottom of loop in this case, just top)

- change the end condition `rbx` to 2
- first iteration: `rbx - rdi = 2 - 1 = 1`, jump to `loop_start` - `rax` becomes 2
- second iteration: `rbx - rdi = 2 - 2 = 0`, continue to `loop_end`
- so with `rbx = 2`, we ran 2 times, so in the real case we should run 14 times

- part 1

For how many cycles will the pipeline stall due to hazards resulting from data dependencies?

46



✓ 100%

- line 7: data hazard on `%rax` (happens once only) → 1 stall
- line 8: data hazard on `rcx` (happens only once) → 1 stall (due to the fact that we already added a stall before line 7)
- line 10: data hazard on both `%rbp` (3 stalls) and `%rdi` (1 stall)
 - we take the bigger of the two so we incur 3 stalls (also solve the other one) → we do this for 14 iterations → $3 \times 14 = 42$ stalls
- line 19: data hazard on `%r11` → 2 stalls
- overall: 46 stalls

- part 2:

- since we always take, we are wrong once at the end
- thus we incur 2 squashes

- I11.2: Processor Performance – computing CPI

You have instrumented your y86 program and found that:

- 11% of the instructions are branches.
- 34% of the branches are conditional.
- 8% of the instructions have a data hazard that cannot be avoided by forwarding.
- There are no return instructions.

If you run this on an implementation with data forwarding, but without any branch prediction, how many cycles per instruction (CPI) will you achieve?

Round your final answer to two decimal places.

Ignore startup effect, i.e., the 4 cycles you need to fill the pipeline.

- let's first compute all the extra cycles we need

$$\left(2 \text{ stalls (for branch)} \times \frac{11}{100} \times \frac{34}{100} \right) + \left(1 \text{ stalls (for mem read)} \times \frac{8}{100} \right) = 0.1548$$

- note: data hazards that cannot be avoided by forwarding are memory reads
- so this is the extra instruction per cycle

- if we had a perfect pipeline with no hazard, we'd have CPI = 1, so finally

$$\text{total CPI} = 1 + 0.1548 = 1.15$$

In Class 12

- I12.1: Simulator

- part 1

In Lab 3, you implemented a simulator for the y86. Which of the following architectures would be likely to make your implementation of the simulator run faster than running on a single-core sequential processor (i.e., not pipelined)?

- A pipelined processor
- Nothing is likely to make it run faster (let's assume that it runs in the first place, and you're not just feeling hopeless, because you were not able to implement parts of it!)
- A distributed system
- A multi-core processor
- A hyper-threaded processor

Select all possible options that apply. 

- pipelined register is the only correct answer here
- this is because the simulator was written in a single-thread manner - if you run on a pipelined processor, you would unlock some form of parallelism
- however, the code cannot automatically take advantage of a hyper-thread or multi-processor architecture

- part 2

Would your answer change if we asked about running all the simulators written by all the students in the class (i.e., if we were grading them)? Discuss with your group why or why not.

- No
- Yes

- yes
- if we were running all the processors written by students, we can dedicate a single core to each of them and achieve parallelism that way

- I12.2: GPUs

A Graphics Processing Unit (GPU), originally designed for, shockingly, graphics, is now widely used for running machine learning algorithms. GPUs were originally designed to apply a sequence of operations to every pixel in an image. Given that description, which of Flynn's architectures best describes a GPU?

- SISD
- MISD
- MIMD
- SIMD

- SIMD is the answer
- GPU is applying the same set of instructions (single instruction stream) to each element of an image (multiple data stream)

- I12.3: Matrix Operations

- part 1

I12.3. Matrix Operations

Numerical operations on large (multiple GB) matrices are fairly common. They are so common, that many architectures, including the x86 provide Advanced Vector Extensions (AVX) which let you pack multiple data items into a 256-bit register and then apply the same operation to each value in the register.

This is an example of what sort of architecture?

- SISD
- SIMD
- MISD
- MIMD

- SIMD
- since you are applying the same operation it's single instruction
- since you are applying the operation to multiple registers it's multiple data stream

- part 2

Now, imagine that I have a multiprocessor system. On each system, I use the AVX instructions to apply a different matrix transformation. I stream data from one processor to the next processor so that it undergoes all the different transformations. What architecture have I now implemented?

- SISD
- MIMD
- MISD
- SIMD

- MIMD
- now you have multiple cores which have multiple instruction set (because they're different transformations) - so now it's multiple instruction
- same reasoning for the data stream part

- I12.4: Kernel Build

I12.4. Kernel Build

The Linux kernel consists of thousands (a few tens of thousands) of C files (as well as a few assembly files). To build a kernel from source, you must compile each of those C files and then collect the compiled object files (the `.o` files) together to be linked. This produces the Linux kernel (i.e., operating system). This process is called a *Linux kernel build*.

On a high end laptop, it takes about 15-20 minutes to build a kernel from sources, using only a single core. The Linux build system allows you to create parallel compilations to take advantage of additional cores.

With your team, discuss the merits of performing a Linux kernel build on a single machine versus distributing across a set of machines (identical to the single machine) connected by a network. Then answer the question below.

Which of the following questions might you ask to help you evaluate the tradeoffs?

- Do the files all start out local to a machine or are they retrieved from a remote server?
- How many stages does the processor on the machine(s) have?
- How much memory is on each machine?
- Is the limiting resource on your machine the CPU, memory, or persistent storage?
- How fast is your network?

- the answer is everything but "How many stages does the processor on the machine have"
 - because that does not really matter
 - any performance gain by making 1 processor faster (more stages) will be dwarfed by parallelising it across multiple slower processors

- I12.5: Web Service Architecture

- part 1

I12.5. Web Service Architecture

You are the new CTO for the new startup "Awesome Web Services" (i.e., AWS). The company is just launching its first web service -- while you expect that it will be a smashing success, you do not yet have many users and money is tight.

The service handles user logins, allows them to browse a catalog of services, and then use any of those services online in an interactive manner. The software development team has built a great server that is multi-threaded and runs quite efficiently on a single machine with up to eight cores.

Using Flynn's taxonomy, how would you describe this architecture?

- MISD
- MIMD
- SISD
- SIMD

- there's multiple cores so it's MI
- the service is doing a lot of things so it will need a lot of different data hence MD

- part 2

You realize (as most startups do) that it is far more cost effective to use a cloud provider than to buy your own machines and run your own servers. Assuming that you are successful, would it be better to use larger and larger servers or to just add more small servers? Discuss with your team and then select an answer below.

- Get more machines
- Get a larger machine

- it's always cheaper and more beneficial to scale horizontally (get more machines)

Sauce

- question 5

Question 5: Understanding Push and Pop

y86 Corp just hired some new architects. They have decided that the `pushq %rA` instruction is unnecessary. They claim that it can be correctly implemented by the following sequence of instructions.

```
irmovq 8, %r8  
rmovq %rA, -8(%rsp)  
subq   %r8, %rsp
```

You have been asked to critique this implementation. Assuming that we do not allow operations of the form `pushq %rsp`, which of the following statements are true?

- Will leave the condition codes in the same state as the `pushq` instruction does. X
- Uses (wastes) an additional register ✓
- Will leave the correct value in %rsp ✓
- Will write to the correct memory location ✓

Select all possible options that apply. ?

○ 75%

1. false: `pushq` DOES NOT set the condition code - but this new implementation WILL
2. true: in `pushq` you only use `%rsp` and `%rA` - here you need `%r8` as well to inc/dec the stack pointer
3. true: it decrements the stack pointer as necessary
4. true: it will write to the stack pointer after it decrements (correct)

- question 6

Question 6: Pointers in y86

Assume that `mlf` is a `uint64_t` variable in a C program. We might initialize such a variable in a y86 program with the following:

```
mlf: .quad 0x6adf
```

Use the dropdowns below to write an instruction that would properly initialize `%rax` to contain a pointer to `mlf`.

Select None for any unused register values; you may use a symbolic name in the offset/dest field or if no value is necessary there, leave it blank.

Instruction	rA	rB	Offset/Dest
irmovq	✓ 100%	None	✓ 100% %rax ✓ 100%
27359	?	x 0%	

- the simulator is pretty smart we could have just gone `irmovq mlf, %rax`
 - so `%rB = %rax` and `offset/dest = mlf`
- question 7

We have spent a lot of time this term iterating over structures in C; now let's do it in y86.

Consider the function, `structIter(size_t structSize)`, which iterates over a buffer containing an array of structures. The buffer starts at the label `data`.

The function returns when it encounters the value 0 in memory. (We call the 0 a sentinel value).

For each structure, `structIter` calls a function `helper` (not shown). `helper` takes a single parameter, which is the address of the current structure.

A former CPSC 313 student ran out of time and wrote the following function, but left out a large part. Drag and drop instructions from the left to correctly complete the function, **which should use proper y86 calling conventions**.

```
1 structIter:  
2     rrmovq %rdi, %rbp  
3     irmovq data, %rbx  
4  
5     loop:  
6         mrmovq 0(%rbx), %rax  
7         # ...  
8         # your own code  
9         # ...
```

```

10    jmp loop
11
12 done:
13    ret
14
15 data:
16    # values here
17    # ...
18    0 # sentinel value

```

- options (the constructed solution is NOT correct):

Drag from here: <div style="border: 1px solid #ccc; padding: 5px; margin-bottom: 10px;">je loop</div> <div style="border: 1px solid #ccc; padding: 5px; margin-bottom: 10px;">addq %rbp, %rbx</div>	Construct your solution here: ? <div style="border: 1px solid #ccc; padding: 5px; margin-bottom: 10px;">pushq %rbx</div> <div style="border: 1px solid #ccc; padding: 5px; margin-bottom: 10px;">rrmovq %rbx, %rdi</div> <div style="border: 1px solid #ccc; padding: 5px; margin-bottom: 10px;">call helper</div> <div style="border: 1px solid #ccc; padding: 5px; margin-bottom: 10px;">popq %rbx</div> <div style="border: 1px solid #ccc; padding: 5px; margin-bottom: 10px;">addq %rax, %rax</div> <div style="border: 1px solid #ccc; padding: 5px; margin-bottom: 10px;">je done</div>
---	--

- my solution

```

1 structIter:
2     rrmovq %rdi, %rbp          # rbp = size(struct)
3     irmovq data, %rbx
4
5 loop:
6     mrmovq 0(%rbx), %rax
7     # my stuff starts
8     addq %rax, %rax           # does %rax * 2 (would stay 0 if it was sentinel)
9     je done                   # check if %rax is sentinel or not - if yes, done
10    rrmovq %rbx, %rdi         # set up parameter for helper
11    call helper               # call helper
12    addq %rbp, %rbx           # structPtr += 1 (since %rbp is the struct size)
13    # my stuff ends
14    jmp loop
15
16 done:
17    ret

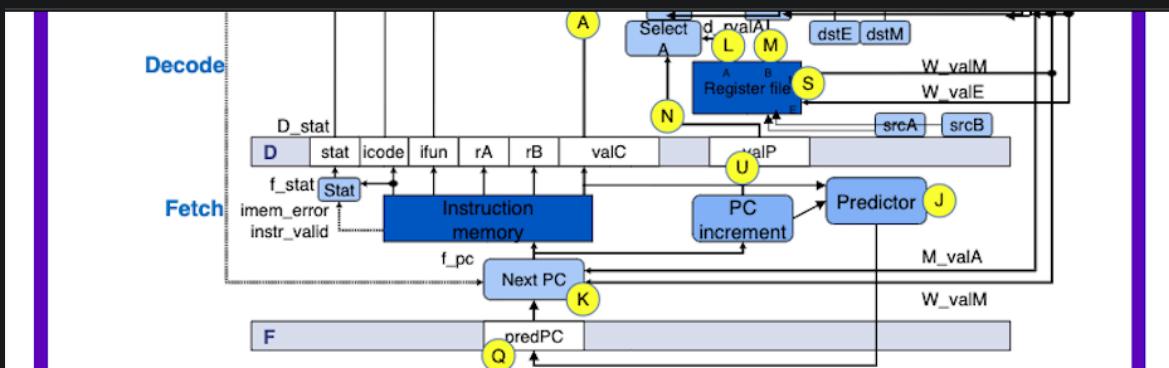
```

```

18
19 data:
20 # values here
21 #
22 0 # sentinel value

```

- basically the line 6 is accessing the int that is at address `%rbx` - if this is zero (the sentinel node), we will jump to done
 - need to add `%rax` to itself to set the condition codes
- then we set up the parameter for `helper()` then call it
- I THINK the push/pop stuff with `%rbx` is a trick and we don't actually have to do that - because `%rbx` is a caller saved register so its value is guaranteed to stay the same after the function returns
- question 9



Consider a **new version of the call instruction** that takes a register as an argument. The programmer puts the address of where to jump into the register prior to the the call instruction. When the call instruction is executed the instruction will push the return address onto the stack and transfer control to the address contained in the specified register. The usage would look like:

```
call %rbx
```

Assume that there is no data hazard on `%rbx`.

With a slight modification, our existing pipelined processor can be made to execute this instruction by routing a signal between two of the points on the above diagram. Indicate the routing of this signal by specifying the starting and ending locations of the signal using the labels from the diagram. If there is more than one way of implementing this choose the one that produces the fewest bubbles/stalls.

Starting Point: **M**

?

x 0%

Ending Point: **U**

?

x 0%

- this question is half missing so it's a bit difficult to understand
- answer key:

Correct answer

Starting Point: L Ending Point: J

It might be tempting to select R as the starting point but observe that because we need to push the return address onto the stack valA must be set to valP and because the stack is involved %rsp will be put in valB. As for the destination it has to be the Predictor. We can't send it to Q as there is no circuitry there to make a decision.

- basically, since the stack pointer and return address is involved
 - valA will have to be valP and valB has to be the %rsp stack pointer
 - this means that R[%rbx] has to get forwarded before it goes into the logic block selectA which will pick valA to be valP
 - the end will have to be the Predictor because there's no way to go from L to Q (no wires available)
 - and Predictor will forward this value to predPC anyways because it's not a conditional jump
- question 10

Consider two different implementations of a y86 processor; one is pipelined and the other is not. Let's say that the sequential implementation has a cycle time that is 4 times that of the pipelined implementation.

Which of the following statements are true?

- The sequential implementation will always achieve a CPI of 1. ✓
- The sequential implementation is likely to have better throughput than the pipelined one.
- The pipelined implementation can have instructions from multiple threads in-flight at the same time. ✗
- The pipelined implementation will always achieve a CPI less than 1.

Select all possible options that apply. ?

○ 75%

1. this is correct an a fact - there are no stalling and the cycle count is different where 1 cycle is 1 instruction executed → this will get you CPI of 1
2. false: pipelined registers almost always have better throughput - that's the whole point
3. false: our pipelined implementation is single-threaded - so we can have multiple instructions from 1 thread in flight at the same time
4. false: in our case, pipelined implementation will at best achieve CPI of 1

- question 11

Consider the following sequence of y86 instructions.

```
1      irmovq    $0xa0, %r12
2      irmovq    $0x15, %r13
3      addq     %r12, %r13
4      jl      skip
5      irmovq    $0x1d9, %r10
skip:
6      rrmovq    %rdi, %r10
```

Assume a y86 pipelined implementation **with forwarding** and a branch prediction algorithm that predicts taken for forward branches and not taken for backward branches. Will the instruction on line 6 stall?

- after adding, **%r13 = big positive number** so we should not jump
 - but the predictor predicts jump
 - though this is lowkey irrelevant
- we go to **skip** (wrongfully) - but the instruction itself doesn't use any registers that past instructions wrote to (even if it did, there's data forwarding so it would require a memory read)

Corrections

Question 1: Designing a Pipeline

- question

Question 2: Designing a Pipeline

Imagine we have 7 phases in a processor. Each must complete before the next one starts; none can be divided up into finer phases. Their durations in order are: **272ps 351ps 327ps 331ps 399ps 278ps 393ps**

We want to build a pipeline that executes the phases above in order, in which each pipeline stage is composed of one or more consecutive phases. Our pipeline registers take **20ps** to operate.

You design a pipeline with exactly 6 stages purely to maximize throughput.

1. How many **ps** long will one cycle be?

642 ps ? ✘ 0%

2. What will its throughput be in giga-instructions per second, rounded to two decimal places?

1.56 GIPS ? ✓ 100%

This question is complete and cannot be answered again.

- part 1

- we only have 5 stages but there are 6 phases → we will have to combine some
- we're only able to combine consecutive phases → we want to combine the 2 smallest consecutive phases, in this case 271 and 351 ps (first 2)
- combine them and you get 623 ps and then add 20 ps and get 643 ps total

- (unfortunate typo here for me)
- part 2
 - since the time is given ps, we can just do $10^3/t$
 - $10^3/643 \approx 1.56$ GIPS

Question 3: Why Parallelism?

- overall point
 - clock speed has stopped getting any faster 10 years ago → we cannot speed up the clock speed
 - however, transistors have consistently gotten smaller so we're able to fit more of them on the chip
 - say a core takes up 10 billion transistors, if transistors get smaller and we can now fit 20 billion transistors on the chip, it makes sense to make 2 cores instead
 - basically, we have more transistors than we know what to do with so we make more cores to utilize parallelism
- question

Question 3: Why Parallelism?

We have discussed various ways to support parallelism in hardware. Which of the following is a motivating factor in supporting and exploring hardware parallelism? (You can click the "?" icon for partial credit information.)

Chips' clock speeds have continued to grow enormously.
 Chips' transistor counts have continued to grow enormously.

Networks of cheap computers can provide large quantities of on-demand computation.
 Each separate core in a multi-core chip design requires significant individual design and production effort.
 We often have room on chips for extra circuit elements beyond the minimum to implement a processor.

Select all possible options that apply.

66%

This question is complete and cannot be answered again.

1. False - clock speed has stopped growing
2. True - chip transistors have continued to grow since they get smaller and we can fit more of them in the chip AND since we can fit more of them in a chip, we want to make more cores with them
3. True - it's cheaper to scale horizontally than vertically
4. False - (while this is true - this is not a reason **for** parallelism)
5. True - basically the 2nd statement, since transistors get smaller, we have more room (to fit more transistors) to make more cores

Question 4: Calling Conventions

- question:

Question 4: Calling Conventions

Consider the following function prototype.

```
int func(int size, int* arr);
```

The assembly implementation of this function is given below, which correctly follows the y86 calling conventions. Note that there are no bugs in this code - it's a fully correct implementation of the original C function. Remember that on the y86, int is a signed 64-bit integer type.

```
func:  
    irmovq 1, %r8  
    irmovq 8, %r9  
    xorq %rax, %rax  
  
loop:  
    andq %rdi, %rdi  
    je done  
  
    mrmovq 0(%rsi), %rdx  
    irmovq 13, %rcx  
    subq %rdx, %rcx  
    jle done  
  
    addq %rdx, %rax  
    subq %r8, %rdi  
    addq %r9, %rsi  
    jmp loop  
  
done:  
    ret
```

What will be the value of `ans` after each of the calls below? Hint: This question will be *much easier* if you first write the C implementation of the function for yourself!

- translate into C, it should look something like

```
1 int func(int size, int* arr) {  
2     int sum = 0;  
3     while(size != 0 && *arr < 13)  
4     {  
5         sum += *arr;  
6         size--;  
7         arr++;  
8     }  
9     return sum;  
10 }
```

- or in Python

```

1 def func(size, arr):
2     sum = 0
3     for i in range(len(arr)):
4         if arr[i] >= 13:
5             break
6         sum += arr[i]
7     return sum

```

- part 1

```

int arr[] = { 7,3,8,2,5,1 };
int ans = func(6, arr);

```

ans = 26



✓ 100%

- you just add everything and you get 26

- part 2

```

int arr[] = { 3,1,13,4,6 };
int ans = func(5, arr);

```

ans = 4



✓ 100%

- you are supposed to add the first 5 elements but `arr[2] >= 13` so we break there
- only get to add up the first 2 elements which is 4

- part 3

```

int arr[] = { 1,8,7,5,3,2,13,6 };
int ans = func(4, arr);

```

ans = 21



✓ 100%

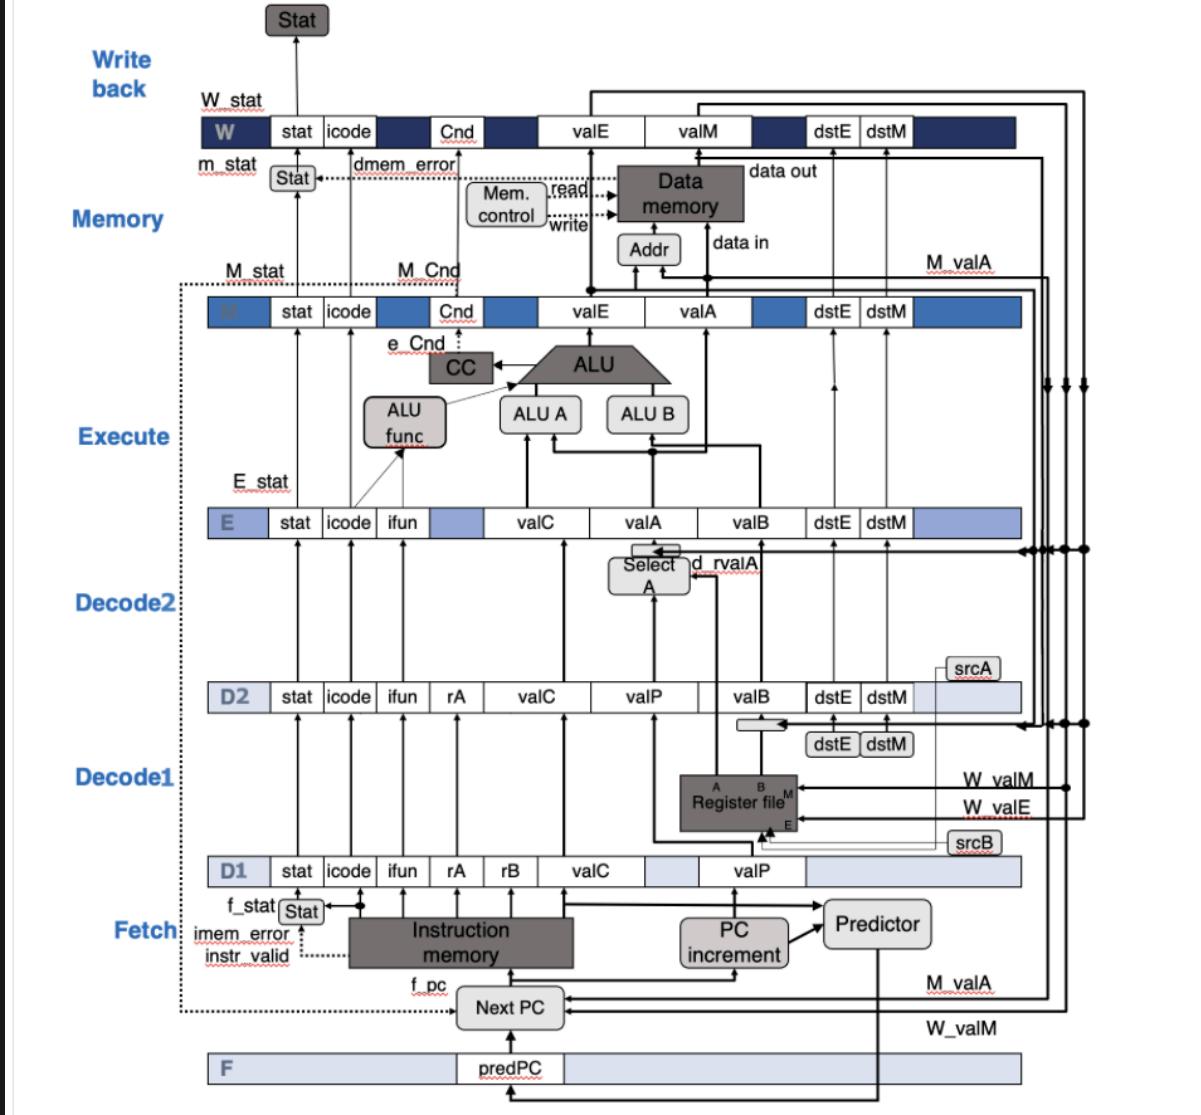
- add up the first 4 elements (none of which are ≥ 13) so normal summation
- add it up and you get 21

Question 5: Pipeline Stalls with a different pipeline

- question

Question 5: Pipeline Stalls with a different pipeline

Imagine we find we can improve performance by modifying the Decode stage. The pipeline below breaks Decode into two stages Decode1 and Decode2. Decode1 prepares `valB` while Decode2 prepares `valA`. The following figure shows what such a pipeline might look like, in the style of our existing pipelined processor diagram:



- part 1

Consider the following program running on such an implementation:

```
Ln_1: irmovq %rax, %rax
Ln_2: andq %rax, %rdi
Ln_3: rmmovq %rdi, 0x65(%rdi)
Ln_4: modq %rdi, %rcx
Ln_5: mrmovq 0x3c(%rcx), %rdx
Ln_6: subq %rcx, %rdx
Ln_7: subq %rdx, %r9
Ln_8: mrmovq 0x14(%r9), %r10
Ln_9: mulq %rdx, %rax
Ln_10: rrmovq %r10, %rdx
```

How many stalls will be inserted by a "PIPE with Stalling" implementation of this new pipeline:

1. between instruction 2 and instruction 3?

4



✓ 100%

2. between instruction 5 and instruction 6?

4



✓ 100%

3. between instruction 9 and instruction 10?

0



✗ 0%

- since this is a stalling pipeline, you need to wait for the previous one to fully finish
- 1. need to wait for line 2 to fully finish while 3 is in fetch → 4 stalls
- 2. need to wait for line 5 to fully finish while 6 is in fetch → 4 stalls
- 3. you actually need to wait for line 8 to fully finish (because %r10)
 - note: also remember to count the stalls you've already taken

- part 2

Assume we forward data similarly to our "PIPE with Data Forwarding" implementation, updated to properly handle the new stage. These forwarding wires and logic are **not necessarily in the diagram**, but the pipeline registers are complete. How many stalls will be inserted in this implementation:

1. between instruction 2 and instruction 3?

1



✓ 100%

2. between instruction 5 and instruction 6?

2



✓ 100%

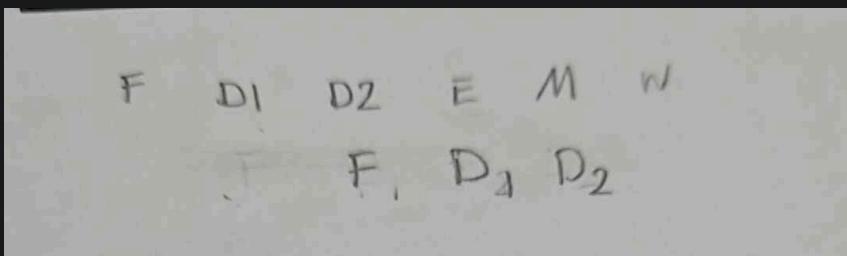
3. between instruction 9 and instruction 10?

0



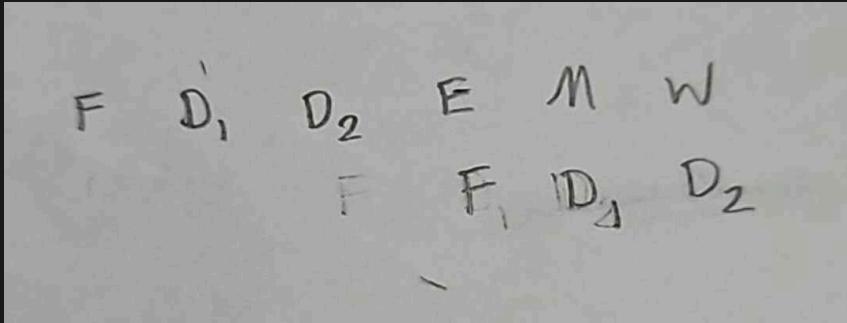
✓ 100%

1. line 3's D_2 stage has to match up with line 2's E stage (because we need %rdi as valB in line 3)



- note that we also need %rdi as valA but it so happens that valB stage came first

2. line 6's D_1 stage has to match up with line 2's M stage (because we need `%rdx` as `valB` in line 6)



- part 3

Our processor can cause a `dmem_error` on an invalid memory access, like trying to read from a "protected" area of memory. We do not want to allow writes to the register file or to memory either from such an instruction or any subsequent instruction.

Assume:

- An instruction that causes `dmem_error` causes no reads or writes during or after the cycle its error is detected.
- When that instruction completes its final stage, it squashes all other stages and halts the machine. (Writes happening during that final cycle **still occur**; no writes at all occur after that cycle.)

Continue assuming we use the pipeline above with Forwarding. Which illegal writes can happen? (You may click the "?" icon below for partial credit information.)

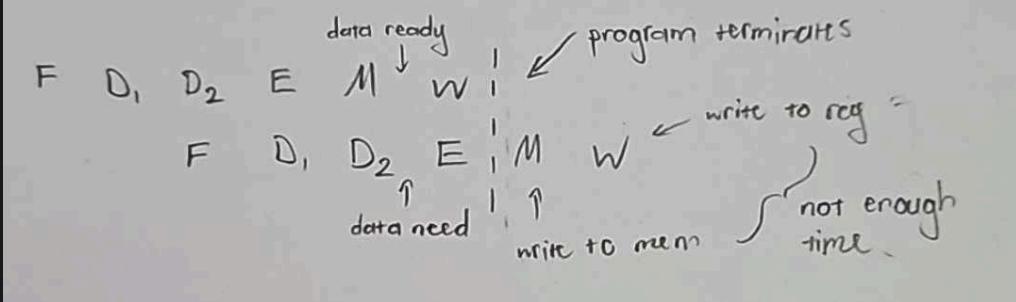
- No illegal writes can occur.
- The illegal instruction can write to the register file.
- The illegal instruction can write to memory.
- A subsequent instruction can write to memory but **not** based on any value from the illegal instruction.
- A subsequent instruction can write to the register file but **not** based on any value from the illegal instruction.
- A subsequent instruction can write to memory including based on a value from the illegal instruction. X
- A subsequent instruction can write to the register file including based on a value from the illegal instruction. X

Select all possible options that apply. ?

50%

This question is complete and cannot be answered again.

- now, it's true that a failed read CAN forward the illegal data to subsequent instructions
- however, the question also states that once the illegal instruction finishes, the program ends
- diagram



Question 6: y86 Pipeline Branch Prediction Choices

- question

Question 6: y86 Pipeline Branch Prediction Choices

Consider the following y86 program

```

1:          irmovq 1, %rax
2:          irmovq 12, %rbx
3:          irmovq 1, %rcx
4:
5:          loop_start:
6:          rrmovq %rax, %rdi
7:          addq   %rcx, %rax
8:          rrmovq %rbx, %rbp
9:          subq   %rdi, %rbp
10:         jg     loop_start
11:
12:         loop_end:
13:         addq   %r13, %r8
14:         divq   %r12, %r9
15:         xorq   %r11, %r10
16:         addq   %r13, %r12
17:         subq   %r13, %r9
18:         mulq   %r8, %r11
19:
20:
```

- the key thing to note here is that the first time into the loop, you possibly incur less stalls, but as the loop iterates you incur more stalls
- part 1

For how many cycles will the pipeline stall due to hazards resulting from data dependencies?

39



✓ 100%

How many cycles will be squashed/cancelled?

2



✓ 100%

```

[Line  7] loop_start: rrmovq %rax, %rdi
  1 stall due to data hazard on %rax
[Line  8] addq %rcx, %rax
  1 stall due to data hazard on %rcx
[Line 10] subq %rdi, %rbp
  3 stalls due to data hazard on %rbp
  3 stalls due to data hazard on %rbp (sometimes also %rdi) (11 times)
[Line 11] jg loop_start
  2 squashes due to branch prediction failure
[Line 18] subq %r13, %r9
  1 stall due to data hazard on %r9

```

Question 7: New Semantics, Old Processor?

- question

Question 7: New Semantics, Old Processor?

Consider the following new y86 instructions with their semantics:

Instruction	Semantics	Example
new1 %rstart[%rinc], V	$m[8 \cdot r[rinc] + r[rstart]] \leftarrow r[V]$	new1 %rbx[%rbp], 5
new2 C(V, %rs), %rd	$r[rd] \leftarrow m[V \cdot r[rs] + C]$	new2 97(8, %rbp), %rbx

Each of these requires changes to the y86 processor, but we are **only** interested in changes to the internals or "ports" (available inputs/outputs) on Data Memory, the Condition Code registers, the ALU, the Register file, and Instruction memory (all the "dark grey" parts of our pipelined processor diagram except **Stat**).

For each instruction, what is the **first** stage of the pipeline in which such a change is needed to implement the given semantics? (It should be clear from the change(s) needed in which stage(s) to count changes to the Register file.)

- new1: X 0%
- new2: X 0%

- new1
 - we need to perform both an addition and a multiplication in a single instruction, which requires changes to the ALU or an additional ALU
 - so the answer is **Execute**
- new2
 - we need to be able to extract a second constant from instruction memory
 - so the answer is **Fetch**
 - (note: later, we will need more ALU functionality to both multiply and add in a single instruction → but question ask for the earliest modification)

Quiz 3

In Class Assignment

In Class 13

- it's a weird interactive thing - try it out if you want to do it again

In Class 14: Calculating Cache Parameters

- I14.1: How big is the cache?

I14.1. How big is the cache?

Assume a system with memory addresses that are 28 bits wide with a direct-mapped cache with characteristics:

- Cache lines are 4 bytes.
- The index consists of 8 bits.

How large is the cache? (in KB) integer KB

Save & Grade Single attempt **Save only** Additional attempts available with new variants ?

- note: the 28 bits wide memory address thing is a red-herring
- the fact that there are 8 index bits implies that there are 2^8 cache slots
- so we have $256 \text{ slots} \times 4 \text{ bytes/slots} = 1024 \text{ bytes} = 1 \text{ kb}$

- I14.2: How many bits in the tag?

I14.2. How many bits in the tag?

Assume a system with memory addresses that are 28 bits wide and a direct mapped cache with the following characteristics:

- The cache is 8 KB.
- Cache lines are 64 bytes.

How many bits of the address are devoted to the cache tag? integer bits

Save & Grade Single attempt **Save only** Additional attempts available with new variants ?

- first we need to figure out the number of bits required for the index and the offset
 - offset: cache lines are 64 bytes long $\rightarrow \log_2(64) = 6$
 - index: we need to first know how many slots are in the cache
 - number of slots: $(8 \times 1024) \text{ bytes} / 64 \text{ bytes per slot} = 128 \text{ slots}$
 - therefore we need $\log_2(128) = 7$ bits for the index
 - so now we have 28 bits per address total, 6 going to the offset, 7 going to the index, so tag is 15 bits long
- I14.3: How many cache lines?

I14.3. How many cache lines?

Assume a system with memory addresses that are 32 bits wide and a direct mapped cache with the following characteristics:

- The cache is 64 KB.
- Cache lines are 64 bytes.

How many lines are in the cache?

integer	cache lines	?
---------	-------------	---

Save & Grade *Single attempt*

Save only

Additional attempts available with new variants ?

- ordering of these questions are a little weird
- the total size of the cache is $64 \times 1024 = 65536$ bytes
- and since we have 64 bytes/slots, the number of slots is $65536/64 = 1024$ slots

- I14.4: How many bits in the offset?

I14.4. How many bits in the offset?

Assume a system with memory addresses that are 24 bits wide and a direct mapped cache with the following characteristics:

- The cache is 16 KB.
- Cache lines are 8 bytes.

How many bits of the address are devoted to the cache offset?

integer	bits	?
---------	------	---

Save & Grade *Single attempt*

Save only

Additional attempts available with new variants ?

- the cache lines are 8 bytes long, so we need enough bits to represent that
- so the number of offset bits is $\log_2(8) = 3$ bits

- I14.5: How many bits are there in an address?

I14.5. How many bits are there in an address?

Assume a direct mapped cache with the following characteristics:

- The cache is 32 KB.
- Cache lines are 4 bytes.
- There are 11 bits in the tag.

How many bits are in an address?

integer	bits	?
---------	------	---

Save & Grade *Single attempt*

Save only

Additional attempts available with new variants ?

- tag: we know it's 11 bits
- index: $(32 \times 1024 \text{ bytes})/4 \text{ bytes per slot} = 8192 \text{ slots} \rightarrow \log_2(8192) = 13$ bits
- offset: $\log_2(4) = 2$
- in total we need 26 bits

- I14.6: How many index bits?

I14.6. How many index bits?

Assume a system with memory addresses that are 24 bits wide and a direct mapped cache with the following characteristics:

- The cache is 16 KB.
- Cache lines are 32 bytes.

How many bits of the address are index bits?

integer bits



Save & Grade *Single attempt*

Save only

Additional attempts available with new variants

- again, we need the number of slots $16 \times 1024 / 32 = 512 \rightarrow \log_2(512) = 9$ bits

- I14.7: Computing cache access time

I14.7. Computing cache access times

You are told that an application has a cache miss rate of 35%. If it takes 8 ns for a cache hit and 56 ns for a cache miss, what will be the average access time for this application? Round your answer to two decimal places.

number (2 digits after decimal)

ns



Save & Grade *Single attempt*

Save only

Additional attempts available with new variants

- cache miss rate of 35% also means there's a cache hit rate of 65%

$$\text{access time} = 0.65 \times 8 + 0.35 \times 56 = 24.8 \text{ ns}$$

- I14.8: Computing cache hit/miss rate

I14.8. Computing cache hit/miss rates

Let's say that you have a 8 KB cache with 32 byte cache lines. Given a 64 byte array, if you access every second byte in the array just once, what will your miss rate be (expressed as a percent, rounded to the nearest integer)?

integer

%



Save & Grade *Single attempt*

Save only

Additional attempts available with new variants

- note: for this question, sometimes they ask for miss or hit rate - PAY ATTENTION
- per cache line, you miss once at the beginning, and hit on all subsequent accesses (note that we're only accessing every other array - so only touching 16 elements in the 32 bytes cache line)

$$\text{miss rate} = \frac{1 \text{ miss}}{16 \text{ accesses}} = 0.0625 = 6.25\%$$

- here, since they ask us to round to a whole percentage, this is 6%

In Class 15: Cache Replacement Policies

- I15.1: Experimenting with replacement (LRU)

I15.1. Experimenting with replacement (LRU)

In this problem, we're going to experiment with deriving access patterns that produce different kinds of behavior with LRU cache replacement.

Let's assume that we have a tiny 4-element, fully-associative cache that uses LRU replacement. The cache stores single byte values (i.e., a line is a single byte).

For each question, enter your accesses as a comma-separated sequence of letters or numbers. For example:

1,1,1,1,1,1,1,1,1,1

- part 1

Part 1

Produce a sequence of 10 accesses that suffers only compulsory misses and uses every entry in the cache.

1,2,3,4,1,2,3,4,1,2

?

✓ 100%

- so the first 4 suffers compulsory misses
- but then every thing after just uses stuff that's already in the cache so there are no more misses
- since we use 4 distinct elements, they will take up different entries in the cache

- part 2

Part 2

Produce a sequence of 10 accesses using more than 5 unique data items that produces 0 hits.

1,2,3,4,5,6,7,8,9,10

?

✓ 100%

- everything here is a compulsory miss - we don't reuse any data once we load it into the cache

- part 3

Part 3

Produce a sequence of 10 accesses using exactly 5 unique data items that produces 0 hits.

1,2,3,4,5,1,2,3,4,5

?

✓ 100%

- so the first 4 are compulsory miss, we then have a cache of $\{1, 2, 3, 4\}$
- we then ask for 5 - another compulsory miss, it'll kick out 1 because that's the least recently used one
→ cache of $\{2, 3, 4, 5\}$
- then we ask for 1, not there, kick out 2 and put 1 back in → $\{3, 4, 5, 1\}$
- basically we keep asking for elements that we just kicked out → there will be no hits

- I15.2: LRU Cache Replacement

I15.2. LRU Cache Replacement

Assume you are given a tiny 4-element, fully-associative cache that uses LRU replacement. The cache stores single byte values (i.e., a line is a single byte). If the cache is empty at the beginning of execution, how many hits will the following string of accesses produce?

G,B,E,D,A,C,B,A,E,E,G,C,C,A,E,D,D,D,G

integer

hits



[Save & Grade](#) [Single attempt](#)

[Save only](#)

Additional attempts available with new variants [?](#)

- first 4 are all misses \rightarrow cache = {G, B, E, D}
 - 5: A is a miss, evict G \rightarrow cache = {B, E, D, A}
 - 6: C is a miss, evict B \rightarrow cache = {E, D, A, C}
 - 7: B is a miss, evict E \rightarrow cache = {D, A, C, B}
 - 8: A is a hit \rightarrow cache = {D, C, B, A} (A at the end because it was most recently used)
 - 9: E is a miss, evict D \rightarrow cache = {C, B, A, E}
 - 10: E is a hit \rightarrow cache = {C, B, A, E}
 - 11: G is a miss, evict C \rightarrow cache = {B, A, E, G}
 - 12: C is a miss, evict B \rightarrow cache = {A, E, G, C}
 - 13: C is a hit \rightarrow cache = {A, E, G, C}
 - 14: C is a hit, cache stays the same
 - 15: A is a hit \rightarrow cache = {E, G, C, A}
 - 16: E is a hit \rightarrow cache = {G, C, A, E}
 - 17: D is a miss, evict G \rightarrow cache = {C, A, E, D}
 - 18: D is hit, cache stays the same
 - 19: D is a hit, cache stays the same
 - 20: G is a miss, evict C \rightarrow cache = {A, E, D, G}
 - so totally that up - we have 8 hits
- 15.3: LFU Cache Replacement

I15.3. LFU Cache Replacement

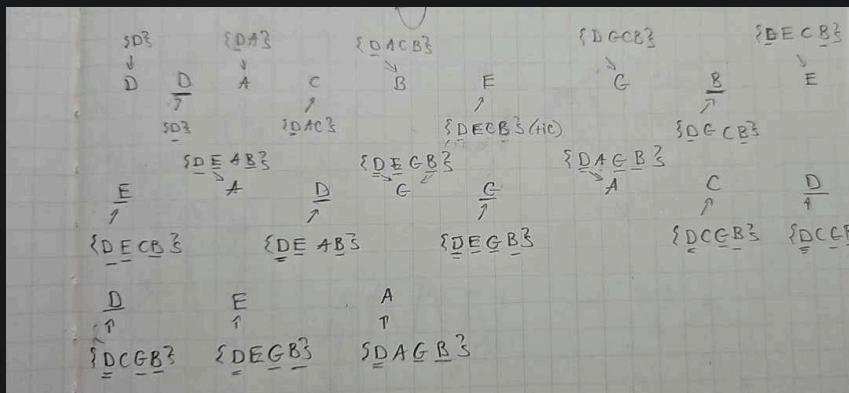
Assume you are given a tiny 4-element, fully-associative cache that uses LFU replacement. The cache stores single byte values (i.e., a line is a single byte). Assume the following:

- In case of a tie, replace the item in the earliest position in the cache (e.g., if the slots are numbered 0, 1, 2, 3, then in a tie between slots 1 and 3, place the new item in slot 1).
- If there are multiple slots empty in the cache, assume that the slot at the earliest position is filled in first.

If the cache is empty at the beginning of execution, how many hits will the following string of accesses produce?

D, D, A, C, B, E, G, B, E, E, A, D, G, G, A, C, D, D, E, A

- there are 7 hits
- diagram



- I15.4: 2-way set associative cache, 1 access

I15.4. 2-way set associative cache, 1 access

The following table shows the state of a 8B, 2-way set associative cache that has 4 sets and 1 byte per cache block. Least recently used blocks per set before any following memory accesses are marked with X.

Set Index	Way 0		Way 1	
	Tag	LRU	Tag	LRU
00	01		11	X
01	01	X	10	
10	11		00	X
11	11	X	00	

Suppose a program accesses the 4-bit address **1101**.

Fill in what tags are stored in the cache after the memory access by typing the binary bits for each tag (be sure to specify ALL bits). Also indicate which blocks per set are least recently used with an 'X'.

- first we need to figure out the cache structure
 - since there are 4 sets (as opposed to slots) → we need $\log_2(4) = 2$ bits for index

- it's 1 byte per cache block → we need $\log_2(1) = 0$ bits for offset
- so the tag is whatever is left
- accessing address **0x1101**
 - index = **0x01** and tag = **0x11**
 - in set index **0x01**, tag **0x11** isn't in either way 0 or way 1, so we need to replace
 - we decide to replace the LRU (in this case the stuff in Way 0), put in our stuff
 - need to update the LRU for that cache index as well
- final answer

Set Index	Way 0		Way 1	
	Tag	LRU	Tag	LRU
00	01		11	X
01	11		10	X
10	11		00	X
11	11	X	00	

- I15.5: 2-way set associative cache, 4 accesses

I15.5. 2-way set associative cache, 4 accesses

The following table shows the state of a 8-byte, 2-way set associative cache that has 4 sets and 1 byte per cache block.

In each set, the least recently used block is marked with an X.

Set Index	Way 0		Way 1	
	Tag	LRU	Tag	LRU
00	01	X	11	
01	01		11	X
10	00	X	11	
11	01		00	X

Suppose a program accesses the following series of 4-bit addresses, in the order given:

0010, 0111, 0100, 0100

and uses the LRU replacement policy.

Fill in the table below with the state of the cache after a program has accessed each of the four addresses described above. Be sure to:

- Fill in all bits in each tag.
- Be sure to place one X in each set to indicate which Way holds the LRU.

- first: **0x0010**
 - go to index **0x10** and look for tag **0x00** - it's there, so this is a hit
 - just need to update and make tag **0x11** the LRU now
- second: **0x0111**
 - go to index **0x11** and look for tag **0x01** - there's also there
 - update to make tag **0x11** the LRU
- third and fourth: **0x0100** and **0x0100**
 - go to index **0x00**, look for tag **0x01** - it's there
 - make tag **0x11** the LRU
- final answer

Set Index	Way 0		Way 1	
	Tag	LRU	Tag	LRU
00	01		11	X
01	01		11	X
10	00		11	X
11	01		00	X

In Class 16: Amdahl's Law and an Introduction to Write Caching

- I16.1: Applying Amdahl's Law

I16.1. Applying Amdahl's Law

Assume that a program spends its time only in ALU operations and Memory operations.

You are told that an application spends 59% of its time in the memory system and takes 388 seconds to complete. If we make the memory system 1.5 times faster, how long will it take for the program to complete?

Round your answer to two decimal places.

number (2 digits after decimal)

seconds



Save & Grade Single attempt

Save only

Additional attempts available with new variants ?

- we have

$$\begin{aligned}
 T_{\text{new}} &= T_{\text{old}} \left(1 - \alpha + \frac{\alpha}{k}\right) \\
 &= 388 \left(1 - 0.59 + \frac{0.59}{1.5}\right) \\
 &= 311.69
 \end{aligned}$$

- I16.2: Applying Amdahl's Law

I16.2. Applying Amdahl's Law

Assume that a program spends its time only in ALU operations and Memory operations.

You are told that an application spends 62% of its time in the memory system and takes 921 seconds to complete. What speedup will we produce if we make the ALU 2 times faster?

Round your answer to two decimal digits.

 speedup

Additional attempts available with new variants

- trick: we are not speeding up memory access - but rather the ALU operations, so $\alpha = 1 - 0.62 = 0.38$
- first calculate the new time T_{new}

$$\begin{aligned}T_{\text{new}} &= T_{\text{old}} \left(1 - \alpha + \frac{\alpha}{k}\right) \\&= 921 \left(1 - 0.38 + \frac{0.38}{2}\right) \\&= 746.01\end{aligned}$$

- then, calculate the ratio to get speed up

$$\begin{aligned}\text{speedup} &= \frac{T_{\text{old}}}{T_{\text{new}}} \\&= \frac{921}{746.01} \\&= 1.23\end{aligned}$$

- or, you can use the formula

$$\begin{aligned}\text{speedup} &= \frac{1}{\left(1 - \alpha + \frac{\alpha}{k}\right)} \\&= \frac{1}{1 - 0.38 + \frac{0.38}{2}} \\&= 1.23\end{aligned}$$

- I16.3: Applying Amdahl's Law

I16.3. Applying Amdahl's Law

Assume that a program spends its time only in ALU operations and Memory operations.

You are told that an application spends 10% of its time in the memory system and takes 676 seconds to complete.

If we want our program to complete in 422 seconds, how many times faster do we need to make the ALU?.

 times faster

Additional attempts available with new variants

- so here, we can speed up the ALU so $\alpha = 0.9$

- need to rework the formula a bit

$$T_{\text{new}} = T_{\text{old}} \left(1 - \alpha + \frac{\alpha}{k}\right)$$

$$\begin{aligned} 422 &= 676 \left(1 - 0.9 + \frac{0.9}{k}\right) \\ k &= \frac{676(0.9)}{422 - 67.6} \\ &= 1.716 \end{aligned}$$

- I16.4: Counting Cache Misses

I16.4. Counting Cache Misses

For a 64KB (65536 bytes) direct-mapped cache with 16-byte blocks, consider the following code snippet.

```
#define SIZE 6907
uint64_t A[SIZE];      //uint64_t is 8 bytes
uint64_t total = 0;
for (int i = 0; i < 10; i++) {
    for (int j = 0 ; j < SIZE ; j++) {
        total += A[j];
    }
}
```

For the purposes of this question, consider accesses to array A, only. Assume that array A is allocated starting at address `0x1000000000`. Assume that no elements of A are in the cache when this code begins running.

How many misses (on array A) will this code produce?



If we change the array to contain 21917 elements, how many misses (on array A) will this code produce?



[Save & Grade](#) [Single attempt](#)

[Save only](#)

Additional attempts available with new variants

- part 1

- let's worry about the inner loop first (but assuming this is our first iteration of the outer loop)
 - every time you access an element, it loads the next 16 bytes i.e next 2 integers, so you get 1 miss, 1 hits on every 2 access
 - meaning for every 2 integers, you get 1 miss, there are 6907 integers, so you have half of those are misses
 - so it'll be something like $6907/(16/8) = 3453.5 \approx 3454$ misses
 - (alternatively, you can think of it as the entire array of `int` is 6907×8 bytes, then we do $55256/16 \approx 3454$ as this is the number of lines we'll need - take a miss on all of those)
 - note: always round up for these things
 - then the outer loop

- the cache is 65536 bytes large - we are accessing the same 6907 integers over and over → that's $6907 \times 8 = 55256$ bytes
- so the entire array A can fit into the cache, so we only have to go to the data source in the first iteration of the outer loop
- for every iteration after that, we'll get all cache hits
- so overall solution is just 3454
- part 2
 - inner loop
 - same thing above
 - it's $21917/(16/8) \approx 10959$ misses
 - note: it's important that you round **at this step**
 - note: this is a very version of this question because the array is twice as big as the cache so you don't have to deal with the case where some of the array overlaps while some does not
 - outer loop
 - total size of the array A is now $21917 \times 8 = 175336$ bytes so it no longer fits in our cache
 - since it doesn't fit, every iteration we incur this cost
 - so total is $10959 \times 10 = 109590$ misses

- I16.5: Introducing to Write Caching

- variation 1

I16.5. Introduction to Write Caching

Consider a cache backed by a data source. If the cache uses a writeback (write allocate) policy, on a write miss, from which place(s) will data be read?

the cache
 the data source
 none of the above

Select all possible options that apply. ?

✓ 100%

In which place(s) will data be written?

the cache
 the data source
 none of the above

Select all possible options that apply. ?

✓ 100%

[Try a new variant](#)

- since it's a write miss - focus on the write allocate policy
- since it's write allocate - we read the data from the data source INTO the cache, then we edit the data **WITHIN** the cache

- note that since it's a write, you are not reading anything from the cache, you're only reading from the data source to put into the cache
- variation 2

I16.5. Introduction to Write Caching

Consider a cache backed by a data source. If the cache uses a writethrough (no write allocate) policy, on a write hit, from which place(s) will data be read?

the cache
 the data source
 none of the above

Select all possible options that apply.

100%

In which place(s) will data be written?

the cache
 the data source
 none of the above

Select all possible options that apply.

100%

[Try a new variant](#)

- it's a cache hit so we focus on the write back policy - in this case it's writethrough
- **note: checking if something exists in the cache does not count as "reading"**
 - so in this case, no reading was done
 - the data gets written to the cache AND memory (right away)

- variation 3

I16.5. Introduction to Write Caching

Consider a cache backed by a data source. If the cache uses a writeback (write allocate) policy, on a write hit, from which place(s) will data be read?

the cache
 the data source
 none of the above

Select all possible options that apply.

100%

In which place(s) will data be written?

the cache
 the data source
 none of the above

Select all possible options that apply.

100%

[Try a new variant](#)

- cache hit so focus on writeback policy
- again, we don't do any reading here because it already exist in the cache
- the data is only written to the cache (and in time, it will get written to memory)

Practice Quiz

R3.1: Caches - Location Ranking Based on Properties

- **important point**: closer to the CPU means A) faster, B) smaller size, C) smaller cache lines, and D) more expensive
- variation 1

R3.1. Caches - Location Ranking Based on Properties

C1 and CA are two different caches on the same system. C1 is slower than CA. Select all the statements below that are most likely to be true in a properly designed system.

C1 is farther from the CPU than CA 

 C1 is smaller than CA

C1 has the same per-byte cost (in \$) as CA

C1's cache line size is less than or equal to that of CA

None of these are likely to be true

Select all possible options that apply. 

 **100%**

[Try a new variant](#)

- variation 2

R3.1. Caches - Location Ranking Based on Properties

C and C* are two different caches on the same system. C is cheaper (in terms of dollars per-byte) than C*. Select all the statements below that are most likely to be true in a properly designed system.

C is smaller than C*

C is closer to the CPU than C*

accesses to C take more time than accesses to C* 

C's cache line size is less than or equal to that of C*

None of these are likely to be true

Select all possible options that apply. 

 **100%**

[Try a new variant](#)

R3.2: Caches - Implementation Properties to Bits

- question

R3.2. Caches - Implementation Properties to Bits

Patrice and Margo have designed a new x86-compatible processor that has an 8-way set associative, 64 KB level 1 data cache for data, with a cache line size of 16 bytes.

How many offset bits does this cache require?

How many **slots** are in this cache?

How many index bits does this cache require?

If the processor has a 48-bit physical address, how many tag bits are there?

- part 1: offset bits

- the cache line is 16 bytes $\rightarrow \log_2(16) = 4$ bits for the offset

- part 2: # of slots

- you first get the size (64 kB) in bytes and divide it by the cache line size

$$(64 \times 1024 \text{ bytes})/16 = 4096 \text{ slots}$$

- part 3: index bits

- since this is not direct-mapped, we need to find the **number of sets**

$$4096 \text{ slots}/8 = 512 \text{ sets}$$

$$\log_2(512) = 9 \text{ bits}$$

so we need 9 bits for the offset

- part 4: tag bits

- so we have 48 bits, 9 goes to the index, 4 goes to the offset, so 35 bits goes to the tag

R3.3: Interpreting cache metadata

- question

R3.3. Interpreting cache meta data

Consider the following depiction of a direct-mapped cache with 8 cache slots, 4-byte blocks, and a main memory that has 12-bit addresses.

Slot Number	Tag	Offset			
		0	1	2	3
0 (0b000)	0b0010110	26	41	0e	fe
1 (0b001)	0b0011011	0f	7e	21	68
2 (0b010)	0b0001001	cf	00	4e	4e
3 (0b011)	0b0101100	04	36	45	89
4 (0b100)	0b1101101	d8	01	76	f1
5 (0b101)	0b0000110	47	f8	ca	20
6 (0b110)	0b0000011	73	d7	45	20
7 (0b111)	0b0011101	80	33	50	b6

- part 1

What is the address, in binary, of the byte in slot 7 at offset 2?

0b



- slot 7 → index 7 → index bits are 011
- offset 2 → offset bits are 10
- the tag bits at slot 7 currently is 0011101
- we know the address is (tag bits + index bits + offset bits) → 0b001110111110

- part 2

Will an access for a single byte at address 0b000010110000 hit in the cache?

- Yes
- No

- skipping 2 offset bits (counting right to left), next 3 bits are the index bits → in this case they are 0b100 (slot 4)
- now check the tag that's in slot 4 (0b1101101) vs the current tag (0b0000101) - they are not the same, so it'll be a cache miss

- part 3

Will an access for a single byte at address 0b00000111010 hit in the cache?

- Yes
- No

- do the same thing as above, we're looking at slot 7 here
- the tag matches up - it'll be a cache hit

- part 4

What value (in hex) will be retrieved from the cache when the processor reads a single byte from address 0b000011010101? If the cache does not hold the requested data, then enter "none".

0x



- looking at slot $0b101 = 5$
- tag matches up so it's a cache hit
- offset is $0b01 = 1$ which is element $0xf8$

R3.4: Determining where data is placed in a cache

- question

R3.4. Determining where data is placed in a cache

Consider a cache with 256 cache sets and a 256-byte cache line size. Which of the following addresses could possibly occupy set 125?

- 0x975088ad927d
- 0x9d792ee3ebac
- 0xe16d43c77d7d
- 0xd089a6077d90

Select all possible options that apply.



Save & Grade Single attempt

Save only

Additional attempts available with new variants

- offset bits: $\log_2(256) = 8$ bits = 2 hexits
- index bits: $\log_2(256) = 8$ bits = 2 hexits
- so we will be looking at hexits at index 2 and 3 (counting from the rightmost bit and calling that index 0)
- we're looking for something with index bits set to $125 = 0x7d$
- 3rd and 4th one fits this criteria

R.3.5: Adding a cache to a processor

- question

R3.5. Adding a cache to a processor

Annoyed by slow memory accesses, Patrice and Margo have decided to add a cache to the PAM processor. For the rest of this problem, ignore writes (you can assume all cache lines are clean when they are evicted).

Before adding the cache, it takes 288 ns to read/write from memory. The new cache has a hit time of 5 ns, and its total miss time (including both transferring the data from memory into the cache and then reading it from the cache) is 288 ns.

What read hit rate will a program need to achieve if we want to produce a memory system speedup of 8 compared to the memory system before we introduced the cache?

Round your answers to two decimal places. If there is insufficient information to answer the question, enter -1.

- first, a speed up of 8 means that

$$\begin{aligned}\text{speedup} &= \frac{T_{\text{old}}}{T_{\text{new}}} \\ 8 &= \frac{288 \text{ ns}}{T_{\text{new}}} \\ T_{\text{new}} &= 36 \text{ ns}\end{aligned}$$

- now we want a hit rate x such that

$$\begin{aligned}T_{\text{hit}}(x) + T_{\text{miss}}(1 - x) &= 36 \text{ ns} \\ 5(x) + 288(1 - x) &= 36 \text{ ns} \\ -283x &= -252 \\ x &= 0.89\end{aligned}$$

- so you'd need a hit rate of 0.89 to get a speed up of 8

R3.6: Adding a second level cache

- question

R3.6. Adding a second level cache

Patrice and Margo are at it again: this time they are adding a second level cache to the PAM processor. The existing cache is called the L1 (for level 1) cache, and the new cache is called the L2 (level 2) cache.

The caches has the following properties:

- The L2 cache is larger than the L1 cache.
- Both caches are writeback, write-allocate.
- The L2 cache is the data source for the L1 cache.
- Main memory is the data source for the L2 cache.
- The L1 cache is a subset of the L2 cache (i.e., if an item is in the L1 cache, it must also be in the L2 cache).
- Data can only be loaded into the L1 from the L2; it cannot be read directly from the data source.
- In both the L1 and L2, dirty cache lines are only written to the data source, when the processor needs to evict the line.
- Both caches are direct mapped.

The performance of the caches is as follows.

- The L1 hit time is 2 ns.
- It takes 24 ns to transfer a cache line between the L1 and L2 caches (both for moving a line from L2 to L1 and for moving a line from L1 to L2).
- It takes 72 ns to transfer a cache line between main memory and the L2.

Answer the following questions. Round your answers to two decimal places. If there is insufficient information to answer the question, enter -1.

◦ part 1

1. How long does a write take if it hits in the L1?

Hit in L1:	number (2 digits after decimal)	ns	?
------------	---------------------------------	----	---

- write-back so the write just have to happen in the cache
- in this case, it just have to happen in the L1 cache
- so it's 2 ns

◦ part 2

2. How long does a read take if it misses in the L1, causing eviction of a **clean** line, but the read request hits in the L2?

L1 Miss (clean):	number (2 digits after decimal)	ns	?
------------------	---------------------------------	----	---

- since the eviction is clean, we don't have to do anything in terms of writing it back
- the line is in L2 - so we transfer from L2 to L1 → 24 ns
- then we have to read from the L1 → 2 ns
- so it's 26 ns

◦ part 3

3. How long does a read take if it misses in the L1, causing eviction of a **dirty** line, but the read request hits in the L2?

L1 Miss (dirty):	number (2 digits after decimal)	ns	?
------------------	---------------------------------	----	---

- first, we need to transfer the dirty line to L2 → 24 ns

- then we're essentially handling a miss in L1 with a clean line
 - the line is in L2 - so we transfer from L2 to L1 → 24 ns
 - then we have to read from the L1 → 2 ns
- so it's 50 ns overall
- part 4

4. How long does a read take if it misses in the L1, causing eviction of a **clean** line, and the read request also misses in the L2, causing eviction of a **clean** line?

L1 miss (clean), L2 miss (clean):	number (2 digits after decimal)	ns	?
-----------------------------------	---------------------------------	----	---

 - since the line is clean, no work done in terms of write back
 - we go to the data source, read it into the L2 → 72 ns
 - transfer from L2 to L1 → 24 ns
 - read from L1 → 2 ns
 - total is 98 ns
- part 5

5. How long does a read take if it misses in the L1, causing eviction of a **clean** line, and the read request also misses in the L2, causing eviction of a **dirty** line?

L1 miss (clean), L2 miss (dirty):	number (2 digits after decimal)	ns	?
-----------------------------------	---------------------------------	----	---

 - no work done on write-back for L1 eviction
 - for L2, you need to evict and write-back to memory → 72 ns
 - then we're essentially handling a read miss in L2 with a clean line
 - we go to the data source, read it into the L2 → 72 ns
 - transfer from L2 to L1 → 24 ns
 - read from L1 → 2 ns
 - total is 170 ns
- part 6

6. Consider only the memory instructions of a program. 55% of the memory instructions hit in the L1. Of those that miss, 92% hit in the L2. Assuming that all evictions evict clean cache lines, what is the average access time?

Average access:	number (2 digits after decimal)	ns	?
-----------------	---------------------------------	----	---

 - first let's define some terms

$$L1_{hit} = \text{percentage we hit the L1 cache} = 0.55$$

$$L2_{hit} = \text{percentage we hit the L2 cache} = (1 - 0.55)(0.92)$$

$$L2_{miss} = \text{percentage we miss the L2 cache} = (1 - 0.55)(0.08)$$
 - for the L2 stuff, just think of it in terms of we only go to the L2 cache if we miss in the L1 cache
 - calculations

$$\begin{aligned} T_{\text{total}} &= (L1_{\text{hit}} \times T_{\text{L1 hit}}) + (L2_{\text{hit}} \times T_{\text{L1 miss, L2 hit}}) + (L2_{\text{miss}} \times T_{\text{L1 miss, L2 miss}}) \\ &= 0.55(2) + (1 - 0.55)(0.92)(26) + (1 - 0.55)(0.08)(98) \\ &= 15.392 \text{ ns} \end{aligned}$$

- the times are from the previous questions

Corrections

Question 1: Designing a Pipeline

- this was from the midterm
 - question

Question 1: Designing a Pipeline

Imagine we have 7 phases in a processor. Each must complete before the next one starts; none can be divided up into finer phases. Their durations in order are: **418ps 338ps 328ps 422ps 280ps 422ps 367ps**

We want to build a pipeline that executes the phases above in order, in which each pipeline stage is composed of one or more consecutive phases. Our pipeline registers take **40ps** to operate.

You design a pipeline with exactly 6 stages purely to maximize throughput.

1. How many **ps** long will one cycle be?
706 ps ? ✓ 100%
2. What will its throughput be in giga-instructions per second, rounded to two decimal places?
1.42 GIPS ? ✓ 100%

This question is complete and cannot be answered again.

- you have to combine consecutive cycles so you want to combine the two that's the smallest
 - in this case 338 and 328 is the smallest consecutive phases that you can combine
 - so now the stages are 418, 706, 422, 280, 422, 367
 - and since it's parallel, it'll take as long as the longest stage so it's 706 ps
 - from there you can calculate the throughput by doing $10^3 / 706 \approx 1.42$ GIPS

Question 2: Caches - Bits to Implementation Properties

- question

Question 2: Caches - Bits to Implementation Properties

A processor has 32-bit memory addresses, the bits of which are broken into tag, index, and offset as described in the following table.

Tag	Index	Offset
19 bits	9 bits	4 bits
XXXXXXXXXXXXXX	XXXXXXXXXXXX	XXXX

- part 1

How large is each cache line in this cache? bytes ? ✓ 100%

- as we know, how big the cache line is influence the number of offset bits, and vice versa
- here, there are 4 offset bits $\rightarrow 2^4 = 16$ byte cache line

- part 2

How many **sets** are in this cache? ? ✓ 100%

- now this is related to the number of index bits
- there are 9 index bits $\rightarrow 2^9 = 512$

- part 3

The size of this cache is 32 KB.

What is the associativity of this cache? -way ? ✓ 100%

- $32\text{ KB} = 32768\text{ bytes}$
- so we know we have 512 sets \rightarrow each set have $32768/512 = 64$ bytes/set
- now we know every line has 16 bytes $\frac{64\text{ bytes}}{\text{set}} \times \frac{\text{lines}}{16\text{ bytes}} = \frac{4\text{ lines}}{\text{set}}$
- which means that it is 4-way associative

Question 3: Ordering Memory

- question

Question 3: Ordering Memory

Each line below corresponds to some level in the memory hierarchy. For each level you are provided with some details about the cost, size, or performance of the memory at that level.

Use all the memories to build a hierarchy of memories, placing the memory closest to the CPU at the top and the memory farthest from the CPU at the bottom.

Drag from here:

Construct your solution here: ?

Latency: 4 ns, Size: 16 KB

Cost: 50 cents per KB, Size: 2 MB

Cost: 51 cents per MB, Latency: 24 ns

Latency: 90 ns, Size: 8 GB

Cost: 11 cents per GB, Size: 8 TB

- the general idea is
 - the caches get smaller as you get closer to main memory
 - the caches get more expensive as you get closer to main memory
- (keep in mind that some of these are a little ambiguous because they have differing metric - some in terms of size, some in term of cost → answer key states that there may be multiple correct orders)

Question 4: Compare Cache Configurations (associativity)

- question

Question 4: Compare Cache Configurations (associativity)

You are told that a system has a 32-bit address space and a 32 KB cache with 128 byte cache lines.

If the cache is direct-mapped, how many bits will be used for the index?

8



✓ 100%

If we change the cache to be 8-way set-associative, and make no other changes to the cache, how many bits will be used for the index?

5



✓ 100%

If we change the cache to be fully-associative, and make no other changes to the cache, how many bits will be used for the index?

0



✓ 100%

This question is complete and cannot be answered again.

- part 1
 - we need to figure out how many cache slots there are → this is $32768 \text{ bytes}/128 = 256$ slots
 - to represent 256 slots, we need $\log_2(256) = 8$ bits
- part 2
 - when we make the cache 8 way associative, we have now 8 times less cache slots → $256 \text{ slots}/8 = 32$ sets
 - to represent this we need $\log_2(32) = 5$ bits
- part 3
 - fully associative means any cache line can go into any slot - there is no need for an index → 0 bits needed
 - (or you can think of it as 1 huge set - we need $\log_2(1) = 0$ bits)

Question 5: Interpreting cache meta data

- question

Question 5: Interpreting cache meta data

Consider the following depiction of a 2-way set associative cache with 8 sets, 4-byte cache lines, and a main memory that has 12-bit addresses.

Set Number	Way-0					Way-1				
	Tag	Offset				Tag	Offset			
		0	1	2	3		0	1	2	3
0 (0b000)	0b0111100	1b	e0	b3	05	0b0110011	26	c1	7f	af
1 (0b001)	0b1101100	59	63	6b	3e	0b0010010	ed	ae	f5	f0
2 (0b010)	0b0001110	da	d1	a4	b4	0b0111100	02	ed	d4	29
3 (0b011)	0b0110010	89	bf	25	a9	0b1111000	66	d8	89	bc
4 (0b100)	0b0100101	17	c2	7e	57	0b1011001	a6	3c	3e	f4
5 (0b101)	0b1100101	ff	c3	fb	71	0b0011001	af	d4	3b	b5
6 (0b110)	0b0011101	59	8a	0a	fb	0b0011011	fd	23	fc	6c
7 (0b111)	0b0000110	82	7f	30	a5	0b0001100	d3	cb	3d	bb

- part 1

What is the address, in binary, of the byte in way 0 of set 2 at offset 2?

0b 000111001010



✓ 100%

- you kinda have to reconstruct this, so we know the following from the picture
 - tag: 0b 0001110 (set 2, way 0)
 - index: 0b 010 (because it was set 2 - hence index 2)
 - offset: 0b 10 (because it was offset 2 - only 2 bits because there's only 4 total offset)
- putting them all together and you'll get 0b 000111001010

- part 2

Will an access for a single byte at address 0b110010110010 hit in the cache?

Yes

No

✓ 100%

- breaking down the address
 - offset: 10
 - index: 100
 - tag: 1100101
- the tag at index 100 is 0b0100101 or 0b1011001 so it does not match the tag we have → it'll be a miss

- part 3

Will an access for a single byte at address 0b011001001100 hit in the cache?

Yes ✓

No

✓ 100%

- breaking down the address
 - offset: 00
 - index: 011
 - tag: 0110010

- at index 011, the tag is 0b0110010 which matches → so it will hit

- part 4

What value (in hex) will be retrieved from the cache when the processor reads a single byte from address 0b110010110101? If the cache does not hold the requested data, then enter "none".

0x c3



✓ 100%

- breaking down the address
 - offset: 01
 - index: 101
 - tag: 1100101

- if we got to index 101, the tag there is 0b1100101 so it matches → we get a hit

- then in the cache set (way-0), offset 1 is c3

Question 6: Masking Out the Index

- question

Question 6: Masking Out the Index

Consider a cache with 512 cache sets and a 32-byte cache line size.

Give a mask in hexadecimal that will keep only the index bits for this cache from an address.

0x 3FE0



✓ 100%

Consider the following specific address in hexadecimal: 0x31dc10c6bbc0. Give **only the index bits** from this address and give the bits **in binary**. (So, for example, if there were 40 address bits of which 15 were index bits, your answer would have exactly 15 bits.)

0b 111011110



✓ 100%

- part 1

- we know that we'll need $\log_2(32) = 5$ bits for the offset
- we also know we'll need $\log_2(512) = 9$ bits for the index
- so we need something like `0b (insert leading 0s) 0011 1111 1110 0000`
- which turns out to be `3FE0`

- part 2

Consider the following specific address in hexadecimal: `0x31dc10c6bbc0`. Give **only the index bits** from this address and give the bits **in binary**. (So, for example, if there were 40 address bits of which 15 were index bits, your answer would have exactly 15 bits.)

0b	111011110	?	✓ 100%
----	-----------	---	--------

- since both the offset and index takes up 14 bits, this means we only need the last **4 hex digits** (`0xbbc0`)
- we convert this to binary: `0b 1011101111000000`
- then we cut off the last 5: `0b 10111011110`
- then we keep the last 9: `0b 111011110`
- that's the index bits

Question 7: Cache misses the best and worst

- question

Question 7: Cache misses the best and worst

Consider a **fully associative cache with 4 entries**. Cache slots are numbered 0 to 3.

When placing something in the cache, if there are multiple cache slots empty, we place the item in the cache slot with the lowest number

Suppose that the cache line eviction policy is to randomly choose a victim for eviction.

Using the naming conventions similar to those used in class, assume that the following cache lines are accessed in the following order:

A,A,D,D,C,B,E,B,E,B,C,E

Assume that one is extremely lucky, and the randomly selected victim turns out to be the best possible choice. That is, after a sequence of accesses the number of misses is the minimum possible for the given access sequence.

How many hits will this access string produce?

7	?	✓ 100%
---	---	--------

This question is complete and cannot be answered again.

- we use Belady where whenever we need to evict something, we look far into the future and evict that one that we will not need again for the longest time
- (the other variations will ask for the worst eviction policy - which mean that you look into the future, and evict the one that you'll need again very soon)
- note: **be careful about whether they're asking about hits or misses**

Question 8: Adding a second level cache

- question

Question 8: Adding a second level cache

Patrice and Margo are at it again: this time they are adding a second level cache to the PAM processor. The existing cache is called the L1 (for level 1) cache, and the new cache is called the L2 (level 2) cache.

The caches have the following properties:

- The cache line size for the two caches are identical.
- The L2 cache is larger than the L1 cache.
- The L2 cache is the data source for the L1 cache.
- Main memory is the data source for the L2 cache.
- The L1 cache is a subset of the L2 cache (i.e., if an item is in the L1 cache, it must also be in the L2 cache).
- Data can only be loaded into the L1 from the L2; it cannot be read directly from main memory.
- Both caches are direct mapped.

The performance of the caches is as follows.

- The L1 hit time is 3 ns.
- It takes 49 ns to transfer a cache line between the L1 and L2 caches (both for moving a line from L2 to L1 and for moving a line from L1 to L2).
- It takes 82 ns to transfer a cache line between main memory and the L2.

For this problem, ignore writes (you can assume all cache lines are clean when they are evicted).

Answer the following questions. Round your answers to two decimal places. If there is insufficient information to answer the question, enter -1.

- part 1

1. How long does a read take if it hits in the L1?

Hit in L1:

3

ns



✓ 100%

- it's given in the question, a hit in L1 is 3 ns

- part 2

2. How long does a read take if it misses in the L1, causing eviction of a cache line, but the read request hits in the L2?

L1 Miss:

52

ns



✓ 100%

- (remember we're assuming evictions are clean)

- miss in the L1 it hits in the L2, so we transfer data from L2 to L1 → takes 49 ns
- then we read it from the L1 (a "hit" now if you want to think about it that way) → 3 ns
- 52 ns overall
- part 3

3. How long does a read take if it misses in the L1, causing eviction of a cache line, and the read request also misses in the L2, causing eviction of a cache line?

L1 miss, L2 miss: 134

ns

?

✓ 100%

- miss in L1, miss in L2, hit in main memory
- transfer data from main memory to L2 → 82 ns
- transfer data from L2 to L1 → 49 ns
- read data from L1 → 3 ns
- overall it's 134 ns

- part 4

4. Consider only the memory instructions of a program. 61% of the memory instructions hit in the L1. Of those that miss, 83% hit in the L2. What is the average access time?

Average access: 27.55

ns

?

✓ 100%

- this question builds on the previous answers

$$\begin{aligned}\text{average access} &= 0.61(3 \text{ ns}) + (1 - 0.61)(0.83)(52 \text{ ns}) + (1 - 0.61)(1 - 0.83)(134 \text{ ns}) \\ &= 27.5466 \\ &\approx 27.55 \text{ ns}\end{aligned}$$

- it's basically just the percentage of time things happen multiply by the time it'll take when that thing happens

Question 9: Cache miss/hit rates depending on stride length

- question

Question 9: Cache miss/hit rates depending on stride length

Assume you have a cache with 64-byte cache lines and a program that iterates over a large array (significantly larger than the cache) with a stride of size S.

If we decrease S to one-third its size, but keep everything else the same, the number of **misses per cache line** will likely:

Assume that all strides are ≥ 1 and $<$ the size of a cache line.

- almost double
- decrease by about 66%
- stay the same ✓
- decrease by about 50%
- increase by about 66%

✓ 100%

Now we iterate over the array with a stride of size T. If we triple T, keeping the cache line size the same, resulting in the new stride length being **larger than 64 bytes**. How will the **miss rate** most likely change?

- decrease by about 50%
- equal 1 ✓
- Cannot be determined
- almost double
- increase by about 66%

✓ 100%

This question is complete and cannot be answered again.

- part 1
 - since the stride is less than the size of a cache line, this makes things easier
 - let say our cache line is 16 bytes
 - if our stride is 1: in every cache line, we take 1 compulsory miss and get 15 hits
 - if our stride is 2: in every cache line, we take 1 compulsory miss and get 7 hits (ish)
 - if our stride is 3: in every cache line, we take 1 compulsory miss and get 5 hits
 - point is: the hit count will go change as we increase/decrease the stride, but miss count stay the same
 - note that this is NOT the case if the stride starts getting bigger than cache line - it becomes a bit more complicated
- part 2
 - in the case that the new stride length is larger than the cache line size, the miss rate must be 1 (the hit rate must be 0), because we never access more than one item per cache line

Question 10: Adding a cache to a processor

- question

Question 10: Adding a cache to a processor

Annoyed by slow memory accesses, Patrice and Margo have decided to add a cache to the PAM processor. For this problem, ignore writes (you can assume all cache lines are clean when they are evicted).

Before adding the cache, it takes 148 ns to read/write from memory. The new cache has a hit time of 6 ns, and its total miss time (including both transferring the data from memory into the cache and then reading it from the cache) is 154 ns.

What read hit rate will a program need to achieve if we want to produce a memory system speedup of 2 compared to the memory system before we introduced the cache?

Round your answers to two decimal places. If there is insufficient information to answer the question, enter -1.

read hit rate = 0.54



✓ 100%

This question is complete and cannot be answered again.

- a speed up of 2 means that the new time is $148/2 = 74$ ns
- we want a hit rate x such that

$$\begin{aligned} 74 \text{ ns} &= x(6 \text{ ns}) + (1 - x)(154 \text{ ns}) \\ &= 6x + 154 - 154x \\ 148x &= 80 \\ x &\approx 0.54 \end{aligned}$$

- (you could have also used the speed up formula and solved it from there - this just made more intuitive sense to me)

Quiz 4

In class Assignment

In Class 17: Write Caching

- I17.1: Determining Hit or Miss

I17.1. Determining Hit or Miss

Given a 2048 byte direct-mapped cache with 128 byte cache lines, write a C expression, in terms of an address A, a valid bit V, and a tag T (both of which are in the slot indexed by the index bits of the address), that returns nonzero if the given address will generate a cache hit and 0 if it will generate a miss. You may assume that right shifting an address will not cause sign-extension.

Your expression can use parentheses and only the following operators, constants and symbols:

- basically, you had to check if the cache line valid (i.e $V == 1$) and that the tag in the address is the same as the tag T
- you can get the tag from the address by shifting it by an appropriate amount
 - in this case: index is $\log_2(2048/128) = 4$ bits and offset is $\log_2(128) = 7$ bits
 - so we need to right shift by 11 bits
 - (usually we need to AND it at the end so that the front is all 0 but here we magically assume that's not an issue)
- code: `(V == 1 && (A >> 11) == T)`

- I17.2: Determining Need for Writeback

I17.2. Determining Need for Writeback

Given a 2048 byte direct-mapped writeback cache with 64 byte cache lines, write a C expression, in terms of an address A, a valid bit V, a dirty bit D, and a tag T (all three of which are in the slot indexed by the index bits of the address), that returns nonzero if the given address should force a writeback and 0 otherwise. You may assume that right shifting an address will not cause sign-extension.

Your expression can use parentheses and only the following operators, constants and symbols:

- we need to write back when there's a conflict and we are evicting
 - that means the cache line is valid, our tag is the same as the tag currently there (hence, conflict) and the data is dirty so we have to write it back
 - here, to get the tag we shift by 11
- note: conflict means `T != A >> 11`
- code: `(V == 1 && D == 0 && (T != (A >> (5 + 6))))`

- I17.3: Writeback cache

- variation 1

I17.3. Writeback cache

Given a direct-mapped cache with 4 2-byte cache lines and a main memory that has 8-bit addresses.

A load instruction is trying to load from address 0x79. If there is a cache hit, determine what data would be retrieved from the cache (be sure to add the 0x prefix). If there is a cache miss, determine whether a write-back would happen. If there would be a writeback type "writeback". If there would NOT be a writeback type "replace".

Index	Valid	Dirty	Tag	Block offset(0)	Block offset(1)
00	1	0	0 1111	0x83	0x97
01	1	1	0 1101	0x90	0x88
10	1	1	0 1100	0x52	0x80
11	1	1	1 0101	0x15	0x21

Enter all hex digits in your answer. Enter 0x before the number, except for misses (type "writeback" or "replace").

- `0x79 = 0b1111 00 1`
- index = 00, offset = 1, the tag matches up, so we return `0x97`

- variation 2

I17.3. Writeback cache

Given a direct-mapped cache with 4 2-byte cache lines and a main memory that has 8-bit addresses.

A load instruction is trying to load from address 0xf9. If there is a cache hit, determine what data would be retrieved from the cache (be sure to add the 0x prefix). If there is a cache miss, determine whether a write-back would happen. If there would be a writeback type "writeback". If there would NOT be a writeback type "replace".

Index	Valid	Dirty	Tag	Block offset(0)	Block offset(1)
00	1	1	1 0011	0x74	0x78
01	1	0	1 1100	0x86	0x22
10	1	0	0 0010	0x23	0x53
11	0	0	0 0011	0x71	0x52

- $0xf9 = 0b11111\ 00\ 1$
- the tag doesn't match up so this is a miss
- since the data is valid and dirty - we have to write it back to memory

- 17.4: Computing cache write access time

- variation 1 (write back + write allocate)

I17.4. Computing cache write access times

Consider a system with a writeback (write allocate) cache. You are told that an application has a cache write hit rate of 77%. 16% of the time that we take a miss, the cache line being replaced is dirty. If it takes 25 ns to read/write the cache and 75 ns to read/write the data source, what will be the average write time for this application?

Assumptions: on a miss, if the evicted cacheline is clean, the total miss time should include both the time to read data from the data source and the time to update the cache.

Round your answer to two decimal places.

- time for each actions
 - cache hit = 25
 - cache miss and evict clean = 100 (25 to read from the cache, 75 to read from data source and overwrite current line)
 - cache miss and evict dirty = 175 (25 to read from the cache, 75 to write dirty data to source, 75 to read data from source)

- math

$$\begin{aligned}\text{total time} &= \text{hit} + \text{evict clean} + \text{evict dirty} \\ &= 0.77(25) + (0.23)(0.84)(100) + (0.23)(0.16)(175) \\ &= 45.01 \text{ ns}\end{aligned}$$

- variation 2 (write through + no write back)

I17.4. Computing cache write access times

Consider a system with a write through (no write allocate) cache. You are told that an application has a cache write miss rate of 35%. If it takes 5 ns to read/write the cache and 15 ns to read/write the data source, what will be the average write time for this application?

Assumptions: We can write to the cache and the data source in parallel.

Round your answer to two decimal places.

- since it's a write through cache, no matter if writes are hits or misses - we always write to the data source
 - write hit: we write to both cache and data source (but since they can happen in parallel we take the bigger one of these 2)
 - write miss: write straight to data source
- the answer is 15 ns

(skip in-class 18 because it's a try it out type thing - basically you want row-wise access in the inner most loop)

In Class 19: Cache Coherence (MESI)

- I19.1: MESI Protocol States
 - variation 1

I19.1. MESI Protocol States

If we are following the MESI protocol, given a cacheline in state S, what should the new state for the cache line be after a remote write?

- E
- S
- I
- M

✓100%

In the blank below, indicate if this core needs to take another action. If no action is needed, enter **none**.

Otherwise, enter one of **invalidate** or **writeback**

. action= ?

[Try a new variant](#)

- since someone is writing to the same data that you have, the one you have is now STALE
- so you need to go to invalidate and just get rid of what was there before

- variation 2

I19.1. MESI Protocol States

If we are following the MESI protocol, given a cacheline in state E, what should the new state for the cache line be after a local read?

- I
- E
- M
- S

In the blank below, indicate if this core needs to take another action. If no action is needed, enter **none**. Otherwise, enter one of **invalidate** or **writeback**

. action= ?

Save & Grade Single attempt **Save only**

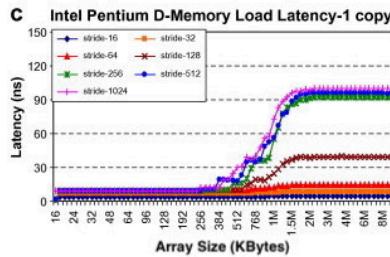
Additional attempts available with new variants ?

- since you're still the only with this data, you stay in exclusive
- so E and None
- (all the variations can be done by tracing the flow diagram)

- I19.2: Deriving Cache Line Size

I19.2. Deriving Cache Line Size

Consider the following data from running an experiment where we perform strided access while iterating over arrays of varying sizes.



What is the size of the cache line of the last level cache?

- 32 bytes
- 64 bytes
- 128 bytes
- 256 bytes
- 512 bytes
- It is impossible to tell

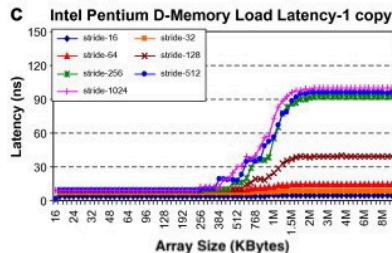
✓100%

- (the last cache is referring to more towards the end of the graph, where performance flattens out)
- since stride size of 256, 512, and 1024 bytes have the same performance - this means that we're likely missing on every access, so we can conclude that line size is 256

- I19.3: Deriving Memory Access Time

I19.3. Deriving Memory Access Time

Consider the following data from running an experiment where we perform strided access while iterating over arrays of varying sizes.

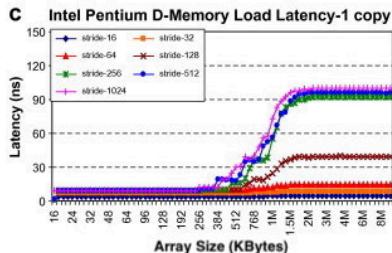


What is the access time of main memory?

- since towards the end, we're missing on all accesses - this means that we're likely just accessing things from main memory every time
 - so the performance towards the end is main memory performance → we see that it's 90 ns
- I19.4: Deriving the size of the last level cache

I19.4. Deriving the size of the last level cache

Consider the following data from running an experiment where we perform strided access while iterating over arrays of varying sizes.

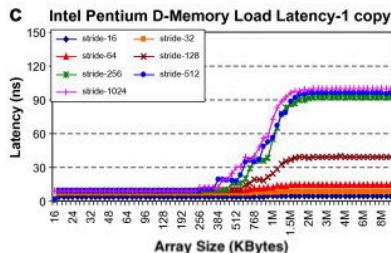


What is the best estimate of the size of the last level cache?

- it starts taking longer at around 256 MB and peaks around 1 MB
 - this means that we start falling out of the cache somewhere between there (array no longer fully fits in the cache so we start having more misses)
- I19.5: Deriving size of main memory

I19.5. Deriving the size of memory

Consider the following data from running an experiment where we perform strided access while iterating over arrays of varying sizes.



What is the best estimate of the size of main memory?

- there's no real indication of the max size of main memory
- since performance is constant until about 8MB - we can say that memory is at least as big as 8 MB (or more)
- (though realistically, we can't determine that from this graph, since if we run out of memory, everything just crashes)

(skip in class I20 because it's writing pseudo code)

In Class 21: Fun with File Descriptors

- I21.1: Single process FD behaviour (offset)
 - variation 1

I21.1. Single process FD behavior (offset)

A process P executes the following code.

```
int fd1 = open("myfile.txt", O_RDWR, 0);
int fd2 = fd1 + 1;
assert(dup2(fd1, fd2) == fd2);

char buf[4096];
ssize_t nbytes = read(fd1, buf, 224);
nbytes = read(fd2, buf, 717);
```

Assume that both files accessed are sufficiently large, that an EOF is not returned, and that read and write operations successfully read/write the number of bytes requested.

What will fd2's offset be? (If there is not enough information to determine the offset enter -1.)

941



✓ 100%

- first `fd1` will return as 3 or something like that (since 0, 1 and 2 are always taken)
- `fd2 = 3 + 1 = 4` (though this doesn't make any sense since we didn't open any file, this fd doesn't exist in the FD table)
- we then call `dup2` which makes `fd2` point to the same thing `fd1` is (the `myfile.txt` file) → so now `fd1/2` share the same offset

- we read 224 bytes then read 717 bytes, that's 941 bytes overall
- since `fd2`'s offset is the same as `fd1`'s offset - the answer is 941
- variation 2

I21.1. Single process FD behavior (offset)

A process P executes the following code.

```
int fd1 = open("myfile.txt", O_RDWR, 0);
int fd2 = open("myfile.txt", O_RDONLY, 0);

char buf[4096];
ssize_t nbytes = read(fd1, buf, 166);
nbytes = read(fd2, buf, 336);
```

Assume that both files accessed are sufficiently large, that an EOF is not returned, and that read and write operations successfully read/write the number of bytes requested.

What will fd2's offset be? (If there is not enough information to determine the offset enter -1.)

336



✓ 100%

- `fd1` and `fd2` are completely separate (i.e they have their own offsets)

- I21.2. Single process FD behavior (value read)

I21.2. Single process FD behavior (value read)

A process P executes the following code.

```
int fd1 = open("myfile.txt", O_RDWR, 0);
int fd2 = open("myfile.txt", O_RDONLY, 0);

char buf[4096];
int filler = 'c';
char writebyte = 'X';
char readbyte;
memset(buf, filler, 4096);
ssize_t nbytes = write(fd1, buf, 160);
nbytes = write(fd1, &writebyte, 1);
nbytes = read(fd2, buf, 69);
nbytes = read(fd2, &readbyte, 1);
```

Assume that both files accessed are sufficiently large, that an EOF is not returned, and that read and write operations successfully read/write the number of bytes requested.

What is the value of readbyte?

- There is not enough information to determine the value
 - X
 - c
- ✓100%**

[Try a new variant](#)

◦ code analysis

```
1 int fd1 = open("myfile.txt", O_RDWR, 0);
2 int fd2 = open("myfile.txt", O_RDONLY, 0);
3 // note that these are separate FD - they DO NOT share offsets
4
5 char buf[4096];
6 int filler = 'c';
7 char writebyte = 'X';
8 char readbyte;
9 memset(buf, filler, 4096); // buf = [c, c, c, c, c, ... 4096 times]
10 ssize_t nbytes = write(fd1, buf, 160); // myfile = [c, c, c, ... 160, other...]
11 nbytes = write(fd1, &writebyte, 1); // myfile = [c, c, ... 160, X, other, ...]
12 nbytes = read(fd2, buf, 69); // read the first 69 bytes (all c's)
13 nbytes = read(fd2, &readbyte, 1); // read the 70th byte - still c
```

◦ variation 2

I21.2. Single process FD behavior (value read)

A process P executes the following code.

```
int fd1 = open("myfile.txt", O_RDWR, 0);
int fd2 = fd1 + 1;
assert(dup2(fd1, fd2) == fd2);

char buf[4096];
int filler = 'c';
char writebyte = 'X';
char readbyte;
memset(buf, filler, 4096);
ssize_t nbytes = write(fd1, buf, 459);
nbytes = write(fd1, &writebyte, 1);
nbytes = read(fd2, buf, 459);
nbytes = read(fd2, &readbyte, 1);
```

Assume that both files accessed are sufficiently large, that an EOF is not returned, and that read and write operations successfully read/write the number of bytes requested.

What is the value of readbyte?

- There is not enough information to determine the value ✓
- c
- X

✓ 100%

- here, `fd1` and `fd2` point to the same thing in the OFT (same offset)
- so we're continuing to read from where `fd1` left off writing → we don't know what might be in the file

- I21.3. FD Behavior with threads

- variation 1

I21.3. FD Behavior with threads

A process P opens the file "myfile.txt" and assigns the returned fd to the global variable `fd1`. P then spawns two threads, T1 and T2. When spawning the threads P passes in `fd1` as the argument to each of the threads. After spawning the threads, P immediately waits for both of them to complete.

The threads T1 and T2 execute the following functions, respectively.

```
void thread1(int fd) {
    int fd2 = fd;

    char buf[4096];
    ssize_t nbytes = read(fd2, buf, 139);
}

void thread2(int fd) {
    int fd2 = open("myfile.txt", O_RDONLY, 0);

    char buf[4096];
    ssize_t nbytes = read(fd2, buf, 736);
}
```

What is the offset for T2's `fd2` when both threads have finished executing? (Assume that all files accessed are sufficiently large, that an EOF is not returned and that the read operation successfully reads the number of bytes requested. If there is not enough information to determine the offset enter -1.)

- in the first thread
 - it created a new FD called `fd2`
 - this is the same as the original `fd` so we are pointing to `myfile.txt`
 - then read 139 bytes from `fd2` (because they are the same)
 - in the second thread
 - we create a new `fd2` by calling `open`
 - this `fd2` is completely separated from `fd` because we called `open` again
 - also threads have separate namespaces so this `fd2` has nothing to do with the one in thread 1
 - so then we read 736 bytes
 - so for T2, the offset of `fd2` is 736
- variation 2

I21.3. FD Behavior with threads

A process P opens the file "myfile.txt" and assigns the returned fd to the global variable `fd1`. P then spawns two threads, T1 and T2. When spawning the threads P passes in `fd1` as the argument to each of the threads. After spawning the threads, P immediately waits for both of them to complete.

The threads T1 and T2 execute the following functions, respectively.

```
void thread1(int fd) {
    int fd2 = fd;

    char buf[4096];
    ssize_t nbytes = read(fd2, buf, 370);
}

void thread2(int fd) {
    int fd2 = 6;
    assert(dup2(fd, fd2) == fd2);

    char buf[4096];
    ssize_t nbytes = read(fd2, buf, 550);
}
```

What is the offset for T2's `fd2` when both threads have finished executing? (Assume that all files accessed are sufficiently large, that an EOF is not returned and that the read operation successfully reads the number of bytes requested. If there is not enough information to determine the offset enter -1.)

- in this case, T2 creates a new `fd2 = 6` but since we call `dup2`, `fd2` points to the same object as `fd` (and `fd` is shared between the threads)
 - point: threads share the same fd and same offset
- here, we don't know the ordering, but it doesn't matter since they're both progressing by offset by some amount and by the end the total amount is going to be `370 + 550 = 920`

In Class 22: File Representation

- I22.1: Computing Rotational Latency and Bandwidth

I22.1. Computing Rotational Latency and Bandwidth

A disk rotates at a speed of 7200 RPM. If a track holds 16 sectors of 4096 bytes each, how long (in ms) will it take to transfer a full track? Round your answer to two decimal digits.

Full track transfer = number (2 digits after decimal) ms ?

What is the maximum bandwidth (in MB/sec) you can get from this device? (Assume you can transfer a full track on every revolution with no seeks.)

Max bandwidth = number (2 digits after decimal) MB/s ?

Save & Grade Single attempt

Save only

Additional attempts available with new variants ?

◦ part 1

- this is basically asking for the time it takes for a full rotation because that's the same as the time needed to read a full track if there are no seeks

$$\begin{aligned}\text{Time for 1 revolution} &= \frac{60 \text{ seconds}}{7200 \text{ rotations}} = \frac{1 \text{ s}}{120 \text{ rotation}} \\ &= \frac{1}{120} \times 1000 \text{ ms} \\ &= 8.33 \text{ ms}\end{aligned}$$

◦ part 2

- a track has 16 sectors which has 4096 bytes, we want to first find the total data per track

$$\text{Data per track} = 16 \text{ sectors} \times \frac{4096 \text{ bytes}}{\text{sector}} = 65536 \text{ bytes} \times \frac{1 \text{ MB}}{2^{20} \text{ bytes}} \approx 0.0625$$

- bandwidth is the amount of data transferred over time and we know that it takes 8.33 ms to get the entire track

$$\text{Bandwidth} = \frac{0.0625 \text{ MB}}{8.33 \times 10^{-3} \text{ seconds}} = 7.5 \text{ MB/s}$$

(another way to think of this is we currently have 0.0625 MB/revolution, can just multiply that by rotations per seconds)

$$\frac{0.0625 \text{ MB}}{\text{revolution}} \times \frac{7200 \text{ revolution}}{60 \text{ seconds}} = 7.5 \text{ MB/s}$$

- I22.2:

I22.2. Computing File Read Performance

A disk rotates at a speed of 7200 RPM and has an average seek time of 6.9 ms. A track holds 128 sectors of 4096 bytes each.

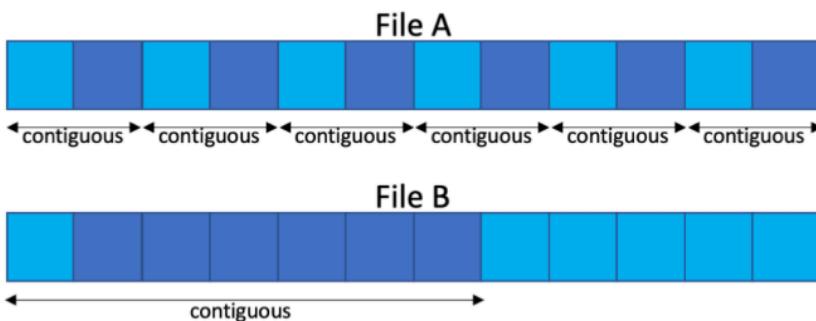
Call a block contiguous if it is allocated immediately after its predecessor on disk. (Note, by this definition, the first block of a file cannot be considered contiguous, nor can any block to which you must seek before accessing it.)

Consider a file of 80 4096-byte blocks, of which 33 are contiguous and the rest are randomly scattered across the disk.

Assume that when blocks are allocated contiguously, they are on the same track.

If we consider the time it takes to read this file sequentially, does it matter if the 33 contiguous blocks appear all together in the file or if they are each contiguous with only one block? That is, will your sequential read performance be different for file A and file B?

Here is an illustration of the conceptual layout of the two files, although the specific distribution of blocks may not exactly match this nor is the specific number of blocks illustrated necessarily 80.



No ✓

Yes

✓ 100%

How long (in ms) will it take to read the entire file sequentially? Round your answer to 2 decimal places.

Sequential read time= 3 ms ? x 0%

How long will it take if we read the file randomly? Round your answer to 2 decimal digits.

Random read time= 2 ms ? x 0%

- part 1
 - the answer is No
 - key thing to consider is how many seeks are needed
 - in the first case, we need 6 seeks (one before each pair of contiguously allocated blocks)
 - in the second, we also need 6 seeks (one at the beginning and then one before each of the non-contiguous blocks (the light blue))
- part 2
 - key point

- for contiguous blocks, you don't need to seek again as they are on the same track
 - so you only pay the transfer time for each of the sector
- non-contiguous blocks, you'll have to seek to the new track
 - AND once you're on the right track, you'll have to spin the disk some time to find the right sector
→ on average it will take half a rotation to get to the right place
 - so it's Transfer Time Per Block + Rotational Latency
- transfer time per block
 - disk spins at 7200 RPM or 120 rotations per second → 8.333 ms per rotation (or per track)
 - there are 128 sector on a track so time to get 1 sector is $8.333/128 = 0.065$ ms (since 1 block is 1 sector)
- rotational latency
 - we said it was half the time for 1 full rotation
 - $8.333 \text{ ms}/2 = 4.166 \text{ ms}$
- formula is something like

$$\begin{aligned}\text{Sequential Reading} &= (\text{transfer time for all block}) + (\text{seek} + \text{rotational delay for non-cont blocks}) \\ &= (80 \times 0.065) + ((80 - 33) \times (6.9 + 4.166)) \\ &= 525.34 \text{ ms}\end{aligned}$$

- the first bracket it's 80 because all of them have to pay it anyways
- (alternatively you could have done 33 in the first bracket, 47 in the second bracket AND add transfer time)
- part 3:
 - key: since it's random, we have to pay the seek and rotational latency every time

$$80 \times (6.9 + 4.166 + 0.065) = 890.54 \text{ ms}$$

- I22.3. Layout Score

I22.3. Layout Score

Given how disks operate, it should be clear that if we are allocating individual blocks (on the order of 4096 byte) to a file, we'll get the best sequential read performance from a disk if we can allocate files contiguously on disk. One way of expressing how well we do this is via something called *layout score*.

We define the layout score as the fraction of blocks that could be allocated contiguously that are allocated contiguously. Since the first block of a file cannot be allocated contiguously (i.e., it has no predecessor), the layout score for a 1-block file is undefined, but for all files larger than a single block, the layout score is $\text{NumberOfContiguousBlocks}/(\text{TotalBlocks} - 1)$.

A file system that allocates files in a single extent has layout score of 1; a file system that randomly throws blocks all over the disk has a layout score of approximately 0. We can use the layout score to compute a crude approximation of performance by assuming that contiguous blocks are accessed at the transfer rate of the drive and all other blocks require a seek plus one-half rotation plus the transfer time.

However, layout score is not a perfect representation of the quality of the allocation of blocks on disk; which of the following might cause us to achieve **better** performance than the layout score suggests?

However, layout score is not a perfect representation of the quality of the allocation of blocks on disk; which of the following might cause us to achieve **better** performance than the layout score suggests?

- Two blocks might be on the same track, even if they are not contiguous, so there would be no seek between them. ✓
- Two blocks might be on the same platter, so even if they are not contiguous, there would be no seek between them.
- Not all seeks take the same amount of time, so if blocks are allocated on nearby tracks, the seektime for them would be smaller than the average seek time.
- Not all seeks take the same amount of time, so if blocks are allocated particularly far away from each other, the seektime between them would be greater than the average seek time.
- Two blocks might be on the same cylinder, even if they are not contiguous, so there would be no seek between them. ✓
- Two disk blocks might appear contiguous but one could be the last block on cylinder i and the next one could be the first block on cylinder $i + 1$

Select all possible options that apply. ?

66%

a thing to keep in mind here is that there are 3 types of answers: something that's true but not relevant to our question, something that's true and relevant to our question, something that's false

1. True and relevant. We get better performance this way because we don't incur seek time
 2. False. If they're on the same platter we still have to seek, it's when they're on the same cylinder that we don't have to seek
 3. True and relevant. We would get better performance this way
 4. True but irrelevant. This would actually get us worst performance - which is not what the question is asking for
 5. True and relevant. We would get better performance because we don't have to seek
 6. True but irrelevant. This would get us worst performance because we have to seek
- I22.4. Computing File Read time using Layout Score

I22.4. Computing File Read time using Layout Score

Let's use layout score to estimate performance. The number you get should be an estimate, because the layout score suggests a fraction of the blocks that are contiguous, and you may find that that doesn't produce an integral number of blocks. For the purposes of a performance estimate, that's a fine approximation.

A disk rotates at a speed of 10000 RPM and has an average seek time of 5.8 ms. A track holds 16 sectors of 4096 bytes each.

Assuming the following:

- Contiguous blocks are allocated on the same track.
- You should ignore the time to read the first block.

Given a file of 18 4096-byte blocks with a layout score of 0.34, how long (in ms) will it take to read the entire file sequentially?

- first let's figure out read time

$$\text{read time for 1 track} = \frac{60 \text{ seconds}}{10,000 \text{ rotation}} \times 10^3 \text{ ms} = 6 \text{ ms/rotation (or track)}$$

$$\text{read time for 1 sector} = \frac{6 \text{ ms /track}}{16 \text{ sector/track}} = 0.375 \text{ ms / sector}$$

- we also know that the rotational latency is half the time of a rotation

$$\text{rotational latency} = 6/2 = 3 \text{ ms}$$

- then we can find out number of contiguous blocks and number of non-contiguous blocks

$$\# \text{ of contiguous block} = (18 - 1) \times 0.34 = 5.78 \text{ blocks}$$

$$\# \text{ of non-contiguous block} = 17 - 5.78 = 11.22 \text{ blocks}$$

- for every contiguous block, we just have to pay the read time; but for every non contiguous block, we have to pay the seek time + rotational delay + read time

$$\text{total} = 5.78(0.375) + 11.22(5.8 + 3 + 0.375) \approx 105.11 \text{ ms}$$

another way to do it straight up is by doing this formula

$$\begin{aligned}\text{total} &= (\text{number of blocks} - 1) \times (\text{transfer time} + (1 - \text{layout score})(\text{seek time} + \text{rotational latency})) \\ &= (18 - 1) \times (0.375 + (1 - 0.34)(5.8 + 3)) \\ &\approx 105.11 \text{ ms}\end{aligned}$$

so basically it's saying every block pays the read time, and the non contiguous block (the $1 - 0.34$) have to pay extra rotational delay and seek time

this works because the layout score is quite literally just a percentage of blocks that are contiguous

- I22.5. Comparing single-extent allocation to block-based allocation

I22.5. Comparing single-extent allocation to block-based allocation

Let's compare the performance of a file allocated as a single extent to one allocated in blocks.

A disk with 4 platters rotates at a speed of 10000 RPM and has an average seek time of 4.2 ms. A track holds 256 sectors of 4096 bytes each.

Given a file of 3840 4096-byte blocks, how much faster (in terms of Speedup) can we read the file if we allocate a single extent compared to reading each block assuming blocks are allocated randomly across the disk?

Assume the following:

- Extents begin on cylinder boundaries.
- In extent-based allocation, we allocate all the blocks on a single track before moving to another track.
- The disk has a track buffer that allows it to read an entire track, starting at any sector on the track.
- It takes the average seek time to move between two adjacent cylinders (this is not a realistic assumption).

Round your answer to 1 decimal digit.

- again from the RPM we know that read time is 6 ms per track and 3 ms of rotational array
- here, the time to read 1 sector is $6/256 = 0.0234375 \text{ ms / sector}$
- calculate the old time (random access)

- since it's random access, we pay a (seek + rotational delay + read time)

$$T_{\text{old}} = 3840(4.2 + 3 + 0.0234375) = 27738 \text{ ms}$$

- calculate the new time (extent based)

- the trick here is that we have to pay as many seeks as the number of cylinders that the file takes up
 - (but we don't have to pay any rotational delay or anything like that - just read time per sector)
- sectors per cylinder: $256 \times 2 \times 4 = 2048$ sectors / cylinder
 - (number of sector on a track, by the number of tracks on a surface, by the number of surface in a disk)
- so our files take up $3840/2048 \approx 2$ cylinders since we have to round up
 - i.e take the ceiling
- putting it together

$$\begin{aligned} T_{\text{new}} &= \text{total seek time} + \text{total read time} \\ &= (2 \times 4.2) + (3840 \times 0.0234375) = 98.4 \end{aligned}$$

- note: don't have to pay rotational delay because the extents are aligned to the cylinder boundaries

- calculating speed up

$$\text{speedup} = \frac{T_{\text{old}}}{T_{\text{new}}} = \frac{27738}{98.4} = 281.89$$

Practice Quiz

R4.1: Impact of Stride on Throughput

- question

R4.1. Impact of Stride on Throughput

Consider a system that has **no caches but does have DRAM** (that is, main memory) and a byte array that fits entirely in DRAM.

Assume that we have an array that is already present in the DRAM.

Which of the following statements best describes the average memory access latency of the system when the elements of that byte array are accessed in order of address, but with different strides?

- As the stride increases, the average latency will remain the same. ✓
- As the stride increases, the average latency will first decrease and eventually reach a point where the average latency remains constant.
- As the stride decreases, the average latency will first decrease and then increase before remaining constant.
- As the stride decreases, the average latency will first increase and then decrease before remaining constant.
- As the stride decreases, the average latency will first increase and eventually reach a point where the average latency remains constant.

Select all possible options that apply. ?

✓ 100%

Try a new variant

- since there's literally no cache, every access will go to memory → hence the same speed

R4.2: Abstract Stride Across Eviction Policies

- question

R4.2. Abstract Stride Across Eviction Policies

Let's say that consecutive letters in our letter-based cache traces indicate consecutive cache lines in memory. (So, B is the line in memory immediately after A, and C is the line immediately after B.)

Consider the following access sequence representing 3 traversals at a fixed stride through an array:

A,A,B,B,C,D,D,E,E,A,A,B,B,C,D,D,E,E,A,A,B,B,C,D,D,E,E

- part 1

What can you say about the stride in this access pattern, in comparison to the line size?

- (a) It is more than one line size.
- (b) It is exactly one line size.
- (c) It (strictly) between one half line size and one line size. ✓
- (d) It exactly one half line size.
- (e) It is less than one half line size.
- (f) There is not enough information to tell.

✓ 100%

- if the stride was exactly half the line size, everything will appear twice
- if the stride was exactly the line size, everything will appear once
- since most things occur twice but C only appears once, this is the answer

- analogy
 1. If your step length is exactly half the size of a sidewalk block, you'll step twice in each block, making each block appear twice during your walk.
 2. If your step length matches the size of a sidewalk block, you'll step once in each block, and each block will only appear once during your walk.
 3. However, if your step length is somewhere between these two, it's like walking with a stride that's a bit longer than half a sidewalk block. Sometimes, you'll take one step inside a block, and sometimes you'll take two steps before moving to the next block. This means that while most blocks are visited twice, some blocks get skipped on the second step.
- part 2

If the cache can hold exactly four lines, which produces the most hits?

(a) A direct-mapped cache. ✓
 (b) A fully-associative LRU cache.
 (c) There is not enough information to tell.
 (d) It's a tie.

✓100%

- if it only holds 4 lines → implies that A and E will have the same index
 - so if we do simple replacement, when E comes, we'll evict A, then A comes again and we evict E, but B,C,D stays
- if we do LRU
 - when E comes we evict A
 - when A comes we evict B
 - B comes we evict C
 - and we keep doing that so we don't get any more hits
- part 3

If the cache can hold exactly eight lines, which produces the most hits?

(a) It's a tie. ✓
 (b) There is not enough information to tell.
 (c) A direct-mapped cache.
 (d) A fully-associative LRU cache.

✓100%

- the cache is big enough where we don't have to evict anything, we just take compulsory misses

R4.3: Comparing State transitions between MSI and MESI

- variation 1

R4.3. Comparing State transitions between MSI and MESI

Consider the MSI and MESI protocols **as presented in class**. We want to know which action(s) could produce **Core 1** transitioning to a different state using MSI than it would using MESI, assuming that the start state for both protocols is the same.

- Core 1 writes a cache line and gets a hit
- Core 2 writes a cache line
- Core 1 reads a cache line, but it is a miss.
- Core 1 writes a cache line, but it is a miss.

Select all possible options that apply.

100%

- the answer is referencing starting in Invalid
 - in MESI, we could potentially go to Exclusive if the data has never been read before
 - in MSI, you could only go to Shared
- another possible answer is Core 2 reads a cache line
 - if Core 2 reads a cache line held in Core 1's cache line that is exclusive, it will transition Core 1's cache line from exclusive to shared

R4.4: Balancing write through and write back performance

- question

R4.4. Balancing write through and write back performance

A processor has a single level cache with the following characteristics:

- A cache hit takes 5 cycles.
- A cache miss takes 14 cycles for reads.
- A cache miss also takes 14 cycles for writes if the cache line being replaced is clean.
- A write directly to DRAM takes 7 cycles.
- Data can be written to the cache and DRAM in parallel.

Assuming that all evicted cache lines are clean, what write hit rate will make the average write time be the same for both write-through (write no-allocate) and write-back (write allocate) caches?

Express your answer as a decimal number between 0 and 1 inclusive, rounded to two decimal places. If there is no hit rate that will allow the average write time to be the same for the two implementations, enter -1.

- let x be the hit rate
- average access time for write back (write allocate) cache

$$x(5) + (1 - x)(14)$$

- first thing is the time it takes for a cache hit
- 2nd thing is the time it takes for a cache miss (since all lines are clean, we don't have to actually do any write back, just write from memory to cache)
- average access time for write through
 - it's just 7 because we can write to cache and memory in parallel

- if it's a hit - we write to both the cache and memory (in parallel - also makes sense that writing to the cache takes less time than writing to memory so we just use 7)
- if it's a miss- we write straight to memory

- formula

$$\begin{aligned}
 5x + (1-x)(14) &= 7 \\
 5x + 14 - 14x &= 7 \\
 -9x &= -7 \\
 x &= 7/9 \approx 0.78
 \end{aligned}$$

R4.5: File systems - I/O Redirection Implementation

- question

R4.5. File systems - I/O Redirection Implementation

In Unix, the shell reads the command line and makes a determination as to whether or not I/O redirection is required for the program it is going to run. One of the possible redirection scenarios is described below.

- ">>&" - append the output of standard error to a file (Note this redirection behaviour is for question purposes only and is not the actual behaviour when used on the command line.)
example: *command >>& fileName*

Semantics:

- Open *fileName*, and if it doesn't exist create it
- Cause the standard error of **command** to be appended to the end of the file with the name *fileName*

You will produce two lines of code below by selecting the right function call and parameters to produce the C code the shell would use to manipulate the file descriptors to perform the I/O redirection for the scenario described above.

Caution the drop down is long so make sure to scroll to see all the options. If a parameter field is **not present, unused/unneeded, or ignored** select **N/A**.

If needed you can assume that the drop down options of *command* and *fileName* are the pointers to the strings that represent the command and file name respectively.

- part 1

First line of code

```
newFD = open ( fileName , O_WRONLY|O_CREAT|O_APPEND , 0644 );
```

- first to create a NEW fd - you need to **create**
- basically you want to open the new file that you will write to
- **O_WRONLY|O_CREATE** means to set it to write only (only valid choice, but could have been write and read too) and **O_CREATE** means to create a new file if that one doesn't exist
- **O_APPEND** just means we write to the very end of the file regardless of where the offset is (not quite sure why we need it - just to be safe I guess)

- part 2

Second line of code

```
dup2    ✓ 100% ( newFD    ✓ 100% ,
STDERR_FILENO    ✓ 100% );
```

- we want STDERR FD to now point to our new file descriptor (mind the order)

R4.6: Relative Disk Drive Performance

- question

R4.6. Relative Disk Drive Performance

Consider a disk drive with a seek time of 7 ms and rotational latency of 16 ms.

A second disk drive has exactly the same configuration and performance characteristics **except that it spins more slowly.**

How will the performance characteristics, listed below, of this second drive compare to those of the first drive?

The throughput of the second drive relative to the first drive will:

be lower ✓ 100%

The seek time of the 2nd drive relative to the 1st drive will:

be the same ✓ 100%

The sector transfer time of the 2nd drive relative to the first drive will:

be larger ✓ 100%

[Try a new variant](#)

- key:

- If the disk spins more slowly, than the RPM for the disk decreases. That means that one rotation takes longer for the second drive.
- If a rotation takes longer, then it also takes longer for sector to pass under the head (i.e., the transfer time will increase) and it takes longer for the sector to move under the head once the head is positioned on the proper track. This extra transfer time and rotational latency will cause fewer bytes to be transferred in a given amount of time so the throughput will decrease. The seek time is how long it takes to move the heads to a track and that is not a function of how fast the disk is spinning.

- part 1

- it spins slower so it takes longer to read the data from the tracks
- it takes longer to transfer the same amount of data → so throughput is lower

- part 2

- seek time has nothing to do with rotational speed
- the arm goes in and out and places itself onto one of the tracks

- so seek time is exactly the same
- part 3:
 - same as part 1, since it spins slower it takes more time to transfer data

R4.7. Fun with File Descriptors

- question

R4.7. Fun with File Descriptors

Suppose a file contains the following sequence of 10 characters starting at offset 1845.

V,T,W,L,S,K,M,J,G,E

Each of the following questions describes actions taken by processes. Assume that all the actions occur in order. That is, when answering question 3, assume that questions 1 and 2 have already occurred. If it is not possible to determine the answer, write Z. If an answer asks for a letter, write it in capital letters.

Two processes A and B start running on the same machine. They both open the file and seek to location 1845.

- part 1

1. Process A reads 1 byte. What character is returned? V ? ✓ 100%

- the offset is at 0, read 1 byte so you get V

- part 2

2. Process A now sleeps for 10 seconds. While A is sleeping process B reads 1 byte. What character is returned by this read? V ? ✓ 100%

- the offset for B is still at 0, so it also reads V

- part 3

3. While A is still sleeping, B next writes the 2 characters P,R to the file. A now wakes up and reads one byte. What character is returned? P ? ✓ 100%

- note that `write` will overwrite what was already there (it will not add)
- so the array after writing will be V,P,W,L,S,K,...
- since A's offset is still at 1, it'll get P

- part 4

4. Process A again goes to sleep and while asleep process B reads 1 byte. What character is returned by this read? T ? ✗ 0%

- after writing the 2 bytes, B's offset was at 3
 - recall array now looks like V,P,W,L,S,K,...
- so now if it reads, it gets L

Corrections

Question 1: MESI Pairings

- part 1

Question 1: MESI Pairings

Imagine two caches C_1 and C_2 at the same level synchronize using MESI. Both have the same configuration (direct-mapped, 16-bit address size, overall cache size, number of sets, line size) and start empty.

With respect to the cache line containing address **0x3487**, which of these pairs of states is possible?

(a) $C_1: S, C_2: S$. ✓
 (b) $C_1: E, C_2: M$.
 (c) $C_1: M, C_2: S$.
 (d) $C_1: I, C_2: I$. ✓
 (e) $C_1: I, C_2: E$. ✓
 (f) $C_1: E, C_2: E$.

Select all possible options that apply. [?](#)

✓ 100%

- since this is both respect to the address itself, we know that both these cache slots will be touching the cache line at **0x3487**
- 1. it's possible for **(S, S)** to happen when the slots are sharing some data
- 2. not possible because C2 also claims to have the data while C1 claims to exclusively have the data
- 3. not possible because C1 has a modified version of the data, while C2 claims to have a shared (if that was a the case, C1 state should be shared too - means the modification has been written back)
- 4. possible because this is how they start
- 5. possible if C2 reads the line exclusively and C1 still hasn't read anything
- 6. not possible because they both claim to have the data exclusively

- part 2

With respect to the metadata for the cache slot in set #10, which of these pairs of states is possible?

(a) $C_1: I, C_2: S$.
 (b) $C_1: I, C_2: E$. ✓
 (c) $C_1: S, C_2: S$. ✓
 (d) $C_1: E, C_2: E$.
 (e) $C_1: E, C_2: S$.
 (f) $C_1: E, C_2: M$.

Select all possible options that apply. [?](#)

○ 33%

- this is basically just saying with respect to the slots in the cache - if we don't know that they both map to the same address, they might just contain different lines
 - in which case, we can't say anything about the combination of states, because by themselves, the individual states are all valid

- so the answer is actually all

Question 2: Computing cache writeback times

- question

Question 2: Computing cache writeback times

Consider a system with a writeback (write allocate) cache. You are told that an application has a cache write miss rate of 35%. 33% of the time that we take a miss, the cache line being replaced is dirty.

If it takes 18 ns to read/write the cache and 144 ns to transfer a cache line between the cache and the data source, what will be the average write time for this application?

Assumptions: on a miss, if the evicted cacheline is clean, the total miss time should include both the time to read data from the data source and the time to update the cache.

Round your answer to two decimal places.

85.03

ns



✓ 100%

- a hit will take 18 ns
- a miss + clean eviction will take 144 (to load from memory) + 18 (to read from cache) = 162 ns
- a miss + dirty eviction will take 144 (to write dirty data) + 144 (load data from memory) + 18 (to read from cache) = 306 ns
- putting that all together

$$\text{average write time} = 0.65(18) + (0.35)(0.67)(162) + (0.35)(0.33)(306) \\ \approx 85.02 \text{ ns}$$

Question 3: Writeback cache

- question

Question 3: Writeback cache

Given a direct-mapped cache with 4 2-byte cache lines and a main memory that has 8-bit addresses.

A load instruction is trying to load from address 0xfe. If there is a cache hit, determine what data would be retrieved from the cache (be sure to add the 0x prefix). If there is a cache miss, determine whether a write-back would happen. If there would be a writeback type "writeback". If there would NOT be a writeback type "replace".

Index	Valid	Dirty	Tag	Block offset(0)	Block offset(1)
00	0	0	0 1001	0x40	0x66
01	1	0	0 0110	0x56	0x20
10	1	1	1 0011	0x64	0x15
11	1	0	0 1001	0x10	0x62

Enter all hex digits in your answer. Enter 0x before the number, except for misses (type "writeback" or "replace").

Answer = replace



✓ 100%

This question is complete and cannot be answered again.

- 0xfe = 0b11111110

- offset: 0
- index: 11
- tag: 11111
- the tag at index 11 is 01001 so it's a miss
- since the valid bit is on and the dirty bit is off, we don't have to write back, we just have to replace (i.e simply just evict the line)

Question 4: Applying Amdahl's Law

- question

Question 4: Applying Amdahl's Law

Assume that a program spends its time only in ALU operations and Memory operations.

You are told that an application spends 24% of its time in the **memory system** and takes 670 seconds to complete.

If we want our program to complete in 394 seconds, how many times faster do we need to make the **ALU**?

2.18	times faster	?	✓ 100%
------	--------------	---	--------

This question is complete and cannot be answered again.

- we have the formula for speedup

$$\begin{aligned} \text{speedup} &= \frac{T_{\text{old}}}{T_{\text{new}}} \\ &= \frac{1}{1 - \alpha + \frac{\alpha}{k}} \\ \frac{670}{394} &= \frac{1}{1 - 0.24 + \frac{0.24}{k}} \\ k &\approx 2.18 \end{aligned}$$

Question 5: Stride Exercises for Oversizing Arrays

- question

Question 5: Stride Exercises for Oversizing Arrays

Imagine we have a cache with the following parameters:

- Total size of 8 KB (2^{13} bytes)
- 16 sets
- 16-way set associativity
- 256 total lines
- 32 byte line size
- LRU eviction policy

To study our cache, we traverse a 8 KB array of `uint8_t` called A exactly 10000 times. We start each traverse at index 0 and then walk through the array with a stride size of 16 bytes. Assume that A is 8 KB-aligned (i.e., perfectly aligned with respect to the cache) and that we make no memory accesses except to A .

(Compute carefully on this problem. No partial credit is available within each part.)

- part 1

If A is exactly 8 KB in size, exactly how many misses will occur in our 10000 traversals?

misses: 256



✓ 100%

- we take on 256 compulsory miss (1 for every line we access)
- since the cache fully fits in the cache, for the next 9999 times, we get all hits
- thus only 256 misses

- part 2

If A is exactly 32 bytes (1 cache line) larger than 8 KB in size, what will be the set number (the index) of the last line in A (the new line)? (This question and the next are each worth half as much as each other question is worth, but they should help you with the subsequent question. Beware of "off by one" errors on these two! A picture may help!)

257000



✗ 0%

- note that this question is asking about the set number index
- it wraps around and we get set number 0

- part 3

If A is exactly $7 * 32$ bytes (7 cache lines) larger than 8 KB in size, what will be the set number (the index) of the last line in A ?

7



✗ 0%

- same as above, this is set number 6 (off by 1 here)

- part 4

Now, let A be exactly 416 bytes (13 cache lines) larger than 8 KB and repeat the original experiment. How many misses will occur in our 10000 traversals? (Don't forget that we are using LRU, what you learned on the previous two parts, and that some sets now have an extra line in them.)

misses: $(256 + (416/32)) * 10000$

?

x 0%

- very hard question
- as above, for the extra 13 cache lines, they will wrap around to the first 13 sets
 - we will call them "affected sets"
- for these affected sets, due to our access pattern, we will miss on every line
 - also due to our eviction policy, we'll always evict the next one we need
 - the hard part is that you're taking a miss on $16 + 1$ lines (because of the extra line)
- for the non-affected sets, they remain in the cache, so all subsequent accesses you get hits
 - so you take 16 compulsory miss per set
- the math

$$\text{affected sets} = 10,000 \times (13 \times 17)$$

$$\text{unaffected sets} = (16 - 13) \times 16$$

$$\text{total} = \text{affected sets} + \text{unaffected sets}$$

- part 5

If we halve the cache's associativity by doubling the number of sets and then repeat the previous experiment, what impact would that have on the number of misses?

- (a) Increases
- (b) Stays the same **x**
- (c) Decreases
- (d) There is not enough information to tell

x 0%

- if you increase the number of sets, the number of lines missed per sets shrinks
 - so the $16 + 1$ becomes like $8 + 1$
- and so the number of misses **decreases**

Question 6: Applying Amdahl's Law

- question

Question 6: Applying Amdahl's Law

Patrice and Margo are building a new version of the PAM processor. They are adding a new cache that makes memory instructions 2.4 times faster. Additionally, they are using logic that makes all non-memory instructions 2.3 times faster.

- part 1

Given an application that spends 18% of its time executing memory instructions, what speedup will the new PAM processor produce?

Round your answer to 2 decimal digits.

speedup = 2.32



✓ 100%

- this is a slightly more complicated version of Amdahl's Law
 - also note that there are no portion of the program that's not sped up (that used to be the $1 - \alpha$ part - that's now 0)
- we can do

$$\begin{aligned}\text{speedup} &= \frac{1}{0 + \frac{\alpha_1}{k_1} + \frac{\alpha_2}{k_2}} \\ &= \frac{1}{\frac{0.18}{2.4} + \frac{0.82}{2.3}} \\ &\approx 2.32\end{aligned}$$

- part 2

If the program ran in 341 seconds on the old version of the PAM processor, how long will it take to run on the new processor?

Round your answer to 2 decimal digits. (Assume the program exactly follows the time profile above.)

New Time = 147.15

seconds



✓ 100%

This question is complete and cannot be answered again.

- we also know

$$\begin{aligned}\text{speedup} &= \frac{T_{\text{old}}}{T_{\text{new}}} \\ T_{\text{new}} &= T_{\text{old}}/\text{speedup} \\ &= 341/2.32 \\ &\approx 147.15\end{aligned}$$

Question 7: A "Pipe" of Two File Descriptors

- question

Question 7: A "Pipe" of Two File Descriptors

A "pipe" is a special structure the operating system can create. It comes with two file descriptors and begins empty. The first file descriptor is write-only and adds bytes to the pipe in the order given. The second is read-only and produces bytes from the pipe in the order given. (So, the pipe is like a real-life "pipe" where the bytes are water flowing into the first file descriptor, through the pipe, and out the second file descriptor.)

Assume that a process has just made a pipe with its first file descriptor (the write-only one) being 3 and its second (the read-only one) being 5 and is otherwise fresh (a plain process with nothing special going on).

Based on what you know about the file system, which of these are true? (Judge each one separately, i.e., ignore the effects they could have on each other.)

- (a) It would make sense to use `dup2(3, 1)` to redirect standard I/O for inserting bytes into the pipe.
- (b) Initially creating a pipe can be done without a system call.
- (c) The process can pass 3 as the first argument to `write`. ✓
- (d) It would make sense to use `dup2(5, 2)` to redirect standard I/O for extracting bytes from the pipe.

If the process forks, the resulting processes will both have access to the pipe using the file descriptors 3 and 5. ✓

- (f) If the process executes `int fd = dup(3); close(3);`, it will lose access to write to the pipe.

Select all possible options that apply. There are exactly 3 correct options in the list above. ?

○ 66%

This question is complete and cannot be answered again.

- so basically, what 5 would look like is a `command.txt` file with a bunch of linux commands, and then when you PIPE the command into the terminal, it reads the commands from the text file instead of the stuff from your terminal
 - conversely, 3 would be a blank file at first and the shell would write to this file instead of writing to the terminal
1. 1 is stdout, this is pointing stdout to the right end of the pipe (fd 3)
 - (inserting into the pipe means to write to the file)
 2. False. This is just not possible
 3. True. If we want to re-direct the output of the shell to our output file - we can write to 3
 4. False, `2` is actually STDERR so not really what we want to do
 5. True. Child processes share fd with its parent
 6. False. Since `fd` now points to the same thing as `3`, when you close 3, you can still reference that file with `fd`

Question 8: FD Behavior with threads

- question

Question 8: FD Behavior with threads

A process P opens the file "myfile.txt" and assigns the returned fd to the global variable `fd1`. P then spawns two threads, T1 and T2. When spawning the threads P passes in `fd1` as the argument to each of the threads. After spawning the threads, P immediately waits for both of them to complete.

The threads T1 and T2 execute the following functions, respectively.

```
void thread1(int fd) {
    int fd2 = fd;

    char buf[4096];
    ssize_t nbytes = read(fd2, buf, 368);
}

void thread2(int fd) {
    int fd2 = fd;

    char buf[4096];
    ssize_t nbytes = read(fd2, buf, 809);
}
```

What is the offset for T2's `fd2` when both threads have finished executing? (Assume that all files accessed are sufficiently large, that an EOF is not returned and that the read operation successfully reads the number of bytes requested. If there is not enough information to determine the offset enter -1.)

1177



✓ 100%

This question is complete and cannot be answered again.

- note that `fd2` in both threads uses `fd`
 - and `fd` was shared from its parents
 - so `fd2` in both threads are pointing to the same thing AND they are sharing the same offset
- we don't really know which one will run or finish first but we know that at the end, it would have moved $368 + 809 = 1177$ bytes

Question 9: Disk Fundamentals

- question

Question 9: Disk Fundamentals

Let's say that we have a disk with 2 platters. Each surface has 1024 tracks. Each track contains 16 sectors.

How many sectors are there on a platter?

32768



✓ 100%

Assuming that sectors are numbered starting at 0, is a seek required between accessing sector 46700 and 16394?

no

yes ✓

✓ 100%

- part 1
 - note that there are 2 face on a platter
 - on one face we have 1024 tracks $\times \frac{16 \text{ sectors}}{\text{tracks}} = 16384 \text{ sectors}$
 - so one 2 face we have $16384 \times 2 = 32768$
- part 2
 - there are no seeks required if the 2 is located on the same cylinder
 - so we have
 - 1 track have 16 sectors
 - on both face of the platter that's 32 sector
 - there are 2 platter so that's 64 sector per cylinder
 - $46700 // 64 = 729$ and $16394 // 64 = 256$ so they are not on the same cylinder

Question 10: Relative Disk Drive Performance

- question

Question 10: Relative Disk Drive Performance

Consider a disk drive with a seek time of 2 ms and rotational latency of 4 ms.

A second disk drive has exactly the same configuration and physical characteristics **except that it can store bits more closely together along a track (with the same number of bits per sector)**.

How will the performance characteristics, listed below, of this second drive compare to those of the first drive?

The average throughput—across many file operations of varying lengths—of the second drive relative to the first drive will: 100%

The seek time of the 2nd drive relative to the 1st drive will: 100%

The sector transfer time of the 2nd drive relative to the first drive will: 100%

This question is complete and cannot be answered again.

- part 1
 - closer bits along a track means shorter sectors, which means faster sector transfer time (you're grabbing more bits under the same distance or time)
 - so throughput would be better
- part 2
 - we didn't say anything about seek time so that'll still be the same
 - (note that if we somehow made the tracks closer together - that perhaps will decrease seek time)

- part 3
 - just like part 1
 - lower sector transfer time because you get shorter sectors so you can finish earlier

Next Quiz + Midterm 2

In Class Assignment

In Class 23: File Representation

- I23.1: Counting IO for meta data blocks in Single Extent Indexes

I23.1. Counting IO for meta data blocks in Single Extent Indexes

Consider a file system that represents files using a single extent.

Assume that it has a blocksize of 8192 and that disk addresses are 4 bytes.

If we represent files using a single extent, and we limit extents to being no larger than a maximum extent size of 597152 blocks, if neither the file's metadata structure nor any other data or metadata for this file has been read yet, how many IOs are required to read block number 72996.

integer IOs:

Additional attempts available with new variants

- for extent based, once you get the inode you will get the address of the start of the file on disk
 - and all the subsequent bytes are stored contiguously
 - once you have the start, it's pretty easy to get to byte offset , just do → so this takes 1 read only
- but you also have to do a read first to get the inode, so the answer is 2
- I23.2: Counting IO for meta data blocks in Multi-level Indexes
 - method 1: counting indexes

I23.2. Counting IO for meta data blocks in Multi-level Indexes

Consider a file system that represents files using a multi-level index. The inode references the root of the index, but the index is only as deep as is necessary to represent a file of a given size (e.g., if the file contains only a single block, then the root of the index references data blocks directly).

Assume that the file system blocksize is 8192 bytes and that disk addresses are 8 bytes.

If we have a multi-level index with a maximum of 4 levels of index blocks and a file consisting of 1048576 blocks, assuming that neither the file's metadata structure nor any other data or metadata for this file has been read yet, how many IOs are required to read block number 294393.

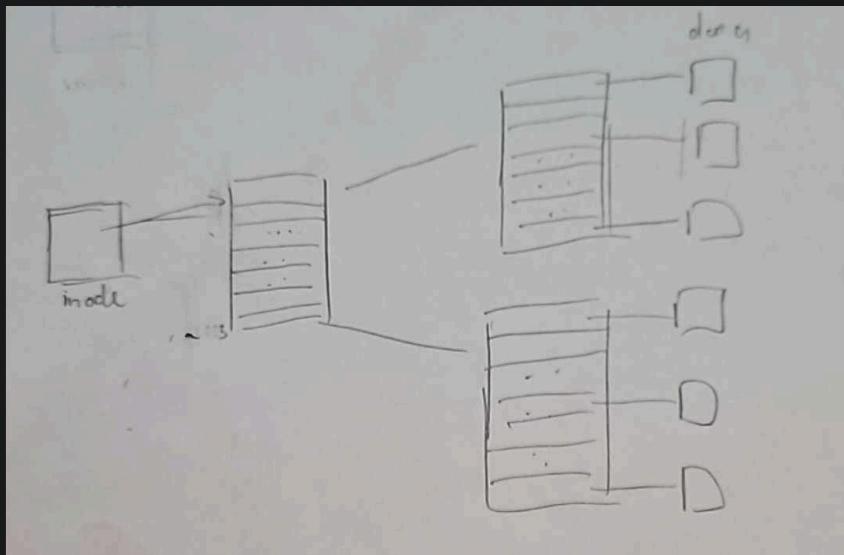
integer IOs: ?

Save & Grade *Single attempt*

Save only

Additional attempts available with new variants ?

- we know inode points to some n -depth indirect block
 - first we need to figure out how many pointers we can fit in an indirect block
 - blocksize is 8192 and disk address is 8 bytes each
 - $8192/8 = 1024$ pointers can fit on a indirect block
 - the file has 1048576, how many indirect block will we need
 - $1048576/1024 = 8$ so we'll need 8 of those pointer tables
 - this mean that our tree will be depth 2 (not counting the inode)
- picture



- we need to read the inode, read address from the double-indirect block, read the address from the direct block AND read the actual data itself
- we'll have 4 reads
- method 2: simple math

I23.2. Counting IO for meta data blocks in Multi-level Indexes

Consider a file system that represents files using a multi-level index. The inode references the root of the index, but the index is only as deep as is necessary to represent a file of a given size (e.g., if the file contains only a single block, then the root of the index references data blocks directly).

Assume that the file system blocksize is 8192 bytes and that disk addresses are 4 bytes.

If we have a multi-level index with a maximum of 3 levels of index blocks and a file consisting of 4194304 blocks, assuming that neither the file's metadata structure nor any other data or metadata for this file has been read yet, how many IOs are required to read block number 368992.

- we just want to be able to quickly determine how many LEVELS this tree has
 - instead of looking at indices, we can look at block count we have $\rightarrow 368992 + 1$ (because 0-index)
 - first - we must find how many addresses (pointers) fit into a block $\rightarrow 8192/4 = 2048$ addresses
 - if we have 1 level only (just an indirect block)
 - that indirect block points to 2048 addresses which is 2048 different data blocks
 - $2048 \ll 368,993$ so we cannot get to our desired block
 - if we have 2 levels
 - the first (double) indirect blocks point to 2048 indirect blocks which point to 2048 data blocks
 - this means we are pointing to a total of 2048^2 data blocks
 - we can see that $2048^2 = 4,194,304 \gg 368,993$ so we can definitely get to our desired block
 - thus we will have 2 levels and require 4 reads
 - (note: here we don't do any subtraction stuff - that's only for hybrid)
- **NOTE:** for these questions it's always $2 + (\text{number of levels})$ - the extra 2 is for reading the inode and then the data block itself
- I23.3: Counting IO for meta data blocks in Hybrid Indexes

I23.3. Counting IO for meta data blocks in Hybrid Indexes

Consider a file system that represents files using a hybrid index.

Assume that it has a blocksize of 1024 and that disk addresses are 4 bytes.

If we have a hybrid index with 8 direct pointers, 1 indirect pointer, 1 double indirect pointer , and 0 triple indirect pointers stored directly in the file metadata structure, assuming that neither the file's metadata structure nor any other data or metadata for this file has been read yet, how many IOs are required to read block number 43655.

integer IOs: ?

Save & Grade Single attempt

Save only

Additional attempts available with new variants ?

- note: the direct block, the indirect block stuff is all in the inode table
- we can store $1024/4 = 256$ addresses per block

- again, we can look at the number of blocks we want $\rightarrow 43655 + 1 = 43656$
 - the first 8 direct pointers won't get us our desired block $\rightarrow 43656 - 8 = 43648$
 - the indirect pointer can get us to $43648 - 256 = 43392$
 - not quite there yet
 - using our double indirect block: $43392 - (256)^2 = -22144$
 - so our block number will be in the double indirect blocks
 - this means that it'll require 4 reads
- I23.4. Manually Mapping logical block numbers in Multi-level Indexes

I23.4. Manually Mapping logical block numbers in Multi-level Indexes

Consider a file system that represents files using a multi-level index.

Assume that it has a blocksize of 16 and that disk addresses are 2 bytes.

The data below represents a collection of disk blocks. The first element in each row is a disk address (physical block number) and the remaining numbers represent the contents of the corresponding disk block.

		Disk block contents							
Disk addresses		19969	10755	48644	56582	25865	31002	59932	26400
47930		25377	30242	11049	44329	32811	48432	39474	47410
38201		38201	47930	14396	18748	36927	27455	62274	45895
11923		15178	49740	26705	25681	18518	2903	63322	27995
60667		49255	12657	49778	55154	28534	40311	60536	31352
36927		25209	56696	45944	42875	55681	58244	53758	12425
62274		5613	24970	27022	8336	44176	54161	8594	22419
5613		14396	60333	27568	20148	46516	39098	53179	10940
14396		18404	56003	54467	41412	14278	11464	54225	53203
18404		27455	64475	18404	57831	5613	48114	60667	11516
27455									27390

If the the index supports a maximum of 2 levels, given a maximum size file, what is the physical disk address of the logical block number 44 in the file with the inode containing 11923 for its index pointer?

Disk Address: 48114 ? ✓ 100%

[Try a new variant](#)

- a very slow way to do this
- 11923 is the first indirect block
 - zero-th entry in that row holds LBN range 0 - 7
 - first entry in that row holds LBN range 8 - 15
 - ...
 - fourth entry in that row holds LBN range 32 - 39

- fifth entry in that row holds LBN range 40 - 47
 - what we want
 - the address at the fifth entry is 27455
- go to row 27455
 - this row holds LBN range 40 - 47 → so LBN 44 is in index 4
 - value at index 4 is 48114
- **note: this one is assuming there is 2 levels always**

- I23.5. Manually Mapping logical block numbers in Hybrid Indexes

I23.5. Manually Mapping logical block numbers in Hybrid Indexes

Consider a file system that represents files using a hybrid index.

Assume that it has a blocksize of 16 and that disk addresses are 2 bytes.

The data below represents a collection of disk blocks. The first element in each row is a disk address (physical block number) and the remaining numbers represent the contents of the corresponding disk block.

Disk block contents									
Disk addresses	39937	58116	22791	32530	30995	56597	57367	39960	
21343	40984	41497	47131	18201	57121	36899	28196	20005	
51982	8998	62246	12336	30257	35635	27443	17206	11836	
41497	59968	37954	24387	62795	20299	32333	1618	52306	
15650	29779	31828	55384	55898	3680	2667	55663	23410	
20005	60649	48505	39803	48769	61059	3981	911	31892	12440
60649	35392	10910	24743	5032	39849	1449	7598	10679	33981
35392	47131	38590	50878	41150	13245	446	48577	59590	51911
47131	56761	25801	41418	17867	1229	6361	28379	7643	28127
56761	18557	14053	29414	2024	60649	37358	14831	39920	6653
18557									

If the inode contains 4 direct pointers, 1 indirect pointers, and 1 double indirect pointers. what is the physical disk address of the logical block number 35 in the file with the metadata containing the following pointers?

Direct pointer	15650
Direct pointer	35392
Direct pointer	56761
Direct pointer	18557
Indirect pointer	21343
Double indirect pointer	51982

- each block table contains 8 adddress
- block number 35 = 36 block count
- first 4 direct pointer → $36 - 4 = 32$
- first 1 indirect pointer → $32 - 8 = 24$

- so it has to be in the double indirect block
 - so we go to 51982
 - first column gets us 8 block $24 - 8 = 16$
 - second column gets us 8 block $16 - 8 = 8$
 - third column get us 8 block $8 - 8 = 0 \rightarrow$ so we know it's in address 47131
- in row 47131, it'll be the last column in this row so 51911
- (note: if we do the block count - when we are counting, don't do 0-index i.e 2nd column literally mean the 2nd column)
- another way to do 23.5

I23.5. Manually Mapping logical block numbers in Hybrid Indexes

Consider a file system that represents files using a hybrid index.

Assume that it has a blocksize of 16 and that disk addresses are 2 bytes.

The data below represents a collection of disk blocks. The first element in each row is a disk address (physical block number) and the remaining numbers represent the contents of the corresponding disk block.

Disk block contents									
Disk addresses	50946	47363	61445	2055	30731	28431	62480	60176	
24543	50946	47363	61445	2055	30731	28431	62480	60176	
50388	54291	8220	61982	24098	30502	3624	32812	21560	
58309	65085	19009	15938	49985	40775	51019	23376	59990	
23376	43863	61018	33627	5980	35167	40544	14946	54371	
15219	7016	36458	27498	18287	57711	63347	28535	21378	
47540	49538	13447	35977	46989	11918	55446	38039	23194	
18287	28061	3746	8354	59045	64166	61861	60072	4271	
28827	50357	22457	21439	63997	53954	27334	26056	26313	
54291	17355	9933	16077	31186	2777	59362	3815	42216	
21439	3818	63723	5869	51441	29427	23285	60919	44541	

If the inode contains 4 direct pointers, 1 indirect pointers, and 1 double indirect pointers. what is the physical disk address of the logical block number 30 in the file with the metadata containing the following pointers?

Direct pointer	58309
Direct pointer	47540
Direct pointer	24543
Direct pointer	50388
Indirect pointer	15219
Double indirect pointer	28827

Disk Address:	5869	?	✓ 100%
---------------	------	---	--------

- here, everything will be 0 indexed
- direct blocks: $30 - 4 = 26 \rightarrow$ not there yet

- indirect blocks: $26 - 8 = 18 \rightarrow$ not there yet
- double indirect blocks: $18 - 8^2 = -46$
 - it will be in the double indirect blocks
 - go to 28827
- now, find which indirect blocks to go to from the double-indirect blocks
 - $18 // 8 = 2$
 - go to the 2nd column (0 index) at row 28827
 - points to 21439, go there
- find which index points to our data block
 - $18 \% 8 = 2$
 - go to 2nd column (0 index) at row 21439
 - data is 5869

Corrections

Question 2: Disk Fundamentals

- question

Question 2: Disk Fundamentals

Let's say that we have a disk with 2 platters. Each surface has 4096 tracks. Each track contains 16 sectors.

How many sectors are there on a platter?

131072
?
✓ 100%

Assuming that sectors are numbered starting at 0, is a seek required between accessing sector 195897 and 195864?

yes
 no ✓
✓ 100%

This question is complete and cannot be answered again.

- part 1
 - there are 2 platter on the disk, each with 2 surfaces and each surface contains 4096 tracks 16 sectors
 - math: $2 \times 2 \times 4096 \times 16 = 131072$
- part
 - you need to figure out how many sectors there are per cylinder, so 1 track across all surfaces (there are 4 total surface) and you can find out how many sectors
 - $2 \times 2 \times 16 = 64$
 - you see what cylinder the sectors fall in
 - $195897 // 32 = 3060$

- $195864 \text{ } // \text{ } 32 = 3060$
- they are on the same cylinders so you don't need a seek

Question 3: Hardware vs. Software

- question

Question 3: Hardware vs. Software

Recall the differences between hardware and software. Which of the following are true about **software**?

It can be fully described by a string of ones and zeros. ✓
 Our pipelined Y86 implementation is an example of it.
 You can touch it.
 You could potentially speed it up by dividing it between threads. ✓
 It can always be trusted to perform correctly.
 You could potentially speed it up by adding a cache. ✓

Select all possible options that apply. 

 100%

This question is complete and cannot be answered again.

- pretty straightforward
- note that the pipelined implementation was a hardware thing
- also note that you can potentially speed up hardware or software using a cache

Question 4: Multicore Processor Access Rate

- question

Question 4: Multicore Processor Access Rate

Consider a multicore processor with 2 cores, with each core having its own direct-mapped cache that uses the writeback and write-allocate policies.

Further consider the following:

- Each cache has 16-byte cache lines (so the rightmost hexadecimal digit of an address is the offset)
- Each cache has 16 slots (so the second-from-the-right hexadecimal digit of an address is the index)
- Memory addresses are 12-bits in size (so the leftmost hexadecimal digit of an address is the tag)
- Each processor reads or writes a single byte
- We're using the MESI protocol (as illustrated in the diagram at the end of this question)

Now, assuming in each case that all caches are initially empty, calculate the number of misses for Core 1 for each of the following 4 access sequences.

- something potentially tricky to watch out for is if a line has the same index bit and gets put in the cache, it might evict something else that was already there
- also it's only asking for misses on core 1

- part 1

Consider this access pattern, but ignore the extra space in the middle:

```
Core 1 reads from 0x47e
Core 1 writes to 0x47e
Core 1 reads from 0x4ca
```

```
Core 1 reads from 0x47e
Core 1 reads from 0x4ca
Core 1 writes to 0xdca
```

Number of misses for Core 1 = 3



✓ 100%

1. core 1 miss (put **0x47e** into the cache)
 2. core 1 hit
 3. core 1 miss (put **0x4ca** into the cache, in a non-conflict spot)
 4. core 1 hit
 5. core 1 hit
 6. core 1 miss (put **0xdca** into the cache)
 - so 3 total misses for core 1
- part 2

Now consider this almost identical access pattern with one additional access made by Core 2 in the middle:

```
Core 1 reads from 0x47e
Core 1 writes to 0x47e
Core 1 reads from 0x4ca
```

```
Core 2 reads from 0x4ca
```

```
Core 1 reads from 0x47e
Core 1 reads from 0x4ca
Core 1 writes to 0xdca
```

Number of misses for Core 1 = 3



✓ 100%

- first 3 is the same A→2 miss and we have **0x47e** and **0x4ca** in Core 1's cache
 - 4. core 2 miss → **0x4ca** is now in S state (still usable by core 1)
 - 5. core 1 hit
 - 6. core 1 hit (since **0x4ca** did not get evicted)
 - 7. core 1 miss
 - still just 3 miss for core 1
- part 3

Consider this access pattern, but ignore the extra space in the middle:

```
Core 1 writes to 0x84c
Core 1 reads from 0x8e4
Core 1 writes to 0x143
```

```
Core 1 reads from 0x143
Core 1 reads from 0x8e4
Core 1 writes to 0x1e6
```

Number of misses for Core 1 = 4



✓ 100%

1. core 1 miss (put `0x84c` into the cache)
 2. core 1 miss (put `0x8e4` into the cache - non conflict)
 3. core 1 miss (put `0x143` into the cache - non conflict)
 4. core 1 hit
 5. core 1 hit
 6. core 1 miss (put `0x1e6` into the cache and evict `0x8e4`)
- 4 misses for core 1

- part 4

Now consider this almost identical access pattern with one additional access made by Core 2 in the middle:

Core 1 writes to `0x84c`
 Core 1 reads from `0x8e4`
 Core 1 writes to `0x143`

Core 2 writes to `0x143`

Core 1 reads from `0x143`
 Core 1 reads from `0x8e4`
 Core 1 writes to `0x1e6`

Number of misses for Core 1 =



✓ 100%

- first 3 is the same A→3 miss and we have `0x84c`, `0x8e4`, `0x143` in our cache
- core 2 miss
 - since `0x143` was in M in Core 1, a remote write would cause us to invalidate that
 - so core 1 cache now only has `0x84c`, `0x8e4`
- 5. core 1 miss (put `0x143` into the cache)
- 6. core 1 hit
- 7. core 1 miss
 - in total there's 5 misses for core 1

Question 5: Ordering Memory

- question

Question 5: Ordering Memory

Each line below corresponds to some level in the memory hierarchy. For each level you are provided with some details about the cost, size, or performance of the memory at that level.

Use all the memories to build a hierarchy of memories, placing the memory closest to the CPU at the top and the memory farthest from the CPU at the bottom.

Correct answer

Correct answer (there may be other correct orders):

- Size: 2 KB, Latency: 6 ns
- Cost: 58 cents per KB, Size: 32 MB
- Cost: 54 cents per MB, Latency: 42 ns
- Cost: 51 cents per GB, Size: 32 GB
- Cost: 13 cents per GB, Latency: 320 ms

- general idea is the closer you are to the memory, you have smaller size, smaller line size, faster, and more expensive
- (I got the last 2 wrong, because 51 cents per GB means it's more expensive than 13 cents per GB but I got that backwards)

Question 6: File Descriptor Basics

- question

Question 6: File Descriptor Basics

Below are a number of statements. Select all that are **always** TRUE.

If two different processes open the same file, their open file table entries will refer to the same vnode.

 If two different processes open the same file, they will get the same FD.

If a process opens the same file twice, the file descriptors returned will map to different open file table entries.

 If two threads in the same process (program) open the same file, they will get the same FD.

Select all possible options that apply.

100%

This question is complete and cannot be answered again.

1. True, there's only 1 vnode per file
2. False, they'll get separate FD if they call `open` separately, each with their own offset
3. True, if they call `open` separately, they'll get 2 different FD which have their own offset so they have to point to different things in the OFT (though the entries in OFT will point to the same vnode)
4. False, the actual value of the FD is dependent on how many files you currently have opened, so it's impossible to say without knowing the code for the threads

Question 7: Investigating Overhead in FS Indexes

- question

Question 7: Investigating Overhead in FS Indexes

When we use a hybrid index structure, we have multiple disk addresses in the inode; some number of each of direct, indirect, double indirect, and triple indirect addresses. When we use a multi-level tree index, we only have one address in the inode, the address of the root of the tree (or, if the file needs only a single data block, the address of that data block - a tree of height 0). In this question we are going to make a more fair comparison between these index structures by allowing both structures the same number of addresses in the inode.

To be specific, consider that there are 11 disk addresses in each inode. When we use a hybrid index, these are broken down as 8 direct, 1 indirect, 1 double indirect, and 1 triple indirect addresses. When we use a multi-level index, there will also be 11 disk addresses in the inode, each of which is the root of a tree; all of these trees will be the same height (or won't exist at all if they are not needed). The maximum height for the multi-level tree is 3.

In these file systems, the disk blocks are 8192 bytes and a disk address is 8 bytes.

- each block can contain $8192/8 = 1024$ addresses
 - **important note:** for the multi-index, if the file size is like 11, all of the entries can be direct pointers to the data block itself
- part 1

Given these parameters, give the size of a file (in disk blocks) that will use more blocks to store the index when using a multi-level index than when using a hybrid index.

1



x 0%

- if we have 1025 blocks, multi-index would require 2 pointer blocks (1 indirect pointer for the first 1024, and 1 more for the extra 1 block)
- hybrid would just need 1 indirect pointer (because first 8 is direct), so it can cover 1018 using the indirect

- part 2

Given these parameters, give the size of a file (in disk blocks) that will use more blocks to store the index when using a hybrid index than when using a multi-level index.

9224



x 0%

- it's best to think about these as "if they'll need index pointers"
 - so a file that can be represented with all direct pointers mean there are 0 index pointers
 - and since both representation will have the same number of direct pointers (to actually point to the data block), these are not too interesting, we'll rather look at the number of index pointers they'll need
- the multi-level index can have 11 direct pointers while the hybrid only has 8 direct pointers
 - so if our file size is 9, multi-level can still use direct pointers (0 index pointers) while hybrid will need to use its indirect pointers (1 index pointers)
 - so this is a valid solution

- part 3

We are also concerned about the number of IOs required to read a particular block given its logical block number or LBN, where LBNs start at 0. For the two questions below, imagine a file where none of the data or metadata is yet stored in memory, and assume the file is large enough to trigger the desired behaviour.

What is the smallest LBN for which the hybrid index can take fewer IOs than the multi-level index to fetch the associated block?

0

?

✓ 100%

- important note: we can assume that the file is as big as needed
 - so if our file is big enough such that the multi-index requires indirect or double indirect blocks (similarly hybrid might need too as well)
 - so in this case, if we access LBN 0
 - for hybrid, we'll ALWAYS use a direct pointer
 - but for multi-index, depending on the size, we MAY use a indirect pointer (or even double-indirect) → and we can make the file size as big as possible
 - part 4
- What is the smallest LBN for which the multi-level index can take fewer IOs than the hybrid index to fetch the associated block?
- 1032
- ?
- ✗ 0%

- a bit related to part 2
- if we can get the hybrid index to use indirect block and multi-level index to use all direct blocks, multi-level will require less IO access
- so if we have size 9, this will be the expected behaviour → this is LBN 8

Question 8: Bumping Cache Parameters

- question

Question 8: Bumping Cache Parameters

Imagine we have a cache with a particular configuration:

- cache size (in bytes of data, not metadata);
- number of bits of tag, index, offset, and overall address;
- associativity;
- number of sets; and
- line size.

We **decrease the associativity**.

This change requires other changes to keep the cache configuration consistent. There may be many strategies to accommodate the change, but some other changes are either unnecessary or counterproductive to every strategy. (A necessary change accommodates the initial change or indirectly does so by accommodating another necessary change. If an example would help, one is given below.)

Select all changes below that are necessary parts of some reasonable strategy. (Choose changes that are necessary to any strategy, even if they are not in the same strategy or lack other changes in that strategy.)

- (a) The number of sets increases.
- (b) The overall cache size decreases.
- (c) The number of tag bits increases.
- (d) The cache line size decreases.

Select all possible options that apply. There are exactly 2 correct options in the list above.

0%

- answer key: If we decrease the associativity, we might:

1. Decrease the overall cache size, while holding everything else constant.
 - in this case, can think of decreasing the associativity as literally decreasing the number of slots we have, so just need less index bits
 - and if we're keeping everything else the same we can just think of it as shrinking the cache
2. Increase the number of index bits and therefore the number of sets. This then requires we decrease the number of tag bits or increase the overall number of address bits.
3. Increase the number of offset bits and therefore the cache line size. This then requires we decrease the number of tag bits or increase the overall number of address bits.

Question 9: Stride Exercises Across a Fixed-Size Array

- question

Question 9: Stride Exercises Across a Fixed-Size Array

Imagine we have a cache with the following parameters:

- Total size of 512 KB (2^{19} bytes)
- 32768 (2^{15}) total lines
- 16 (2^4) byte line size
- Associativity (and therefore number of sets) described in each part below.
- When an eviction policy is needed: LRU eviction policy

To study our cache, we create a global array A of `uint8_t` values of total size 1GB, which is 16-byte-aligned. Then, we call a function like the following:

```
void test(int stride) {  
    for (int i = 0; i < 10000; i++) {  
        for (int j = 0; j < 16397; j++) {  
            read the memory at A[j*stride] // this is pseudocode (correct, but not real code)  
        }  
    }  
}
```

Assume the reads from A in the inner loop are the only memory accesses. Answer carefully and leave no answer blank.

- part 1

Starting from an empty, **direct-mapped** cache, exactly how many misses occur when we call `test(32)`?

misses: 146384



x 0%

- note that this is direct mapped- so we'll have the "affected" vs "unaffected" cache sets problem that we had on the quiz
 - i.e some pairs of lines accessed map to the same slot in our direct-mapped cache → so each time we access one of such a pair, we evict the other one, and we end up missing on every access to such pairs
- **tricky part:** each stride is 2 lines
 - this means that $j = 0$ maps to cache 0, but $j = 1$ maps to cache 2, $j = 2$ maps to cache 4, etc
 - due to the nature of the stride size and the fact that our cache is direct mapped (whole-finely dependent on the index bits), cache 1, cache 3, cache 5, will never be used
 - to make things a bit easier → we can think of this as saying our cache lines is 32 bytes, and decreasing the number of cache slots we have
 - (this works because even if we load in 32 bytes, the last 16 bytes is never used/empty, so the logic is the same)
 - so we say we have **16384 total lines**
- the inner loop tries to access **16397** lines → this is slightly bigger than our cache so we'll have some "spillage"
 - $16397 - 16384 = 13$ cache lines will spill over
 - this "spilled over" cache slots/lines will incur 2 miss (1 for the original line, then 1 more for the extra spilled line - just like the quiz question)
 - and since they keep evicting each other, you will have to incur these 2 misses for every outer-loop iteration

- the other 16384 lines will only incur compulsory misses
- finally, the math

$$\begin{aligned}\text{total misses} &= \text{affected miss} + \text{unaffected miss} \\ &= 10,000(13)(2) + (163884 - 13) \\ &= 276,371\end{aligned}$$

(I was shockingly closed, I just missed the fact that we incur 2 miss per cache lines)

- part 2

Starting from an empty, **direct-mapped** cache, exactly how many misses occur when we call `test(8)`?

misses: 8199

?

✓ 100%

- this is a more classic question, since the stride size is smaller than the line size, we can pay attention to misses per line instead
- we are trying to access $16397 \times 8 = 131176$ bytes
 - (note that this is the number of bytes we're trying to access - not indices, index of the last byte would be - 1 of the quantity above)
 - **this can fit fully into the cache**
 - this quantity will require $131176/16 = 8198.5 \approx 8199$ cache lines
- we take 1 miss per cache line and we only ever take compulsory misses on that first iteration
- so the answer is 8199 misses

- part 3

Starting from an empty, **fully-associative** cache, exactly how many misses occur when we call `test(16)`?

misses: 16397

?

✓ 100%

- for this problem, we no longer has the weird empty lines problem, since it's fully associative, if there's an empty spot, we will fit it in there
- so we're trying to access 16397 cache lines (the inner for loop)
 - we have enough slots in our cache to fit this many lines
 - again, we won't have the weird overlapping/empty slots issue because we always use up the entirety of our cache
- so we fit them all into our cache and every outer loop iteration, we'll have it in the cache → just take the original compulsory misses
- in total we take 16397 compulsory miss

- part 4

Starting from an empty, **fully-associative** cache, exactly how many misses occur when we call `test(32)`?

misses: 163970000

?

✗ 0%

- for this question, the stride is a bit bigger, but again, since we don't care about indices → if there's open space we'll put it there
- so in the same way, we only take compulsory misses in the first iteration
- note: we can't fit the entire array in our cache (i.e $16397 * 32 > 2^{19}$) but this **does not matter**
 - (this is where I went wrong during the exam)
 - because we do not try to access the entire array
 - we try to access 16397 different cache lines within the array, and we only have to remember these specific cache lines
 - so the fact that our cache can fit 16397 cache lines in it is enough
 - (for direct mapped cache, the index actually matters for the mappings so that's why we tend to care about how big the actual array is, we want to be wary of the wrap arounds)
- total is 16397 misses

- part 4: freebie

Imagine:

1. We change the code above to *write* rather than *read* the memory at $A[j*stride]$.
2. The cache is fully-associative.
3. We call `test(1)`.

Correctly finish the following statements:

1. A **✓ 100%** policy will **reduce the number of misses**. However, ...
2. To gain performance benefit from this change we must also use a **✓ 100%** policy.

- key:

The write-allocate policy ensures that we do actually allocate a line in the cache on write. Otherwise, we have no opportunity to exploit the substantial spatial and temporal locality available in this problem.

However, reducing the number of misses but using a writethrough policy doesn't actually confer any practical performance benefit. With a writethrough policy, we'll pay essentially the cost of a miss on every write, anyway. So, yes, many of these accesses will be hits, but they'll be expensive hits tantamount to misses.

With a writeback policy, we instead "batch" our writes until the cache lines are evicted (in this case sometime after the end of the function call). This gives us a huge performance benefit.

Quiz 5

In Class

I24: Directories

- I24.1: Decode Directories (Inode to Name)

I24.1. Decode Directories (Inode to Name)

The following is the hexdump of a directory that uses the following structure for its dirents.

```
struct dirent {
    ino_t d_ino;           // 32 bits
    __uint16_t d_reclen;
    __uint8_t d_type;
    __uint8_t d_namlen;
    char d_name[255+1];
};

#define DT_DIR 4
#define DT_REG 8
#define DT_LNK 10
```

What is the name of the object with inode number 0x66?

You may find it convenient to refer to an online hex to string converter, such as:

<https://onlinestringtools.com/convert-hexadecimal-to-string>.

```
00000000 00 00 10 00 0c 00 00 04 01 2e 00 00 00 01 00 10 00
00000010 0c 00 04 02 2e 2e 00 00 64 00 00 00 14 00 08 08
00000020 4c 61 6e 64 66 6f 77 6c 00 7b f7 18 02 00 10 00
00000030 10 00 04 05 54 61 70 69 72 00 77 6c 00 00 00 00
00000040 14 00 08 0a 50 6f 6c 61 72 20 62 65 61 72 00 18
00000050 65 00 00 00 10 00 0a 07 47 61 7a 65 6c 6c 65 00
00000060 03 00 10 00 14 00 04 0b 4d 6f 63 6b 69 6e 67 62
00000070 69 72 64 00 66 00 00 00 14 00 08 0a 50 6f 6c 61
00000080 72 20 62 65 61 72 00 00
00000088
```

- you could use the `rec_len` field to skip over directory that you know isn't `0x66` inode
- so we look for `66 00 00 00` since it is in little-endian
 - we get `66 00 00 00 14 00 08 0a 50 6f 6c 61 72 20 62 65 61 72 00 00`
 - `66 00 00 00` is the inode
 - `14 00` is the `rec_len`
 - `08` is the type (regular)
 - `0a` means the name length is 10
 - DO NOT CONVERT TO LITTLE ENDIAN since they are just a series of bytes
 - this translates to "Polar bear"
 - `50 6f 6c 61 72 20 62 65 61 72 00 00` is the name
- I24.2: Decode Directories (Inode to Type)

I24.2. Decode Directories (Inode to Type)

The following is the hexdump of a directory that uses the following structure for its dirents.

```
struct dirent {
    ino_t d_ino;           // 32 bits
    __uint16_t d_reclen;
    __uint8_t d_type;
    __uint8_t d_namlen;
    char d_name[255+1];
};

#define DT_DIR 4
#define DT_REG 8
#define DT_LNK 10
```

What is the type of the object whose inode number is 0x100002?

```
0000000 00 00 10 00 0c 00 04 01 2e 00 00 00 01 00 10 00
0000010 0c 00 04 02 2e 2e 00 00 64 00 00 00 14 00 08 0b
0000020 48 65 72 6d 69 74 20 63 72 61 62 00 02 00 10 00
0000030 14 00 04 0a 4d 65 61 64 6f 77 6c 61 72 6b 00 00
0000040 65 00 00 00 14 00 08 08 44 6f 72 6d 6f 75 73 65
0000050 00 6b 00 00 66 00 00 00 14 00 0a 0a 47 75 69 6e
0000060 65 61 20 70 69 67 00 00 00 00 00 00 10 00 08 04
0000070 4d 6f 74 68 00 61 20 70 00 00 00 00 14 00 08 0b
0000080 52 61 74 74 6c 65 73 6e 61 6b 65 00 00 00 00 00
0000090 18 00 08 0f 45 6e 67 6c 69 73 68 20 70 6f 69 6e
00000a0 74 65 72 00 67 00 00 00 10 00 08 04 4d 6f 74 68
00000b0 00 73 68 20
00000b4
```



[Save & Grade](#) Single attempt

[Save only](#)

Additional attempts available with new variants [?](#)

- cheat a bit and look for `02 00 00 10` (little endian)

```
0000020 48 65 72 6d 69 74 20 63 72 61 62 00 02 00 10 00
0000030 14 00 04 0a 4d 65 61 64 6f 77 6c 61 72 6b 00 00
```

- so we can see that `rec_len` is `0x0014`
- and the next byte after that is `d_type` and its `04` meaning it's a directory

- I24.3: Decode Directories (Name to Inode Number)

I24.3. Decode Directories (Name to Inode Number)

The following is the hexdump of a directory that uses the following structure for its dirents.

```
struct dirent {
    ino_t d_ino;           // 32 bits
    __uint16_t d_reclen;
    __uint8_t d_type;
    __uint8_t d_namlen;
    char d_name[255+1];
};

#define DT_DIR 4
#define DT_REG 8
#define DT_LNK 10
```

What is the inode number (in hex) for the file Wildebeest ?

You might find it convenient to refer to an online text to hex converter such as

<https://onlinestringtools.com/convert-string-to-hexadecimal>

```
00000000 00 00 10 00 0c 00 04 01 2e 00 00 00 01 00 10 00
00000010 0c 00 04 02 2e 2e 00 00 64 00 00 00 10 00 08 07
00000020 4c 61 62 20 72 61 74 00 65 00 00 00 10 00 08 07
00000030 48 61 72 72 69 65 72 00 66 00 00 00 14 00 0a 0a
00000040 57 69 6c 64 65 62 65 65 73 74 00 10 67 00 00 00
00000050 10 00 0a 07 57 61 72 62 6c 65 72 00 00 00 00 00
00000060 10 00 08 07 43 68 69 63 6b 65 6e 00 68 00 00 00
00000070 10 00 0a 06 44 6f 6e 6b 65 79 00 00 69 00 00 00
00000080 10 00 0a 05 48 65 72 6f 6e 00 00 00 6a 00 00 00
00000090 0c 00 0a 03 43 61 74 00 6b 00 00 00 10 00 08 07
000000a0 43 68 69 63 6b 65 6e 00
000000a8
```

0x

?

[Save & Grade](#) [Single attempt](#)

[Save only](#)

Additional attempts available with new variants [?](#)

- "Wildebeest" in hex is **57 69 6c 64 65 62 65 65 73 74** (can look for this directly because it's a string - no such thing as endianness)

```
00000030 48 61 72 72 69 65 72 00 66 00 00 00 14 00 0a 0a
00000040 57 69 6c 64 65 62 65 65 73 74 00 10 67 00 00 00
```

- working backwards we see **d_namelen** is **0x0a**
- d_type** is **0x0a**
- d_reclen** is **0x0014**
- and **d_ino** is **0x 00 00 00 66**

I25: Comparing Ext2 and V

- I25.1: ext2 versus V6 max file size

I25.1. ext2 versus V6 max file size

Given a V6 file system and an ext2 file system that use the same number of bytes to represent a disk address, if they have the same block size (which is at least 512 bytes), which can represent a larger file?

- (a) ext2
- (b) V6
- (c) It is impossible to tell

Save & Grade Single attempt

Save only

Additional attempts available with new variants

- we can think of this question in terms of data blocks
 - let's just say the number of address we can have in a block is `num_ptr`
 - at most, v6 can have 7 indirect pointers and 1 double-indirect pointers
 - that's `7 * num_ptr + (num_ptr ** 2)`
 - at most, ext2 can have 12 direct, 1 indirect, 1 double indirect, 1 triple indirect
 - that's `(12 * num_ptr) + (num_ptr) + (num_ptr ** 2) + (num_ptr ** 3)`
 - so it's clear to see that ext2 can hold more data blocks can represent bigger files
- I25.2: ext2 versus V6 ease of contiguous allocation

I25.2. ext2 versus V6 ease of contiguous allocation

Which file system is likely to be better able to allocate files contiguously on the disk?

- (a) ext2
- (b) V6
- (c) It is impossible to tell

Save & Grade Single attempt

Save only

Additional attempts available with new variants

- v6 uses a free list and ext2 uses a bitmap within a block group
 - there's more locality in ext2 (because it's bitmap per block group and you try to allocate files within 1 block group)
 - so ext2 will do a better job storing files contiguously on disk
- I25.3: ext2 versus V6 directory entries (simple)

I25.3. ext2 versus V6 directory entries (simple)

Given a large directory of names, all of which are greater than fourteen bytes in length, which file system will consume more space storing this directory?

- V6
- ext2
- V6 cannot represent such a directory
- ext2 cannot represent such a directory
- Both file systems will consume the same amount of space

Save & Grade *Single attempt*

Save only

Additional attempts available with new variants

- in v6, directory entries can only have names that's up to 14 bytes in length
- so v6 can't even represent this directory

- I25.4: ext2 versus V6 number of IOs

- variant 1

I25.4. ext2 versus V6 number of IOs

Assuming that the inode is already in memory, that disk addresses are 4 bytes, and that the file system block size is 1024, how many IOs will it take in V6 to read logical block 6723 from a file consisting of 39768 blocks?

integer

IOs



Save & Grade *Single attempt*

Save only

Additional attempts available with new variants

- a block can store $1024/4 = 256$ pointers per block
- this file is pretty big - so we'll use the set up where we have 7 indirect and 1 double-indirect
- $6723 - (256 \times 7) = 4481$ so the block doesn't live in any of the 7 indirect, we need to go to double indirect
- since it's in the double indirect, you need to read the inode, the double indirect, the indirect, then the data itself → so 4 reads
- but since the inode is already in memory, you only have to do 3 reads

- variant 2

I25.4. ext2 versus V6 number of IOs

Assuming that the inode is already in memory, that disk addresses are 4 bytes, and that the file system block size is 1024, how many IOs will it take in ext2 to read logical block 387114 from a file consisting of 3529295 blocks?

integer IOs ?

Save & Grade *Single attempt*

Save only

Additional attempts available with new variants ?

- a block can store $1024/4 = 256$ pointers per block
- $387114 - 12 = 387102$ so it doesn't live in the direct block
- $387114 - (256) - (256^2) = 321310$ so it doesn't live in the indirect block or the double-indirect block
- thus it lives in the triple indirect block
 - you need to read the inode, the triple indirect, the double indirect, the indirect, then the data itself → so 5 reads
 - but since the inode is already in memory - you only have to do 4 reads

I28: Translating addresses using a TLB

- I28.1: TLB translations (32-bit addresses)

I28.1. TLB translations (32-bit addresses)

Consider a machine with a 32-bit physical address space with 4096-byte pages running an operating system that supports a 32-bit virtual address space.

Assume that the following table shows you the contents of a TLB. In the box, enter the physical address that will be produced when the TLB tries to translate a read request from a user process for virtual address `0xc9c74051`. If the TLB should produce a fault, enter the word 'fault' in the box.

Virtual page number	Physical page number	Permissions	mode
0x59c8d	0x83db6	RW	user
0x0ad48	0x66a94	RW	user
0xc9c74	0x2be1e	RWX	user
0xa0294	0xba739	RWX	supervisor
0x99b0a	0xb5003	RX	supervisor

?

Save & Grade *Single attempt*

Save only

Additional attempts available with new variants ?

- we need $\log_2(4096) = 12$ bits or 3 hex digits for the offset
- so the VPN is `0xc9c74` → this maps to `0x2be1e`
 - we have the permissions and privilege needed

- need to glue the offset (**051**) at the end of the PBN to get **0x2be1e051**
- note: if we didn't have the permission or privilege required then we just fault
- I28.2: TLB translations (random sizes)

I28.2. TLB translations (random sizes)

Consider a machine with a 30-bit physical address space with 256-byte pages running an operating system that supports a 24-bit virtual address space.

Assume that the following table shows you the contents of a TLB. In the box, enter the physical address that will be produced when the TLB tries to translate a write request from a user process for virtual address **0xfc2488**. **Be sure to format your answer with the correct number of digits, given the size of the physical address space.**

If the TLB should produce a fault, enter the word 'fault' in the box.

Virtual page number	Physical page number	Permissions	mode
0xae82	0x2353c1	RWX	user
0x27d0	0x34a7d4	R	user
0x3dc6	0x1bde9c	R	user
0xebea	0x15eabb	RW	user
0x8e0f	0x1b6363	RWX	user
0xcd7b	0x320d4c	RX	user
0xfc24	0x20a37b	RWX	user
0xd355	0x13a198	none	user

0x



Save & Grade Single attempt

Save only

Additional attempts available with new variants

- we need $\log_2(256) = 8$ bits or 2 hex digits for the offset
- so VPN is **0xfc24** → map to **0x20a37b**
 - we have the permissions required
 - glue the offset (**88**) on at the end and get **0x20a37b88**
- I28.3: TLB translations (unaligned offset)

I28.3. TLB translations (unaligned offset)

Consider a machine with a 16-bit physical address space with 64-byte pages running an operating system that supports a 12-bit virtual address space.

Assume that the following table shows you the contents of a TLB. In the box, enter the physical address that will be produced when the TLB tries to translate a read request from a user process for virtual address 0x664. **Be sure to format your answer with the correct number of digits, given the size of the physical address space.**

If the TLB should miss or produces an exception, enter the word 'fault' in the box.

Virtual page number	Physical page number	Permissions	mode
0x34	0x323	none	supervisor
0x30	0x281	RW	user
0x31	0x090	R	user
0x19	0x1a7	RW	user
0x13	0x3c2	RW	supervisor



Save & Grade Single attempt

Save only

Additional attempts available with new variants

- we need $\log_2(64) = 6$ bits
 - this unfortunately does not translate very well to hex digit, so we'll have to operate in bits
- $0x664 = 0b11001\ 100100$
 - offset is $0b100100$
 - VPN is $0b11001 = 0x19$ (add 0s at the front)
- this VPN maps to $0x1a7$ as the PBN
 - we have the permission and privilege required
 - in binary this is $0b110100111\ 100100$ (added offset at the end)
 - in hex this is $0x69E4$

I29: VM - Page Tables and Page Replacement

- I29.1: Counting page faults

I29.1. Counting page faults

Consider the following C code:

```
uint64_t buffer[1024 * 1024];
uint64_t sum = 0;
for (i = 0; i < (sizeof(buffer) / sizeof(uint64_t)); i++) {
    sum += buffer[i];
}
```

Assume the following:

- When the program begins executing, none of its data is in memory.
- You are concerned only with page fault due to accesses to `buffer`.

If the virtual memory page size is 2048, how many page faults will this program take?

integer



[Save & Grade](#) [Single attempt](#)

[Save only](#)

Additional attempts available with new variants

- from the loop parameter does not matter at all
 - we will be accessing the very last byte of the buffer – so to get that data we will need the array in memory at some point
- the size of `buffer` is $1024 \times 1024 \times 8 = 2^{23}$ (times 8 because each element is a 8-byte int)
- since the data is allocated contiguously, $(1024 \times 1024 \times 8)/2048 = 2^{23}/11 = 2^{12} = 4096$ pages will be needed
 - this means that we'll take 4096 page fault as we're loading it in (since we start with none in memory)

- I29.2: Single-level page table VM systems

I29.2. Single-level page table VM systems

Consider a machine that supports a 19-bit virtual address space with a single flat page table (that fits into a single page). If pages are 1024 bytes, PTEs are 2 bytes, and each PTE requires five bits of metadata, how many bits does it take to represent a virtual page number?

integer



[Save & Grade](#) [Single attempt](#)

[Save only](#)

Additional attempts available with new variants

- pages are 1024 bytes (both in physical and virtual space) means that we need $\log_2(1024) = 10$ bits for the offset
- since our virtual address space is 19 bits, the VPN is $19 - 10 = 9$
- TODID: what about the metadata business
 - metadata is more so for the physical address and is related with the PTE

- I29.3: Single-level page table VM systems

I29.3. Single-level page table VM systems

Consider a machine that supports a 18-bit virtual address space with a single flat page table (that fits into a single page). If pages are 1024 bytes, PTEs are 4 bytes, and each PTE requires five bits of metadata, what is the maximum number of physical pages that this machine can support?

Assume that this architecture, unlike the x86, allows the page number to occupy bit positions in the PTE that overlap with offset bits in the address. (That is, you may assume no unused bits in the PTE.)

integer



Save & Grade *Single attempt*

Save only

Additional attempts available with new variants

- since we're talking about physical pages → look at the PTE (that's what the VA maps to)
 - **note:** understanding this question
 - usually, if our PTE is 4 bytes long, that means our physical address space is 32 bits long
 - however, since we have the offset, our physical address space is actually only $32 - \log_2(\text{page_size})$ - because we then append the offset at the end, then we get 32 bit long address again
 - in this scenario, they are saying we leave no space for the offset bits, meaning we can (technically) use all of the 32 bits to represent PBN (barring the metadata)
 - further: this means that the OS will need a new way to figure out how indexing works
 - so now we our PBN can use 32 bits - but we need 5 bits for the metadata, so it's $32 - 5 = 27$
 - so we can represent 2^{27} page
- I29.4: Single-level page table VM systems (aligned)

I29.4. Single-level page table VM systems (aligned)

Consider a machine that supports a 18-bit virtual address space with a single flat page table (that fits into a single page). If pages are 1024 bytes, PTEs are 4 bytes, and each PTE requires five bits of metadata, what is the maximum number of physical pages that this machine can support?

Assume that this architecture **requires** that the page number is placed in the PTE such that it does not occupy any bit positions that overlap with offset bits. (I.e., the PTE is organized similarly to PTEs on the x86.)

integer



Save & Grade *Single attempt*

Save only

Additional attempts available with new variants

- here, we need to reserve some bits for offsets or metadata
 - you will need $\log_2(1024) = 10$ bits for offset
 - but remember that our x86 PTE points to the start of the page, so we don't actually need the 10 offset bits
 - so the PTE entry can potentially look like `0x101010 xxxxxxxx`
 - where the s are where the offset should be → you can imagine them being all 0s as we are physically pointing to the start of a page

- however, people thought that using those bits as all 0s is very wasteful (after all, you can just right shift and get rid of those bits, same thing as having them as all zeros) → so instead they'll use these offset bits for something else...metadata!
- so here, we can use all these x's as metadata bits, since we have less metadata bits than we do offset bits, some of them will even be unused → point is, we don't have to allocate any more bits for metadata, they can just use the offset bits
 - you only need 5 bits for the metadata so you can use that and leave 5 unused
- 4 byte PTE means you have 32 bits, then we have $32 - 10 = 22$ bits
- so we can represent 2^{22} bits
- I29.5: Single-level page table VM systems

I29.5. Single-level page table VM systems

Sometimes we have unused bits in a virtual address (e.g., the x86 is a 64-bit machine, but supports only a 48-bit virtual address space, leaving 16 bits unused).

Consider a machine with 16-bit virtual addresses with virtual address spaces represented by a single flat page table (that fits into a single page).

If pages are 256 bytes, PTEs are 2 bytes, and each PTE requires five bits of metadata, how many bits in the virtual address are unused?

integer



Save & Grade Single attempt **Save only**

Additional attempts available with new variants ?

- recall that we use some bits in the VPN to serve as an index into the page tables

- note that here we only have 1 level page table while in the x86 there was 4 levels

- our PTE are 2 bytes so 16 bits

- we can store $256/2 = 128$ PTE per page tables
- that means we'll need $\log_2(128) = 7$ bits as an index
- since pages themselves are 256 bytes, we'll need $\log_2(256) = 8$ bits as an offset
- so we have $16 - 7 - 8 = 1$ bit leftover
- so we have 1 unused bit

- note: this is kinda similar to caching now??

I30: The Clock Algorithm

- I30.1: The Clock Algorithm

I30.1. The Clock Algorithm

Assume that memory holds only 4 pages, numbered 0 to 3. The following table lists the pages, in-order, starting at 0 and shows the virtual page number currently resident in that physical page and its use bit. The clock hand is initially at position 2.

Fill in the table below with the correct page numbers and use bits after the following operations have been performed. In each operation, the hex value is a virtual page number; enter virtual page numbers exactly as they appear in the reference stream. Enter a 0 or 1 for the use bit.

0xc0, 0xb1, 0xc0, 0x8b, 0x8b, 0x21, 0xc0, 0x8b

Virtual page number	Use
0xc0	1
0xc1	0
0x8b	1
0xb1	1

- initial state: `0xc0 (1), 0xc1 (0), 0x8b (1), 0xb1 (1)`, hand at position 2 (`0xb1`)
- access `0xc0`
 - new state: `0xc0 (1), 0xc1 (0), 0x8b (1), 0xb1 (1)` (nothing changes)
- access `0xb1`
 - new state: `0xc0 (1), 0xc1 (0), 0x8b (1), 0xb1 (1)` (nothing changes)
- access `0xc0`
 - new state: `0xc0 (1), 0xc1 (0), 0x8b (1), 0xb1 (1)` (nothing changes)
- access `0x8b`
 - new state: `0xc0 (1), 0xc1 (0), 0x8b (1), 0xb1 (1)` (nothing changes)
- access `0x8b`
 - new state: `0xc0 (1), 0xc1 (0), 0x8b (1), 0xb1 (1)` (nothing changes)
- access `0x21`
 - need to evict
 - `0xb1` is in use, so we'll skip it, but set the used bit to 0
 - `0xc0` (because clockwise) is in use, so we'll skip it, but set the used bit to 0
 - `0xc1` is not in use → replace it with `0x21` and mark the use bit
 - new state: `0xc0 (0), 0x21 (1), 0x8b (1), 0xb1 (0)`
- access `0xc0`
 - new state: `0xc0 (1), 0x21 (1), 0x8b (1), 0xb1 (0)`
- access `0x8b`
 - new state: `0xc0 (1), 0x21 (1), 0x8b (1), 0xb1 (0)` (nothing changes)

- I30.2: The Two-handed Clock Algorithm

I30.2. The Two-handed Clock Algorithm

Assume that memory holds only 4 pages, numbered 0 to 3. The following table lists the pages, in-order, starting at 0 and shows the virtual page number (with the current state of their use and dirty bits) currently resident in the given physical page.

The replacement hand is initially at position 0. The write hand is initially at position 2.

On every eviction, you should:

- First attempt to write one dirty page back. It should first consider the page under the dirty hand; if it does not find a dirty page to evict, the hand should return to the position at which it started. If it does find a page to evict, leave the hand on the location that was just written.
- Second, you should select a page for eviction, using the algorithm as presented in the pre-class work.

In other words, you should move the write hand **before** you move the replacement hand.

In a real system, the tricky part of this algorithm is keeping the write hand moving at 'the right' rate. You should ignore this. It is fine if your write hand passes your replacement hand or if your replacement hand passes your write hand. For example, if there are no dirty pages to write back, then the write hand will end up in exactly the same position where it started (and will have passed the replacement hand).

Some corner cases:

- It is possible that the replacement hand will want to replace a page that is dirty. Assume that you do just that, write the page you are about to replace and then read the new page in. Leave the clock hand in the slot that you just replaced.
- When replacement causes a page to be written, advance the write hand to the next slot.

Fill in the table below with the correct page numbers and use bits after the following operations have been performed.

In each operation, the hex value is a virtual page number; enter virtual page numbers **exactly** as they appear in the reference stream (cut and paste is your friend). Enter 0 or 1 for the use and dirty bits.

(very wordy)

write 0xed, write 0xe1, read 0x5f, read 0x1b, read 0xed, write 0x1b

Virtual page number	Use	Dirty
0x2c	0	0
0xe1	1	1
0xe0	0	1
0x1b	1	1

- same as above, just remember
 - when you need to evict something, you move the write hand first, and you write back 1 thing
 - remember as you're traversing over other in used entries with replacement hand, set the used bit to 0
- answer

Correct answer

Virtual page number	Use	Dirty
0xed	1	1
0xe1	0	1
0x5f	1	0
0x1b	1	1

Practice Quiz

R5.1: Unix directory properties

- question

R5.1. Unix directory properties

Which of the following **must** always appear in a Unix/POSIX/Linux directory entry (i.e., that is, a directory entry in any such system)?

- An object's size.
- The time the object was created.
- The attributes of the object (e.g., if the file is read-only).
- The inode number of the entry.
- The type of the directory entry.

Select all possible options that apply. There is exactly 1 correct option in the list above. [?](#)

[Save & Grade](#) [Single attempt](#)

[Save only](#)

Additional attempts available with new variants [?](#)

- you need to have the inode number (cuz the directory entry job to provide name to inode mapping)
- (not in the question) you also need the name of the directory entry

R5.2: Find PBN from extent-based file

- question

R5.2. Find PBN from extent-based file

Consider a file system with 2048-byte blocks that uses the **extent-based** file representation.

Assume that:

- The file below contains many hundreds of blocks.
- All numbers (including your answer) are in decimal.

If the first physical block number of this file is 170, what is the physical block number of the logical block number 3?

PBN= integer



Save & Grade Single attempt

Save only

Additional attempts available with new variants

- since it's extent base, it's contiguous
- so LBN 0 → PBN 170, so LBN 3 → PBN 173

R5.3: Understanding File System Structures

- question

R5.3. Understanding File System Structures

Match each description on the left with the term on the right, by selecting the proper entry in the dropdown boxes.



A file whose contents are structures that can only be written by the file system.

a. directory



Metadata that contains a map from LBNs to PBNs.

b. directory entry



A mapping from a name to an inode number.

c. super block

d. inode

Save & Grade Single attempt

Save only

Additional attempts available with new variants

1. a directory → this is a chunk of byte in memory, and only the OS can access this
 - also we know within the chunk of bytes, there are `dirent struct`
2. this is the inode
3. this is the job of a directory entry
 - another question is "A file containing the name of another file" → this is a Soft Link

R5.4: Ext2 Reachable Fork

- question

R5.4. Ext2 Reachable Blocks

In the ext2 file system, how many data blocks can be reached from a double indirect block in a file system that has 1024-byte blocks and uses 4-byte disk addresses?

integer

blocks



[Save & Grade](#) Single attempt

[Save only](#)

Additional attempts available with new variants

- we know that we can fit $1024/4 = 256$ address per block
- so from this double indirect block, we can reach $256 \times 256 = 65,536$ blocks

R5.5: File System Performance Comparison

- question

R5.5. File System Performance Comparison

File System A has a block size that is larger than that of File System B.

Assume the following.

- The file systems store an identical collection of files.
- These files are of many different sizes.
- Any files mentioned below contain many blocks.
- The file systems have been subject to exactly the same workload.
- The file systems are not particularly good at allocating blocks contiguously.

Which of the following statements are true?

- It will take file System A less time to access and read an individual block (given its physical block number) than it will take file System B.
- File system A is likely to have more internal fragmentation than file system B.
- On average, File System A will have higher throughput than file System B.
- It will take file System A more disk accesses to read the same file than it will take file System B.

Select all possible options that apply. There are exactly 2 correct options in the list above.

[Save & Grade](#) Single attempt

[Save only](#)

Additional attempts available with new variants

2. False. It has bigger blocks so reading a singular block would take longer (more data to swap up)
3. True. Bigger block size means more wasted space (especially at the end of files)
4. True. Recall that $\text{throughput} = \frac{\text{data}}{\text{time}}$. The data amount is the same between the implementations, but the time for A will be less because there are less blocks and likely less seeks
5. False. Like number 4, it'll have less seeks

R5.6: Comparing Ext2 and V6

- question
 - variation 1

R5.6. Comparing Ext2 and V6

In the statements below, assume the following:

- Both V6 and Ext2 are running in the same operating system.
- They are accessing an identical set of files.
- The files are located on identical disks.
- If the statement references a file, both file systems are accessing identical files.

Below are a number of statements. Select all the ones that are always **TRUE**.

- Ext2 likely achieves higher throughput on a specific large file than V6.
- V6 is likely to need longer seeks than Ext2 to read a specific large file
- Ext2 can store larger files than V6.
- None of the Above

Select all possible options that apply. There are exactly 3 correct options in the list above. [?](#)

[Save & Grade](#) Single attempt

[Save only](#)

Additional attempts available with new variants [?](#)

1. True. Has a lot to do with block groups:

- in ext2, files are fit into the same block group
 - and within block group there's locality, so there's likely less seeks
 - while v6 allocates free space wherever
- so far large files, ext2 will take less seeks because the file will be in the same block group

2. True. Same thing, locality in ext2 means less seeks

3. True. ext2 has a triple indirect which is already bigger than V6

- variation 2

R5.6. Comparing Ext2 and V6

In the statements below, assume the following:

- Both V6 and Ext2 are running in the same operating system.
- They are accessing an identical set of files.
- The files are located on identical disks.
- If the statement references a file, both file systems are accessing identical files.

Below are a number of statements. Select all the ones that are always **TRUE**.

- V6 is likely to need fewer seeks than Ext2 to read a specific large file
- Ext2 will likely produce more external fragmentation than V6.
- Ext2 likely achieves higher throughput on a specific large file than V6.
- V6 will use less disk space storing a 10-byte file than Ext2.

Select all possible options that apply. There are exactly 2 correct options in the list above. [?](#)

[Save & Grade](#) [Single attempt](#)

[Save only](#)

Additional attempts available with new variants [?](#)

- False. Ext2 has better locality so will need less seeks
- False. They're both block based systems so they shouldn't have external fragmentation
- True. Ext2 has better locality so likely less seeks and thus higher throughput
- True. This file is quite small - likely to be less than a block, so both Ext2 and V6 will use direct pointers - however, in the inode, there are more entries in Ext2 (which maybe empty - but space are still allocated). So in this case V6 takes up less space (basically because it has a smaller block size)

R5.7: Process Basis

- question

R5.7. Process Basics

Below are a number of statements. Select all the ones that are always **TRUE**.

- The hardware can attempt to translate from a virtual address to a physical address on every instruction.
- Exec creates a new address space.
- The virtual address space can be larger than physical memory.
- All hardware has exactly two privilege levels.
- Virtual memory is implemented exclusively in hardware.

Select all possible options that apply. There are exactly 2 correct options in the list above. [?](#)

[Save & Grade](#) [Single attempt](#)

[Save only](#)

Additional attempts available with new variants [?](#)

1. True. Modern processors do this translation using the Memory Management Unit (MMU) for systems with virtual memory

- note: while the hardware is capable of translating a virtual address to a physical address on every instruction that accesses memory, in practice, this does not happen as often thanks to the efficiency of the TLB and the fact that many instructions operate on data within the same page, allowing them to use the same base physical address with different offsets
2. False. `fork` creates a new address space, `exec` just replaces values within the address space
 3. True. This is a fundamental feature of virtual memory, allowing systems to use more memory addresses than there is physical RAM, often utilizing secondary storage (like a hard disk) to compensate
 4. False. Some have more privilege levels than others (can be more than 2, but at least 2)
 5. False. Virtual memory is a partnership between hardware and software

R5.8: Tracing Fork

- question
- part 1
 - it only prints when `pid == 0` which only occurs when it's in the child process
- part 2, 3, 4

What value of the variable `i` will the first child process have when it begins execution?

? ✓ 100%

How many lines of output will be produced?

? ✗ 0%

How many different processes produce output?

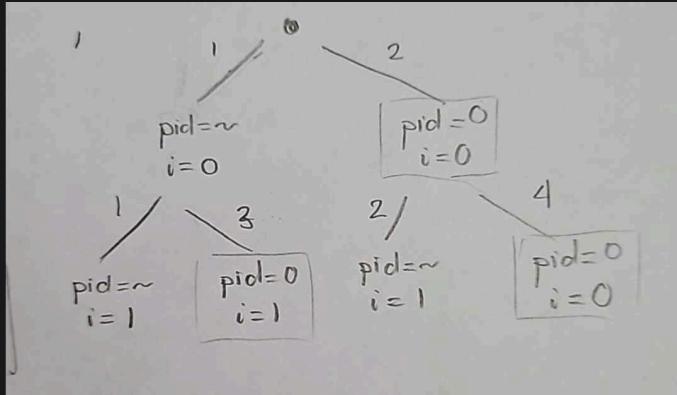
? ✗ 0%

[Try a new variant](#)

- the first process happens when `i = 0`, when the child process gets spawned, it inherits the `i` value as well, hence in the first child process, `i = 0`
- walking through the code
 - first iteration in the parents
 - `i = 0`
 - calls `fork` so we have 1 child process
 - increments `i` and iterate again
 - first child process
 - begins execution immediately after the `fork()` call with `i = 0` (since the value of `i` is copied from the parent at the time of `fork()`)
 - will print `i = 0`

- also has its own loop counter and will proceed to the next iteration, performing its own `fork()` to create a grandchild (child 2)
- second iteration of the parents
 - the parent process continues with `i = 1`.
 - calls `fork()` again, creating another new process (child 3).
- grand child process (child 2)
 - process starts with `i = 1` because it is forked by the first child during the second iteration of the loop
 - it will also print "I am 1\n" because, like all newly created child processes, `fork()` returns 0
- child 3
 - also start with `i = 1` and print "I am 1\n" because `fork()` will return 0 in this new child process
 - (after printing, this child will end because the for loop terminates because `i` is already 1)
- so overall there are 3 child process
 - notice that they print "I am 0/1" so it's only 2 unique things

- another attempt at explaining



- squiggly line next to `pid` means that it's non-zero
- so we can see the `pid == 0` happens 3 times (see boxed)
- those are also 3 different processes (if you go down the left branch - those are same processes aka the parent continuing to be a parent)
 - the numbers next to the edges are counts of processes - so there are 4 total processes (2 child, each with 1 grand child)

(skip R5.9 because it's not that relevant)

R5.10: TLB translations

- question

R5.10. TLB translations

Consider a machine with a 24-bit physical address space with 4096-byte pages running an operating system that supports a 32-bit virtual address space.

Assume that the following table shows you the contents of a TLB.

Virtual page number	Physical page number	Permissions	mode
0x83552	0x6c2	RX	user
0x111a9	0x7e2	RWX	supervisor
0xf3a3c	0x47c	RW	supervisor
0x683ba	0x294	none	user
0xe76cf	0x674	R	user
0xe362c	0x86d	RWX	user

- part 1

Consider the virtual address 0xe76cfce45. What is this address's page offset? Write your answer in hexadecimal.

OFFSET = 0x e45



✓ 100%

- the page is 4096 byte long \rightarrow you need $\log_2(4096) = 12$ offset bits = 3 hexits
- so the first 3 hexit is the offset \rightarrow 0xe45

- part 2

What is the **physical page number** containing this address?

Be sure to format your answer with the correct number of digits, given the size of the physical address space.

PPN = 0x 674



✓ 100%

- this means that the VPN is the remaining of the hexits \rightarrow 0xe76cf
- this VPN maps to 0x674 \rightarrow this is the physical page number (not the actual address)

- part 3

What is the **physical address** of virtual address 0xe76cfce45?

Be sure to format your answer with the correct number of digits, given the size of the physical address space.

PA = 0x 674e45



✓ 100%

- you just have to glue on the offset at the end
- 674 + e45 = 0x674e45

- part 4

Which of the following actions will produce a fault on this address?

- a read request from the operating system for VA 0xe76cfce45
- an execute request from the operating system for VA 0xe76cfce45 ✓
- a write request from the operating system for VA 0xe76cfce45 ✓
- None of the above

Select all possible options that apply.



✓ 100%

- the access permission is Read and for privilege is User
- since all the options involve OS so privilege is not a problem
- however, it's Read only data so we cannot do a Execute and Write

R5.11: Calling Convention - T/F

- variant 1

R5.11. Calling Conventions Basics: T / F

This question is about the construction of y86 stack frames.

When using a base pointer, stack parameters are accessed as positive offsets from %rbp.

True
 False

Save & Grade Single attempt **Save only** *Additional attempts available with new variants*

- False. Local variables are referenced using negative offset from the base pointer, but function variables are referenced a positive offset from the base pointer

- variant 2

R5.11. Calling Conventions Basics: T / F

This question is about the construction of y86 stack frames.

A stack frame that uses a base pointer can access data that cannot be accessed by stack frames that use just the stack pointer.

True
 False

Save & Grade Single attempt **Save only** *Additional attempts available with new variants*

- False. You can access everything using just the stack pointer, just a bit harder because you need to know the size of things

- variant 3

R5.11. Calling Conventions Basics: T / F

This question is about the construction of y86 stack frames.

When not using a base pointer, stack parameters are accessed as negative offsets from %rsp.

True
 False

Save & Grade Single attempt **Save only** *Additional attempts available with new variants*

- False. It's using a positive access

Corrections

Question 1: Directory entry constraints

- question

Question 1: Directory entry constraints

Consider a file system that uses the following structure for its dirents.

```
typedef uint32_t ino_t;
struct dirent {
    ino_t d_ino;
    __uint16_t d_reclen;
    __uint8_t d_type;
    __uint8_t d_namlen;
    char d_name[255+1];
};
#define DT_DIR 4
#define DT_REG 8
#define DT_LNK 10
```

Determine which of the following statements are true.

You may assume that:

- filenames cannot contain nul characters ('\0')
- filenames are not necessarily nul-terminated
- any padding bytes in the record are set to zero
- a directory block contains 4096 bytes

- Typically, a program will need to reorder the bytes in the inode number from little-Endian order prior to using them.
- We determine whether a file is a symbolic link by looking at the type field in the direntry. ✓
- By scanning the bytes in the name itself, we can determine the name length and so derive namelen.
- If reclen were not included in the structure, then given namelen and knowing the C structure padding rules, we can derive reclen. ✓

Select all possible options that apply. There are exactly 2 correct options in the list above. ?

✓ 100%

1. False. If the machine is already a little-endian machine, no reordering is needed - they were made to interpret and understand little-endian numbers
2. True. There's a `d_type` field in the `dirent` struct and we see that `DT_LINK` means that it's a symbolic link
3. False. We don't know how long the name is (it can be shorter than the max length) and the filenames are also not null-terminated, so we don't even know when the name ends.
4. True. Since we know the type we will know the size of each field, and since we know the `nam_len`, we'll know the length of the file name too, all left is to pad to get `rec_len`

Question 2: Relationship between parent and child processes

- question

Question 2: Relationship between parent and child processes

Assume process P is the parent of process C. Which of the statements are always true, immediately after C has been created.

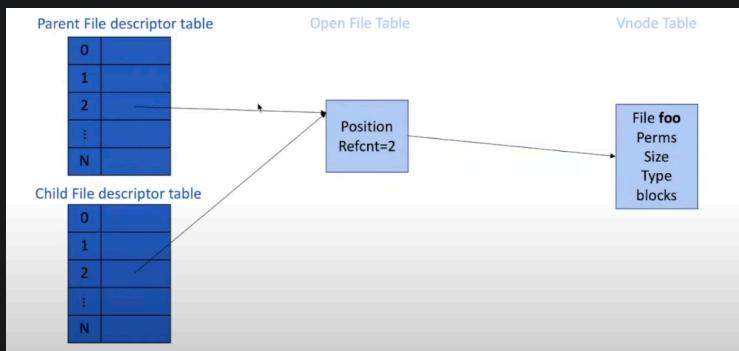
If process P closes a file, the file will be closed in C.
 P and C are executing the same program. ✓
 Any open file table entries that existed before the call to fork have reference count at least 2. ✓
 If process C calls exec, P can see the blocks in the text segment of C.

Select all possible options that apply. There are exactly 2 correct options in the list above. ?

✓ 100%

This question is complete and cannot be answered again.

1. False. Both processes will have their own file descriptor table - so even though they share the same index, one can close and not affect the other (ref count will just decrease)



2. True. Immediately after `fork()`, the child process and the parent process is just continuing from where they were, unless an `exec` is called to make the child run a different process
3. True. Same picture as above, after we call `fork`, the parent's FD table is copied over, and thus so is all the entries (pointers) so now the entries in the OFT have a new thing pointing to it (at least 2 things pointing to it)
4. False. This doesn't really make any sense - processes can't see into each other

Question 3: File System - Links

- question

Question 3: File System - Links

Suppose you have a directory named toys that contains the following files and directories:

Name	Type
beanbag	file
fidget	directory
frisbee	file
top	file

- part 1

What is the link count of the directory toys?

3

?

✓ 100%

- a directory by itself will always have a link count of 2 (the parent, then itself)
- since toys has 1 subdirectory, this subdirectory will point to toys as its parent, thus an additional link count
- so in total we have 3

- part 2

Which of the following statements will always be true after creating a hard link to top named hardlink?

- (a) The number of inodes in use will increase
- (b) The number of inodes in use will stay the same ✓
- (c) The link count of top will decrease
- (d) The link count of top will increase ✓

Select all possible options that apply. There are exactly 2 correct options in the list above. [?](#)

✓ 100%

- False. A hard link means that the newly created directory points to the same inode as another directory → so it reuses an inode and the number of the inode does not increase
- True. Same reason as 1
- False. Since there are now 2 directory entry that points to that inode, the link count will increase
- True. Same reason as 3

- part 3

Which of the following statements will always be true after creating a soft link to beanbag named softlink?

- (a) The link count of beanbag will decrease
- (b) The number of inodes in use will increase
- (c) The number of inodes in use will decrease
- (d) The link count of beanbag will stay the same

Select all possible options that apply. There are exactly 2 correct options in the list above.

100%

1. False. Same reason as 4
2. True. In a soft link, you are basically creating a new file, then putting a path in that file and redirecting people to that new path. Since you are creating a new file - the number of inodes increases
3. False. Same reason as 2
4. True. In a soft link, since you are not referencing a previous inode, but rather creating a new file, the number of inodes will stay the same

Question 4: Adding Fragments to a File System

- question

Question 4: Adding Fragments to a File System

Assume a file system that uses fixed size blocks of size B.

Consider adding the ability to store data from multiple files **in the same block**. We will pick a fragment size, F, such that a fixed number of fragments fit in each block.

For example, if we have a block size of size 4096 and a fragment size of 1024, this would allow us to place four small files (each of size less than 1024) in one block. Alternately, we could place one file with fewer than 2048 bytes in two fragments and then two files with fewer than 1024 bytes in the other two fragments.

Choose the statements that are true of the described fragment-based system. (As always, choose the **best** answers.)

- This is a huge security risk, because a user with data in one fragment could read the data of another fragment that might belong to a different user.
- Fragment-based systems are less likely to be well suited for file systems with fixed-size blocks than those with variable-sized blocks.
- Fragments will be unlikely to cause a file system to run out of space earlier than the system would without fragments.
- It is likely that adding fragments will produce fewer seeks.

Select all possible options that apply. There are exactly 2 correct options in the list above.

0%

- allowing tail end of fails to share blocks (putting files into the empty blocks of other files)
 - say we have 2 files that each require 1.5 blocks → this will in total take up 3 blocks (1 for A, 1 for B, 0.5 of A and 0.5 of B gets put into fragments within block 3)

1. False. It's up to the privileges and permissions that prevents someone accessing data (goes for any file system) - no more usual exposure than it is with any other
2. False. If you can grow and shrink your block size - internal fragmentation wouldn't be a problem (but external would be - but that's unrelated)
 - so we wouldn't even need this fragment shit
3. True. You're using up ends of the block that would otherwise be wasted without fragments
4. True. But a lot of assumptions for these
 - we'll never use up more blocks than we would without fragments
 - so it's possible that we will need less seeks
 - **a ton of deductive reasoning** (since we know first 2 is wrong)

Question 5: Interpreting a Hybrid Inode

- question 1

Question 5: Interpreting a Hybrid Inode

Consider a file system with the following characteristics:

- A hybrid index
- Inodes contain 7 direct pointers and 1 indirect pointer
- Disk block numbers are 4 bytes
- Data blocks are 64 bytes.

Consider the following inode.

other metadata ...	
index 0	451
index 1	221
index 2	338
index 3	565
index 4	430
index 5	316
index 6	516
index 7	725

- part 1

1. What is the maximum size file (in bytes) that we can represent with this index structure? If the answer cannot be determined, enter -1.

max size= 92

bytes



x 0%

- we can store $64/4 = 16$ pointers per block
- using the direct pointers we can reach 7 data blocks, using the indirect pointer we can reach 16 data blocks → so overall can reach 23 blocks
- each data block is 64 bytes, so the max file size is $23 \text{ blocks} \times \frac{64 \text{ bytes}}{\text{block}} = 1472 \text{ bytes}$

- part 2

2. What is the PBN for LBN 3?

PBN= 565 ? ✓ 100%

- this is just index 3 in the inode, which is PBN 565

- part 3

3. What is the PBN of the indirect block for this file?

PBN= 725 ? ✓ 100%

- this is index 7 (but the 8th entry) in the inode because the first 7 entries are direct pointers
- so it's PBN 725

- part 4

4. At what index in the indirect block will you find the PBN for LBN 16?

index= 9 ? ✓ 100%

This question is complete and cannot be answered again.

- we know that this will be in the direct block, so we can do $16 - 7 = 9$
- and within the indirect block, index 9 will get you the PBN for it
- (can also think of it as $9 \bmod 16^0$)

Question 6: ext2 versus V6 directory entries

- question

Question 6: ext2 versus V6 directory entries

Given a large directory of names, all of which are between five and eight (inclusive) bytes in length, which file system will consume more space storing this directory?

- V6
- ext2
- V6 cannot represent such a directory
- ext2 cannot represent such a directory
- Both file systems will consume the same amount of space

✓ 100%

For reference, here are the v6 and ext2 directory entry structure definitions:

```
struct dirent {          // v6
    int d_ino;        // 16 bits
    char d_name[14];
};
```

```
struct ext2_dir_entry { // ext2
    __le32 inode;      /* Inode number */
    __le16 rec_len;    /* Directory entry length */
    __le16 name_len;   /* Name length */
    char name[256];    /* File name */
}
```

This question is complete and cannot be answered again.

- we can think of this in the worst case and in the best case
 - best case is names are 5 bytes each
 - worst case is names are 8 bytes each
- v6
 - it does not matter whether it's best case or worst case - v6 `dirent` struct is padded to take up 16 bytes
- ext2
 - best case: we have $4 + 2 + 2 + 5 = 13$ bytes but we have to pad to be 4-aligned → it'll be 16 bytes
 - worst case: we have $4 + 2 + 2 + 8 = 16$ bytes which does not need padding
- you can see that in any case, v6 and ext2 takes the same amount of data

Question 7: File System Data Structures

- question

Question 7: File System Data Structures

Consider a system on which you have both ext2 and v6 file systems mounted (and no other file systems). For each data structure listed below, select the answer that best indicates the number of **different data structure definitions** there will be for that data structure.

Note: "No answer" is never a correct answer. We include it so that the problem will be gradable even if you do not answer all the questions.

a structure describing an open file table entry ✓ 100%

a structure describing the target of a file pointer, i.e., FILE *

✓ 100%

a structure used for allocating inodes ✓ 100%

a struct stats ✘ 0%

This question is complete and cannot be answered again.

(the options are There will be one, There will be 2, There will be more than 2, Not enough information to know)

1. This structure is usually managed by the OS - and the point is to provide abstraction despite what FS might be mounted, so there should be 1
 - OS maintains a collection of open file entries that is the open file table → so they all have to be the same type, so there's only 1 OF entries
 - (the key here is asking if the OS takes care of it or FS)
 2. This is generally a part of the operating system's standard C library and not specific to any file system. So there should only be 1
 - file type are defined by the OS so it's just 1
 3. We know from class that v6 uses a linked list structure to allocate inodes while ext2 uses a bitmap structure so there will be 2
 - also inodes are laid out on disk - so it's based on the file system
 4. The struct can to express the diff values that the FS will take (i.e how many blocks are in the FS, etc), but we need the OS need to have a consistent structure to differentiate between the two
 - but basically, the OS uses this, so just have 1
- TODO:

QP5.5. File System Data Structures

Consider a system on which you have both ext2 and v6 file systems mounted (and no other file systems). For each data structure listed below, select the answer that best indicates the number of **different data structure definitions** there will be for that data structure.

Note: "No answer" is never a correct answer. We include it so that the problem will be gradable even if you do not answer all the questions.

a structure describing the target of a file pointer, i.e., FILE *

There will be one.

↙ ✓ 100%

a structure describing a directory entry

There will be two.

↙ ✓ 100%

a structure describing a superblock

There will be two.

↙ ✓ 100%

a struct statfs

There will be one.

↙ ✓ 100%

Try a new variant

Question 8: Choose-Your-Own-Address TLB Adventure

- question

Question 8: Choose-Your-Own-Address TLB Adventure

Consider a machine with a 36-bit physical address space with 4096-byte pages running an operating system that supports a 44-bit virtual address space.

Assume that at most one virtual page can map to any physical page (though this is not necessarily true in general). Assume that the following table shows you the contents of the TLB.

Virtual page number	Physical page number	Permissions	mode
0x9c48104c	0x5e4001	R	user
0xa38aca52	0xa6231c	WX	supervisor
0x74659bb8	0x443da5	RWX	supervisor
0x476a3337	0x646c78	RWX	supervisor
0xf81b7108	0x328cd1	RW	supervisor

- part 1

Give a **virtual address** in hexadecimal (of the appropriate width) that definitely **cannot** be read by the OS. (We are asking for a full address, not a page number.)

0x a38aca52000

?

↙ ✓ 100%

- so we want to pick an entry in the TLB that doesn't allow read → the 2nd entry
- this virtual page number if this is **0xa38aca52**
 - to get a full address, you have to add an offset (12 bits or 3 hex digits because $\log_2(4096) = 12$)
 - so we could have chosen any random combination of 3 hexadecimal digit but I went safe and chose **000**
 - so overall we have **a38aca52000**
- part 2

Give a **virtual address** in hexadecimal (of the appropriate width) that **may or may not** be able to be read by the OS because the TLB has insufficient information to tell. (We are asking for a full address, not a page number.)

0x	d4659bb8000	?	✓ 100%
----	-------------	---	--------

- this is basically asking for a virtual address that's not in the TLB
- an easiest thing to do is just take any VPN that's in the TLB and change 1 digit
 - this works because the VPN omits the offset, which spans the entirety of the page
 - so if you change any 1 digit in the VPN, we're at a completely different page

Question 9: Tracing Fork

- question

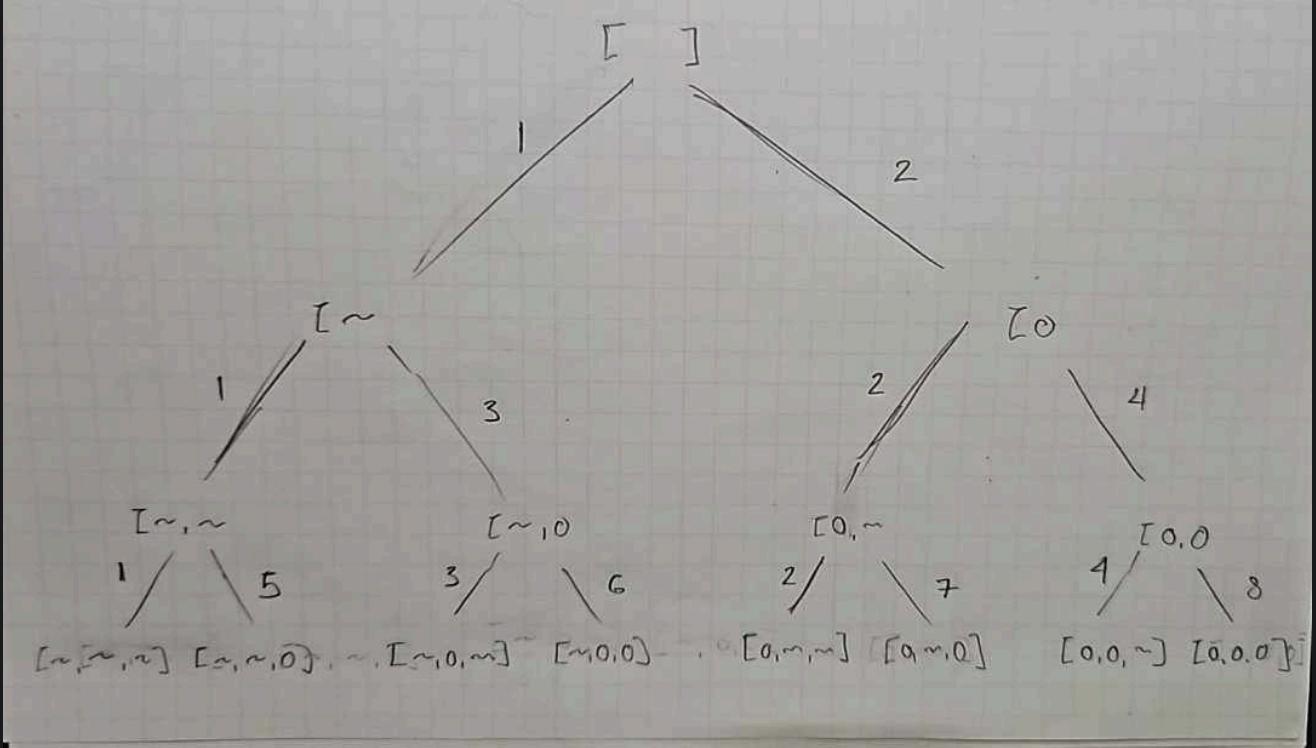
Question 9: Tracing Fork

Consider the following C code.

```
int main(int argc, char **argv) {
    int i = 0, pid[3];
    pid[i++] = fork();
    pid[i++] = fork();
    pid[i++] = fork();

    if (pid[1] == 0) {
        for (int j = 0; j < 3; j++) {
            printf("pid[%d] = %d\n", j, pid[j]);
        }
    }
}
```

- we can draw the following



- the number along the edges are the **count** of processes, **not the pid**
- you can see that squiggly line is non-zero **pid**

- part 1

How many processes will be created during the execution of this program? Include in your count the initial process.

8



✓ 100%

- part 2, 3, 4

How many lines in the output will contain the following string? `pid[0] = 0`

2



✓ 100%

How many lines of output will be produced?

12



✓ 100%

How many distinct processes will produce output?

4



✓ 100%

- not that we enter the for loop (and thus print) 4 times (for process count 3, 6, 4, 8 - because `pid[1] = 0` for them)
- of process 3, 6, 4, 8 that runs → only process 4 and 8 has `pid[0] = 0` so we print it twice
- for the 4 process that enter the `for` loop and prints, they each run 3 times and print 3 times → we print 12 (not all unique but they'll still print)

- again, 4 processes enter the `for` loop so 4 distinct process is producing the output

Question 10: When does the OS get to run

- question

Question 10: When does the OS get to run?

Which of the following events will cause control to transfer from a user program to the operating system?

- A program divides a number by 0 ✓
- A program is in an infinite loop
- A program calls `open` ✓
- A program experiences a cache miss

Select all possible options that apply. There are exactly **2** correct options in the list above. [?](#)

✓ 100%

This question is complete and cannot be answered again.

- True. This causes an exception, which hands control over to the OS (it's a trap)
- False. When the timer runs out, the OS gets to run, but while the program is in an infinite loop, there's no way for the OS to know about this and to intervene
- True. This is a system call
- False. This is a hardware thing

Question 11: Hardware or Software

- question

Question 11: Hardware or Software

Below are a number of different events or actions that could take place in a system. Each one can be categorized as being performed by hardware, the operating system, or user software. Select all items that are categorized as being **performed by hardware**.

- Determining if the running process is the parent process or child process after a `fork()`.
- Determining where in the buffer cache (memory used as a cache of disk blocks) to place a block from a file.
- Determining which L1 cache line to check on a memory read. ✓
- Determining how long the processor stalls on a control hazard. ✓
- Determining which page to evict from the buffer cache (memory used as a cache of disk blocks).
- Determining how long the processor stalls on a data hazard. ✓

Select all possible options that apply. There are exactly **3** correct options in the list above. [?](#)

✓ 100%

This question is complete and cannot be answered again.

- False. We check the `pid` to know if it's a parent or child → happens in software

2. False. Bringing in pages and managing the TLB is a OS thing
 - i.e when there's a page fault, OS takes care of it
3. True. Caching happens on the hardware level
4. True. The hardware waits for the info to be available before proceeding
 - this also happens way too often for us to use the OS every time
5. False. Again, this is a OS thing
6. True. Same as 4